Lecture 15: *Introduction to the information bottleneck approach*

The main results of information theory and coding are more substantial:

source coding Th'm: gives bounds on compressibility of signals given their entropy

channel coding Th'm: describes the limit of the information rate of a noisy channel, and the maximal efficiency of error-correcting codes

Recall our introduction to channel coding:

Recall our introduction to channel coding:

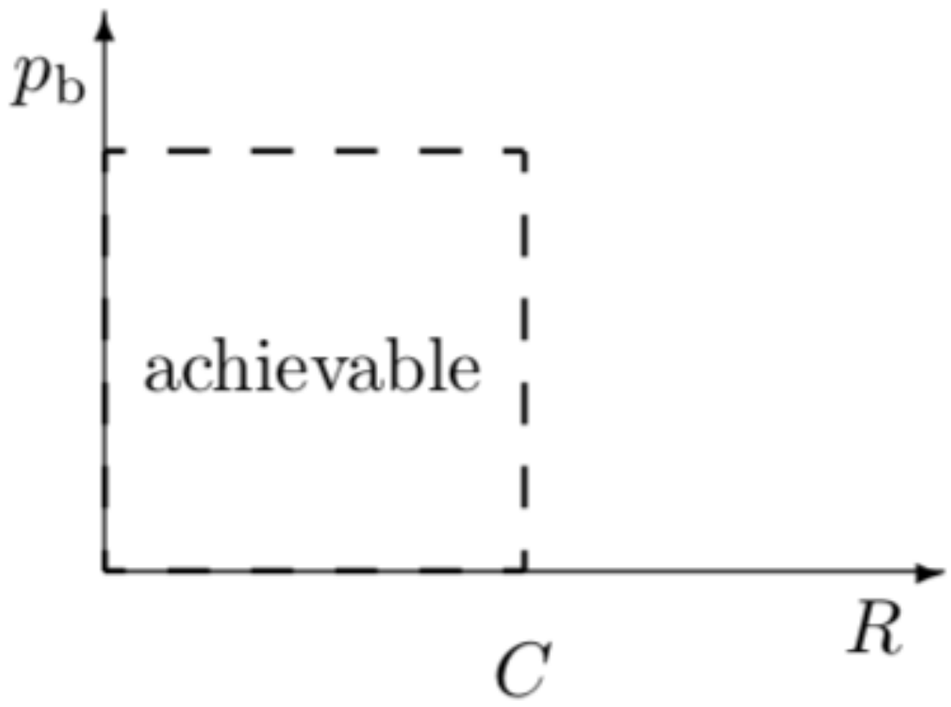

Recall our introduction to channel coding:

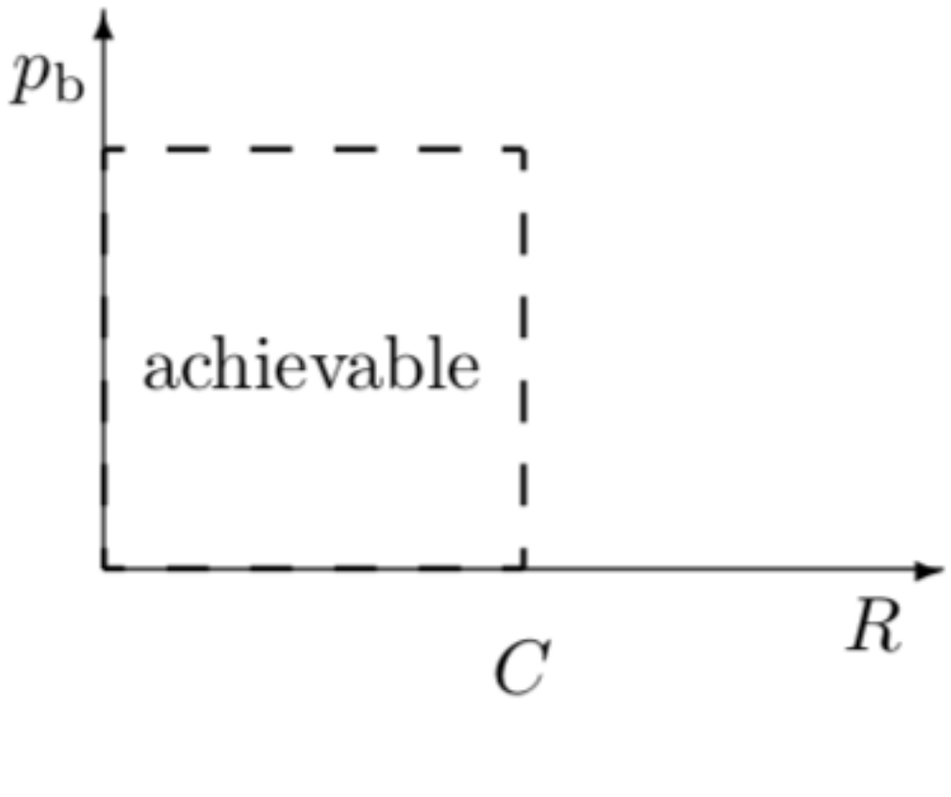

Recall our introduction to channel coding:

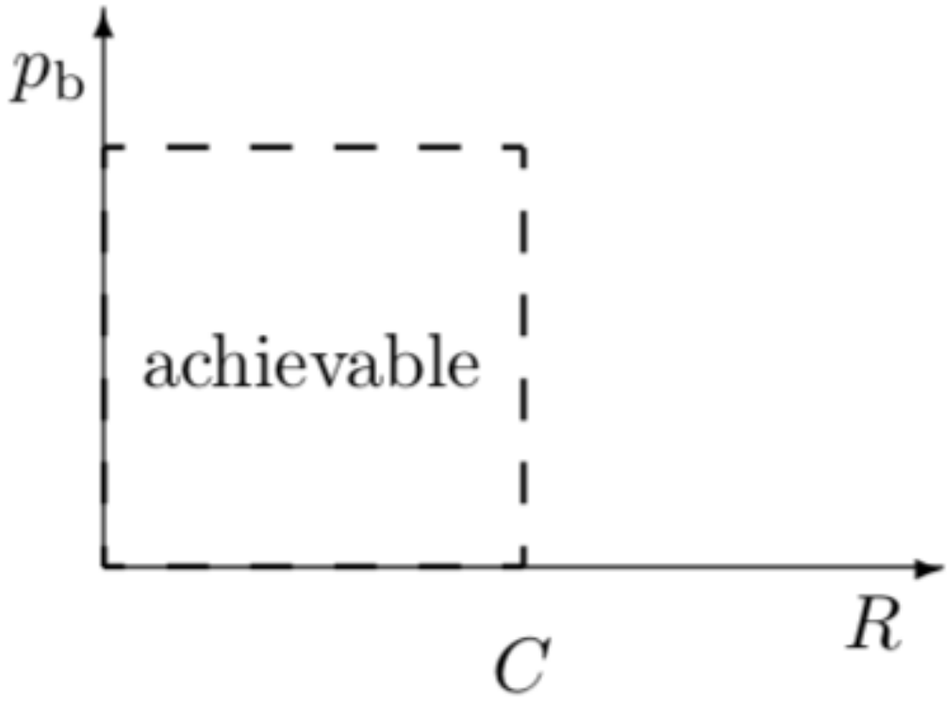

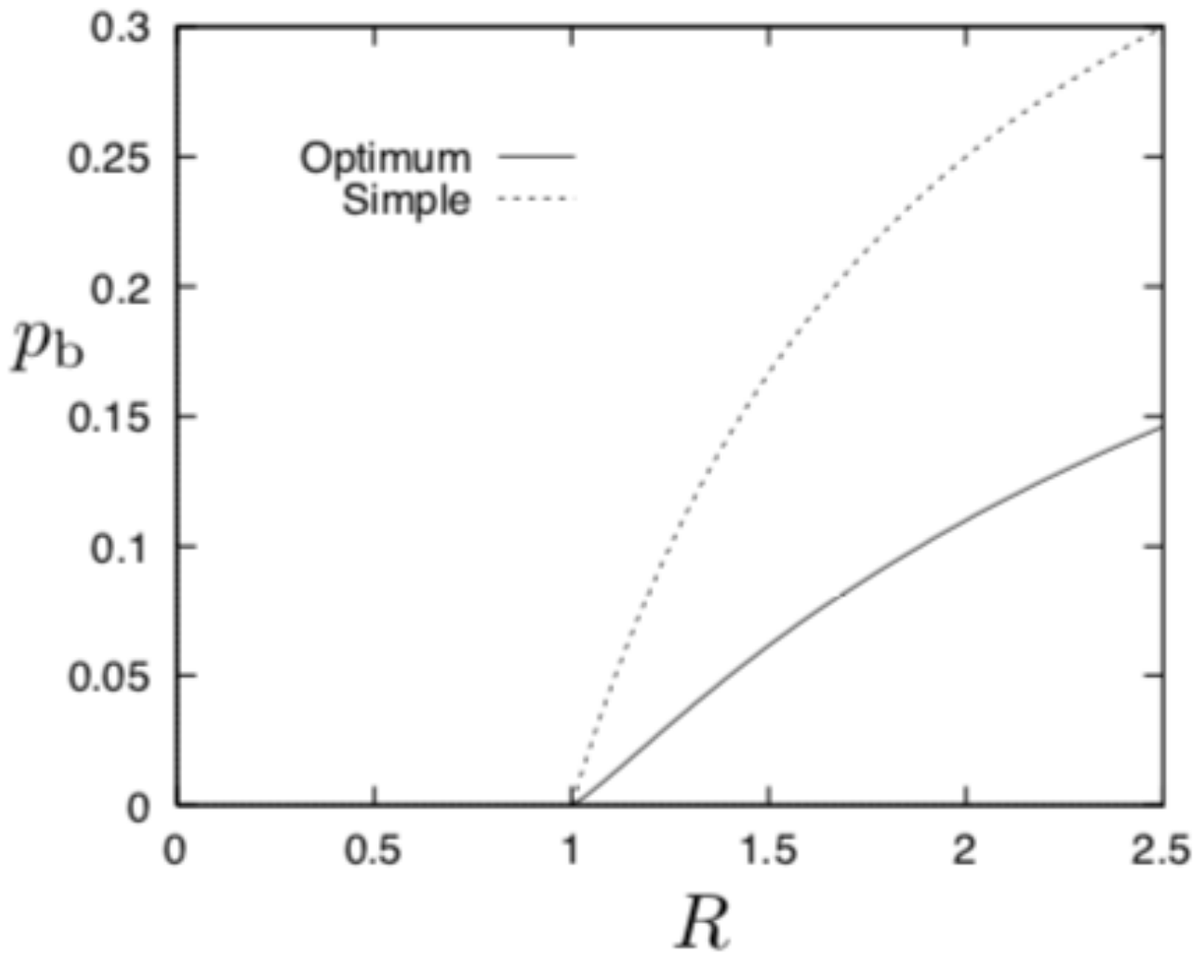

Processing delays mean the brain has to make predictions:

Processing delays mean the brain has to make predictions:

Towards more natural motion stimuli:

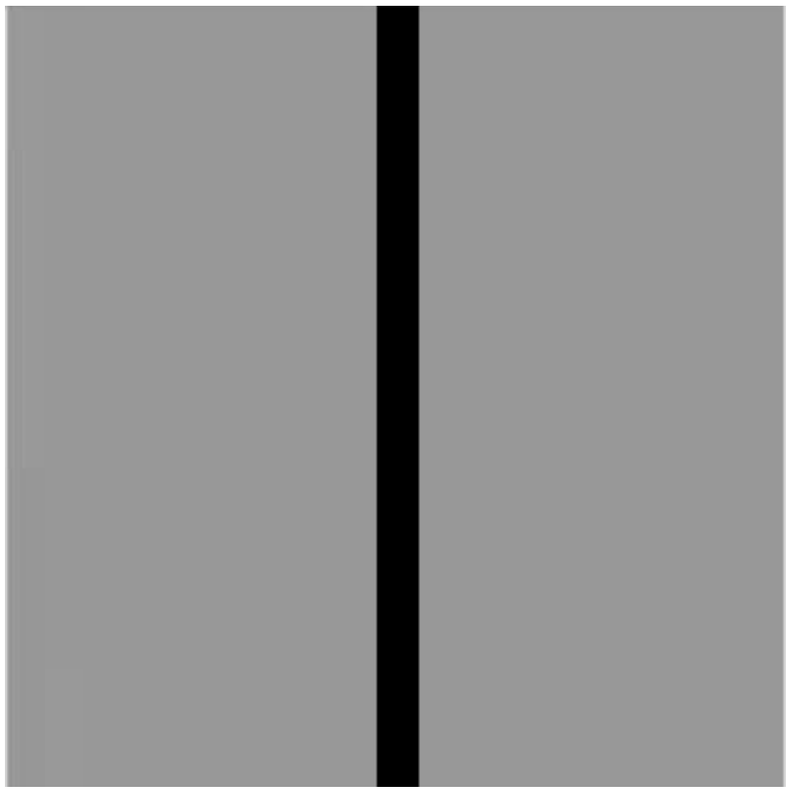

Towards more natural motion stimuli:

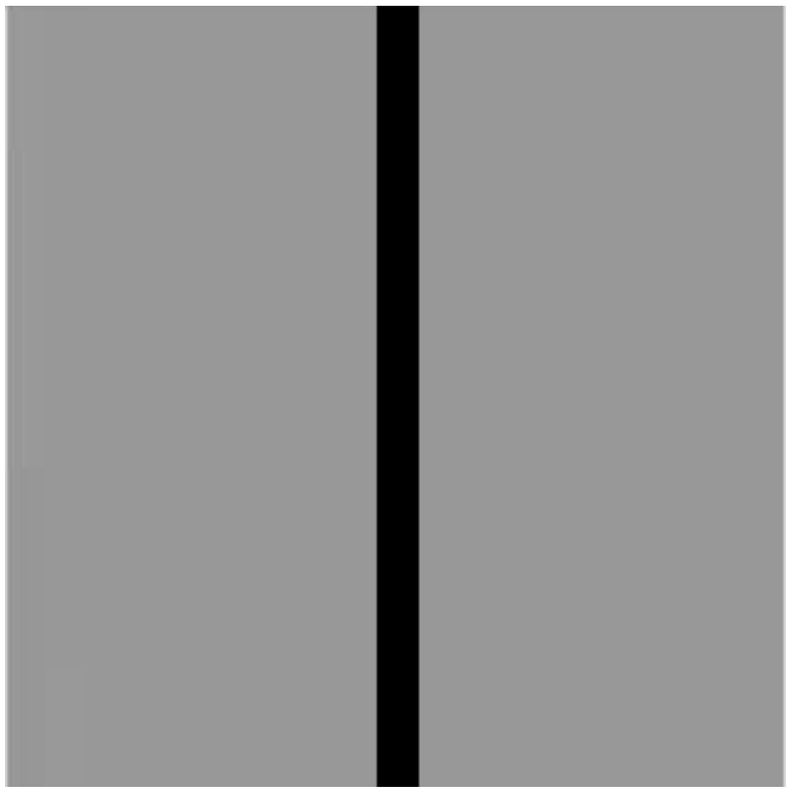

The retina is a complex piece of neural tissue:

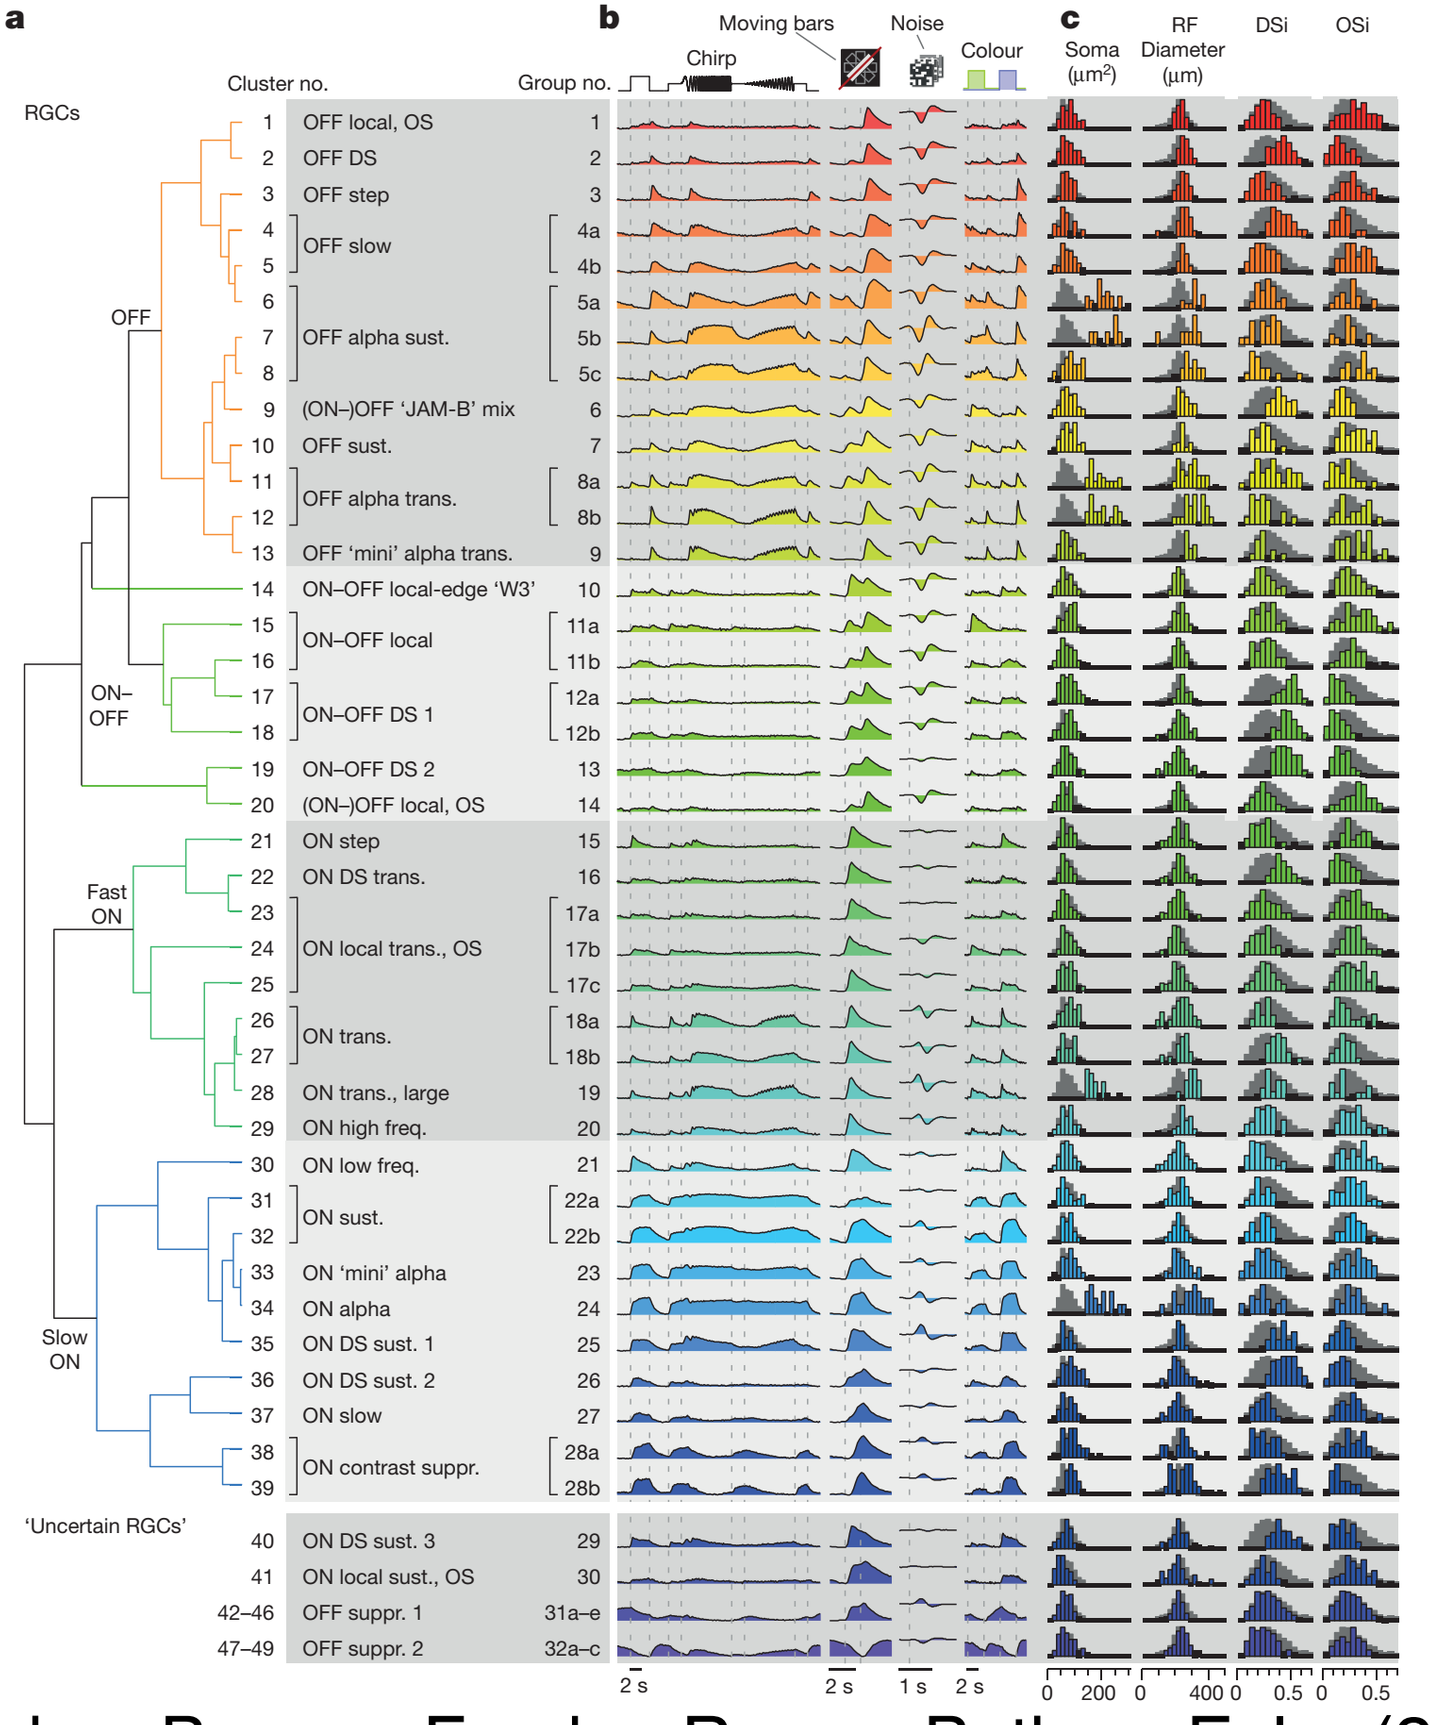

Baden, Berens, Franke, Roson, Bethge, Euler (2016)

The retina does a lot of interesting things:

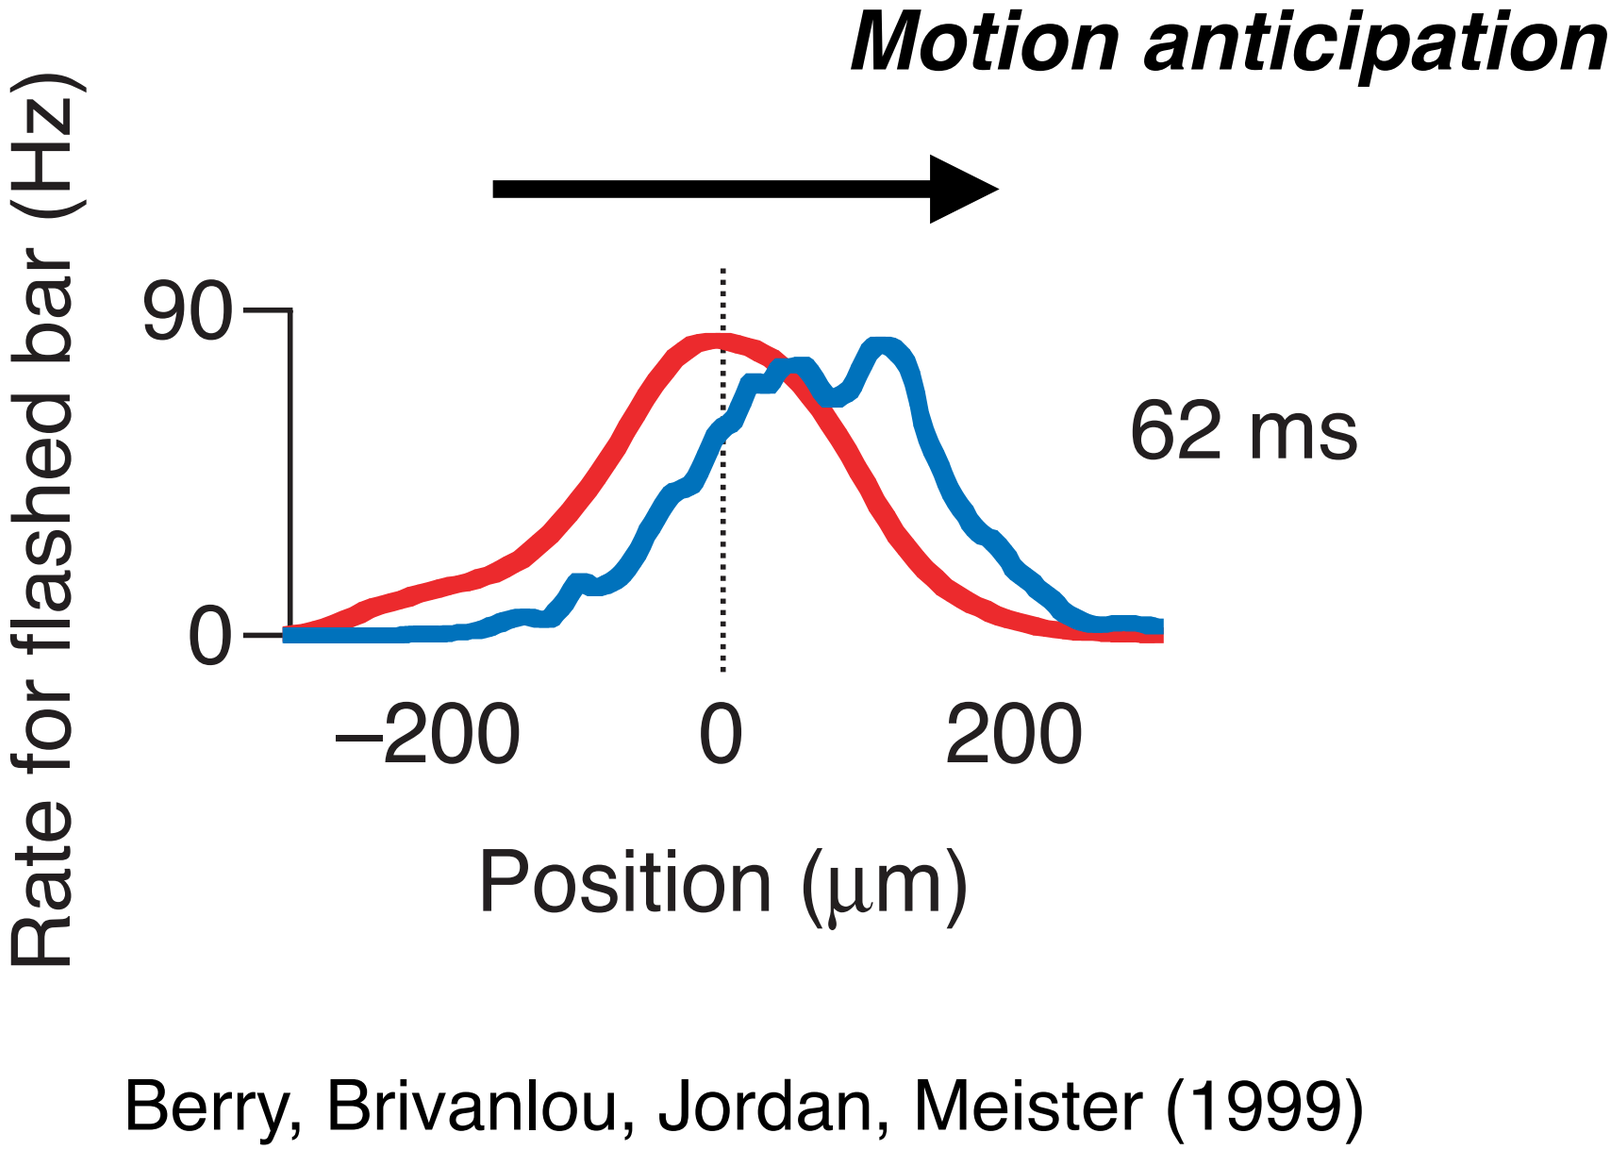

Berry, Brivanlou, Jordan, Meister (1999)

The retina does a lot of interesting things:

Reversal response

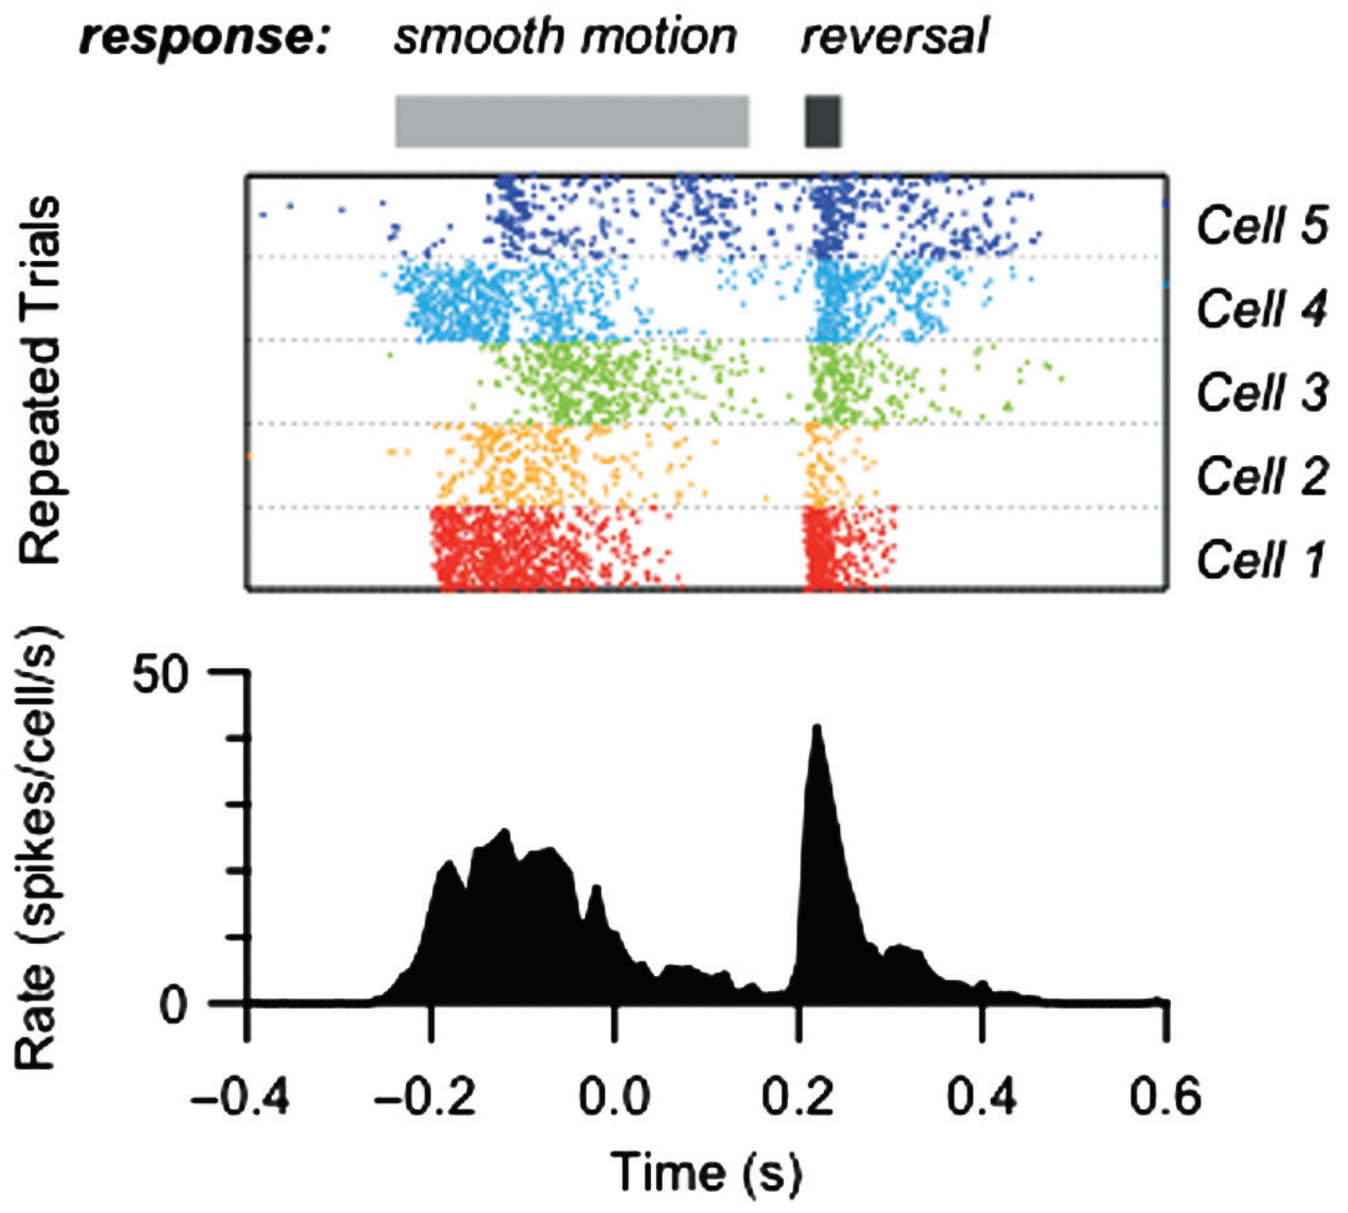

Schwartz, Taylor, Fisher, Harris, Berry (2007)

The retina does a lot of interesting things:

Lag normalization

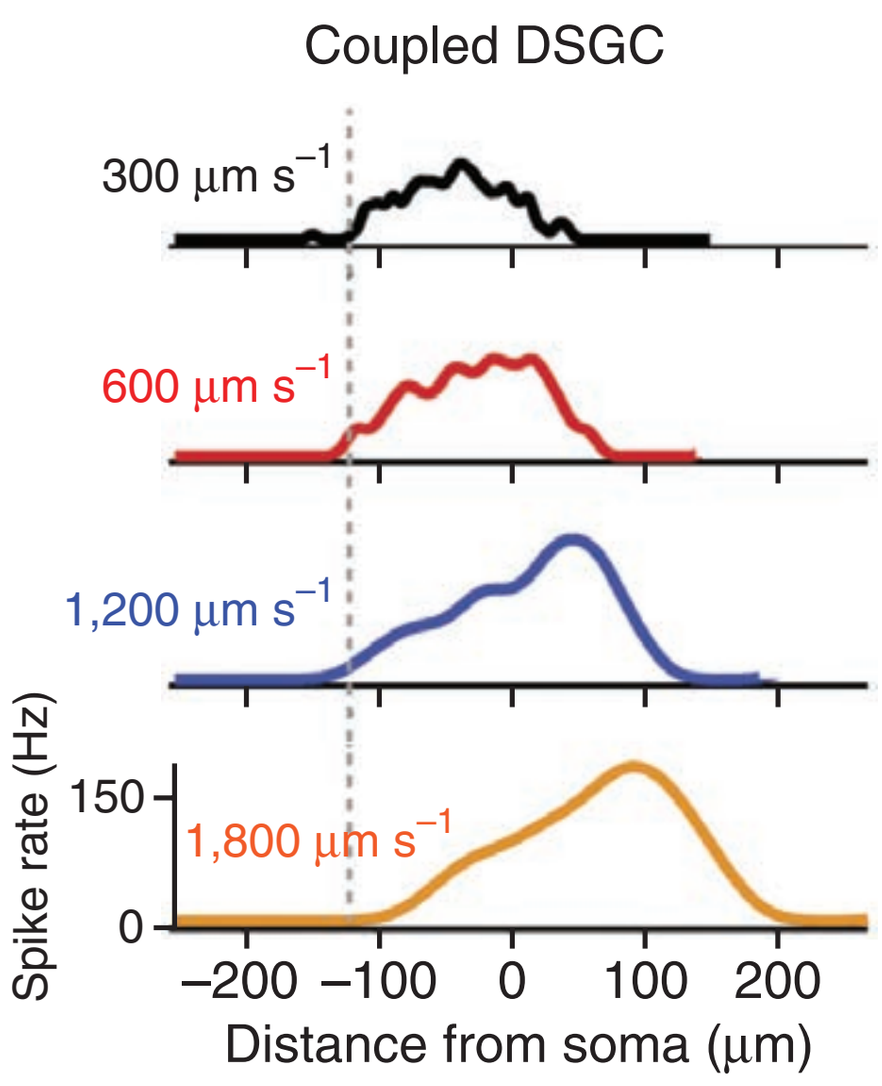

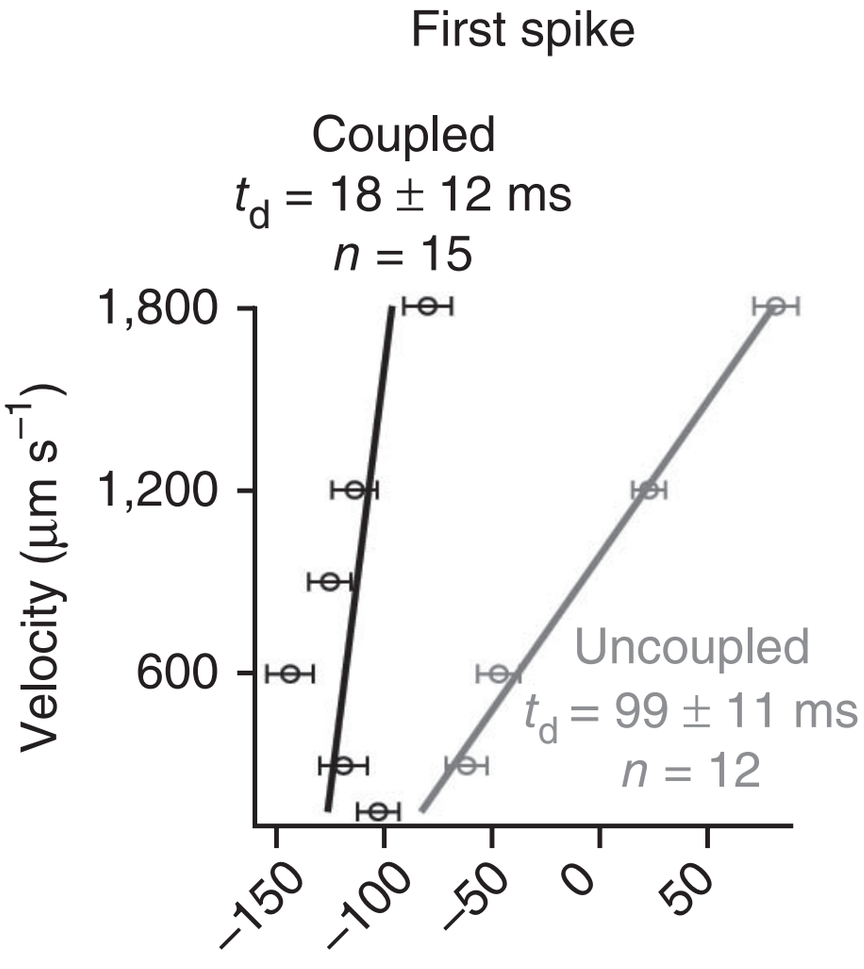

Trenholm, Schwab, Balasubramanian, Awatramani (2013)

A bar stimulus with both predictable and non-predictable motion components:

A bar stimulus with both predictable and non-predictable motion components:

A bar stimulus with both predictable and non-predictable motion components:

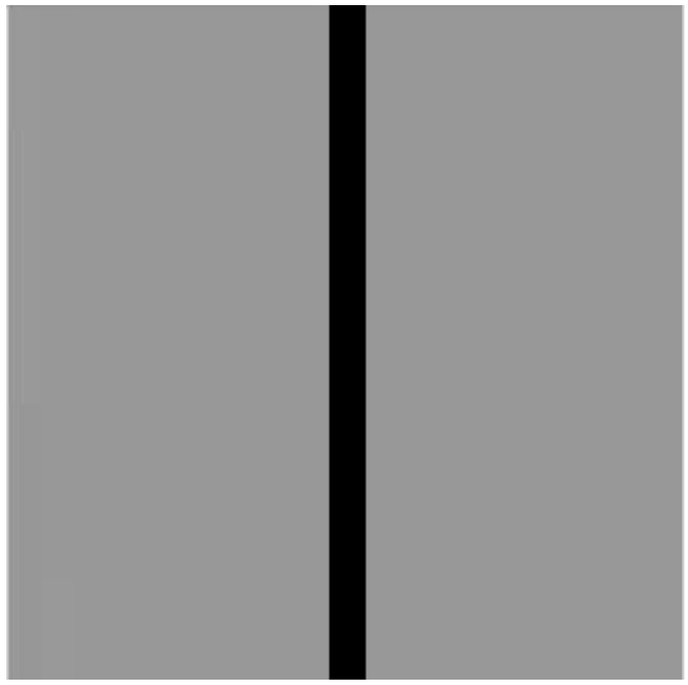

$$\frac{dv}{dt} = -\frac{v}{\tau} + D^{1/2}\Gamma(t) - \omega_0^2 x$$

A bar stimulus with both predictable and non-predictable motion components:

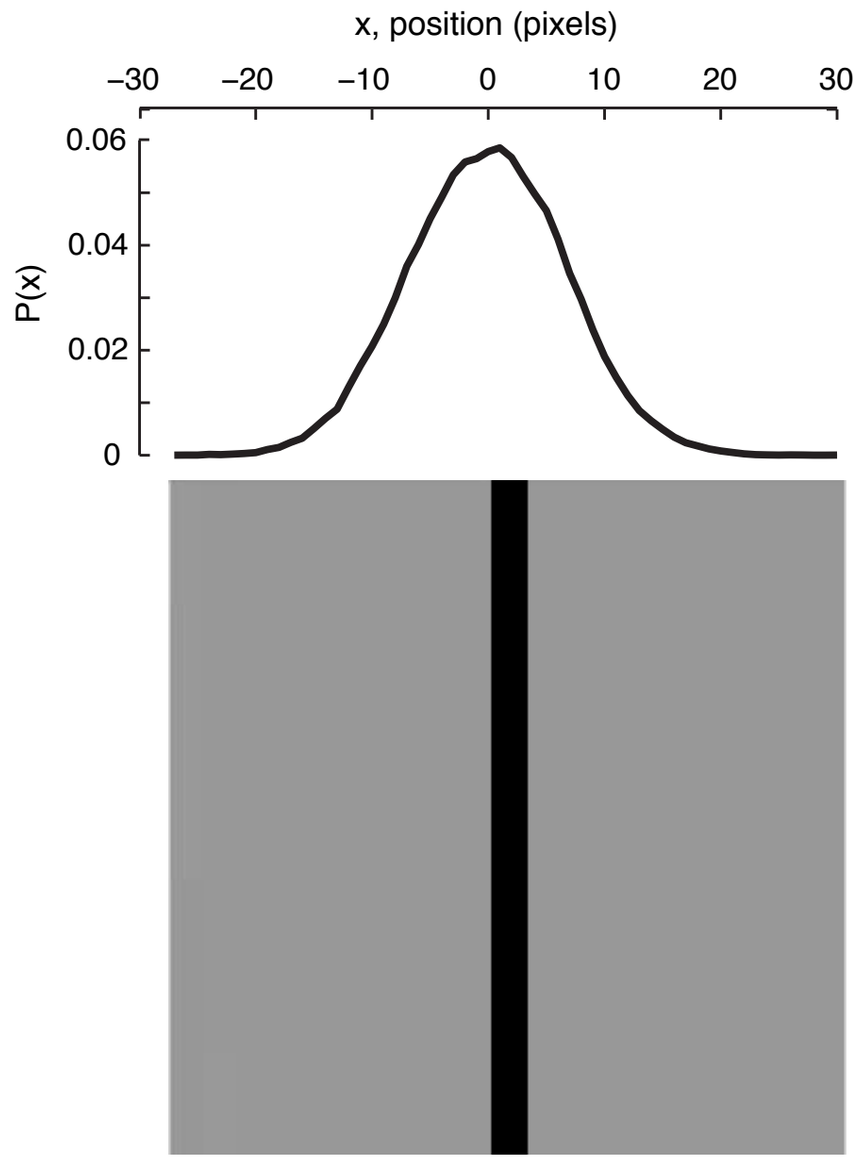

$$\frac{dv}{dt} = -\frac{v}{\tau} + D^{1/2}\Gamma(t) - \omega_0^2 x$$

A bar stimulus with both predictable and non-predictable motion components:

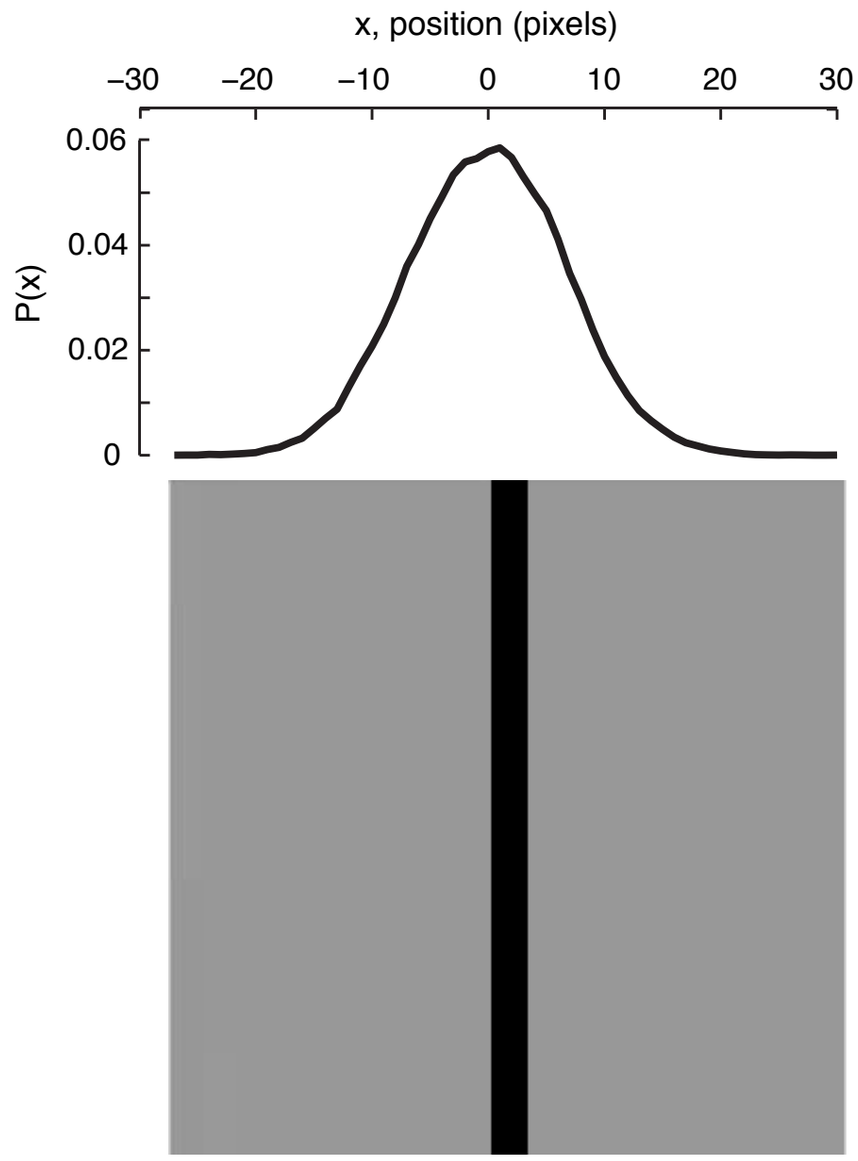

$$\frac{dv}{dt} = -\frac{v}{\tau} + D^{1/2}\Gamma(t) - \omega_0^2 x$$

A bar stimulus with both predictable and non-predictable motion components:

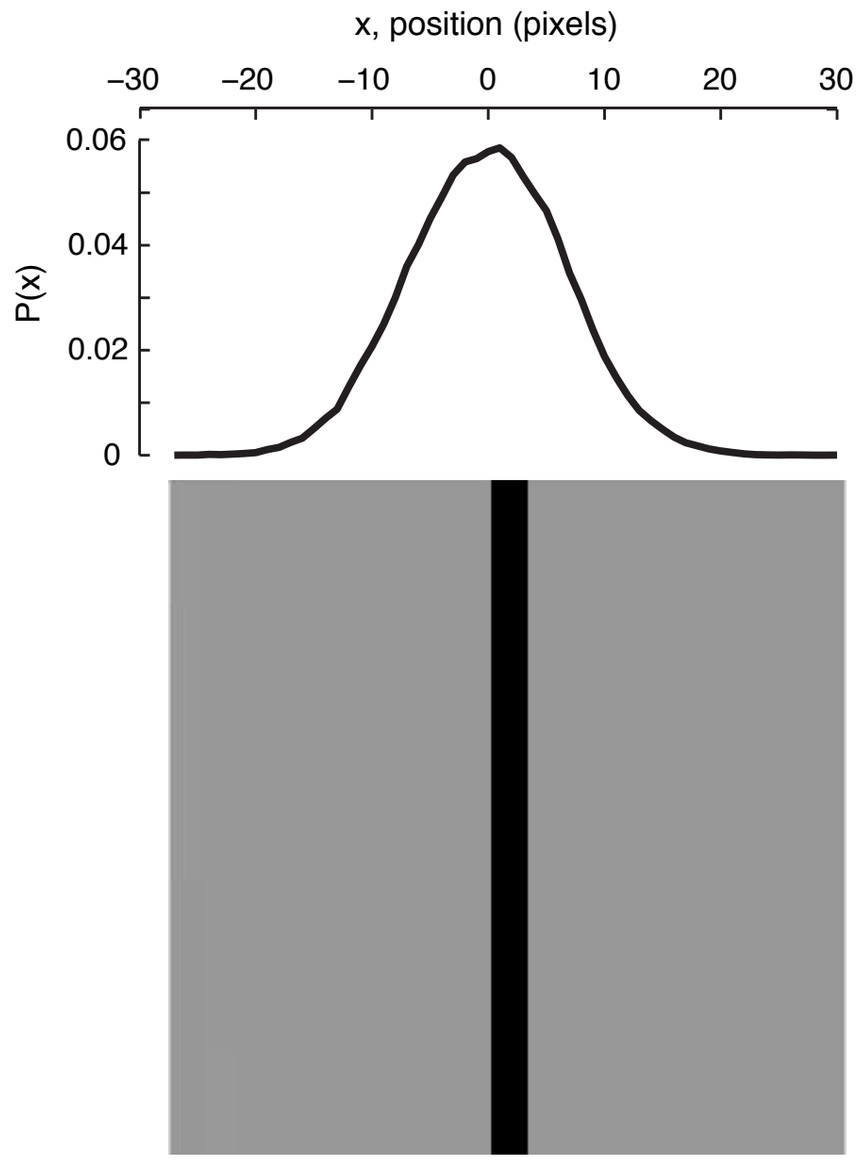

$$\frac{dv}{dt} = -\frac{v}{\tau} + D^{1/2}\Gamma(t) - \omega_0^2 x$$

A bar stimulus with both predictable and non-predictable motion components:

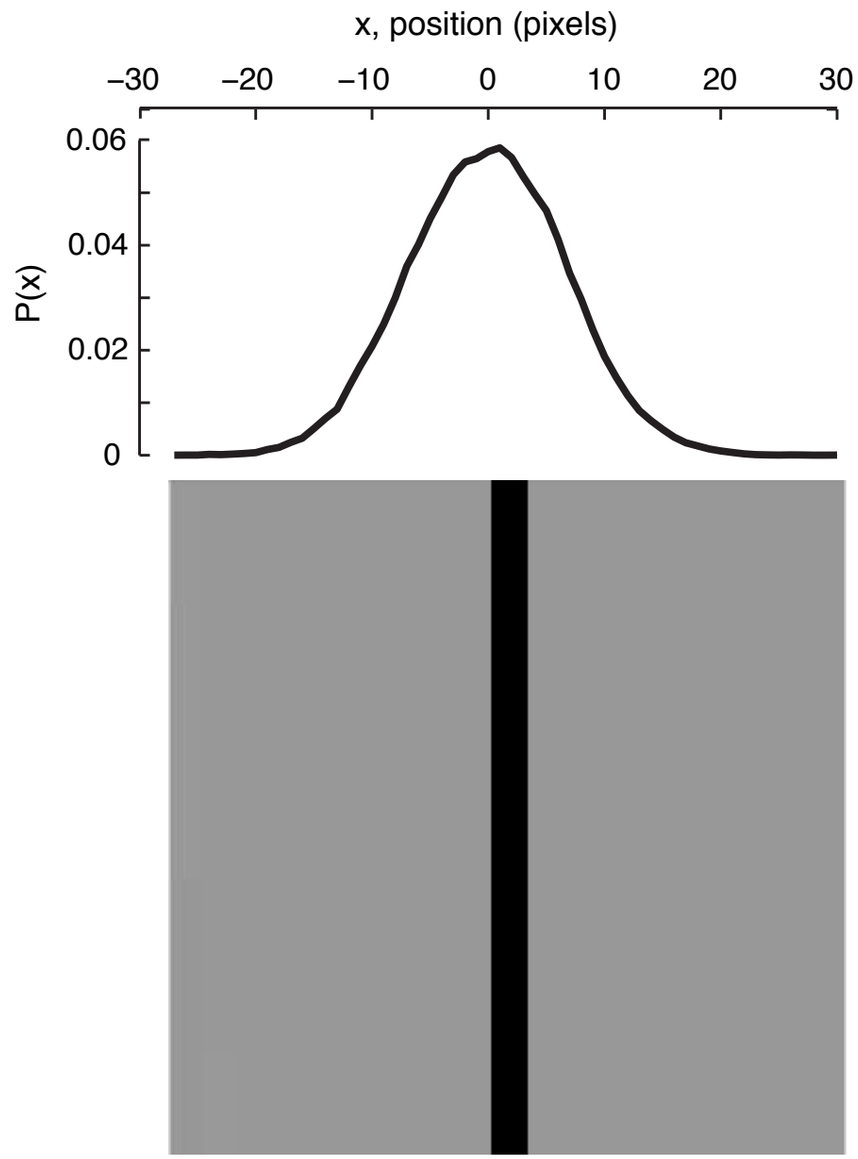

$$\frac{dv}{dt} = -\frac{v}{\tau} + D^{1/2}\Gamma(t) - \omega_0^2 x$$

A bar stimulus with both predictable and non-predictable motion components:

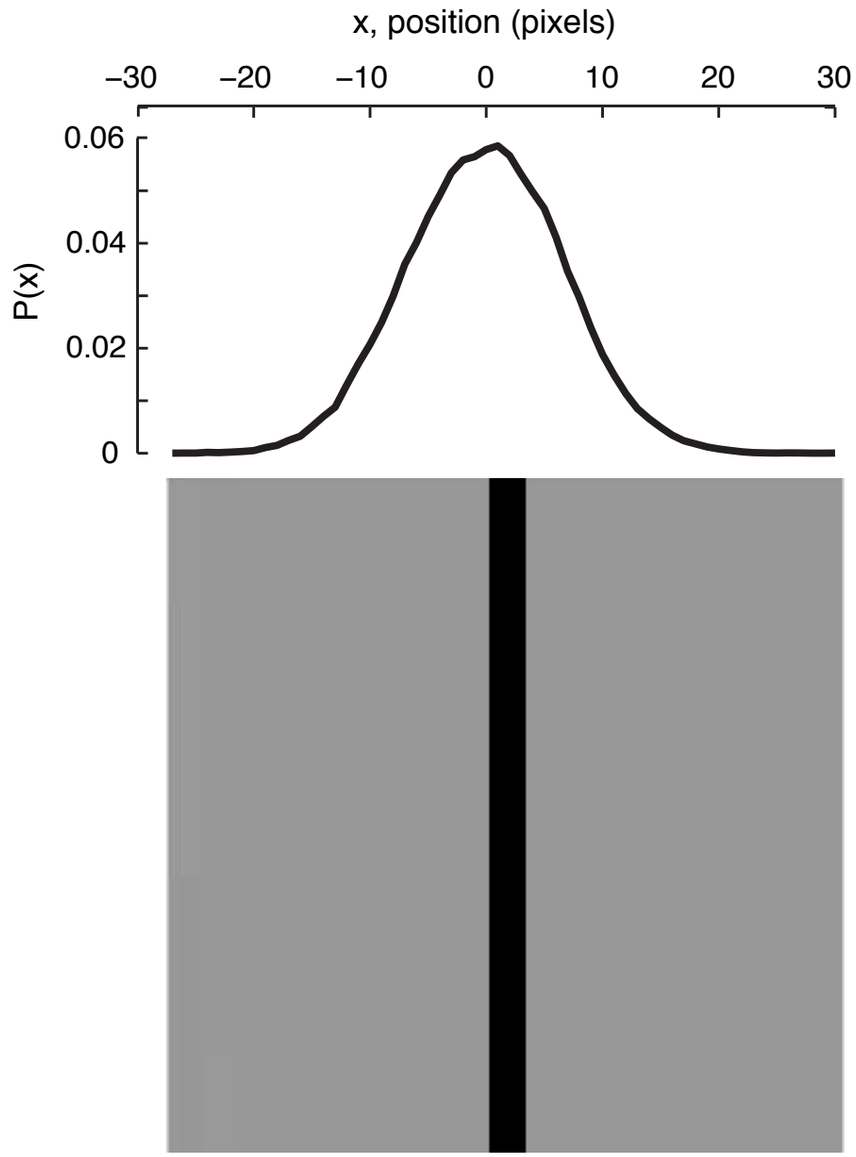

$$\frac{dv}{dt} = -\frac{v}{\tau} + D^{1/2}\Gamma(t) - \omega_0^2 x$$

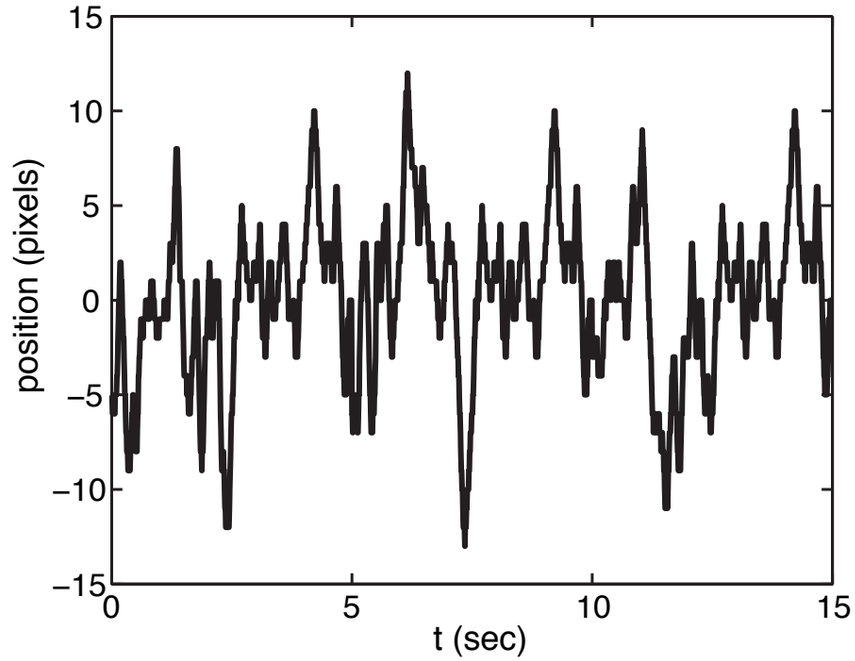

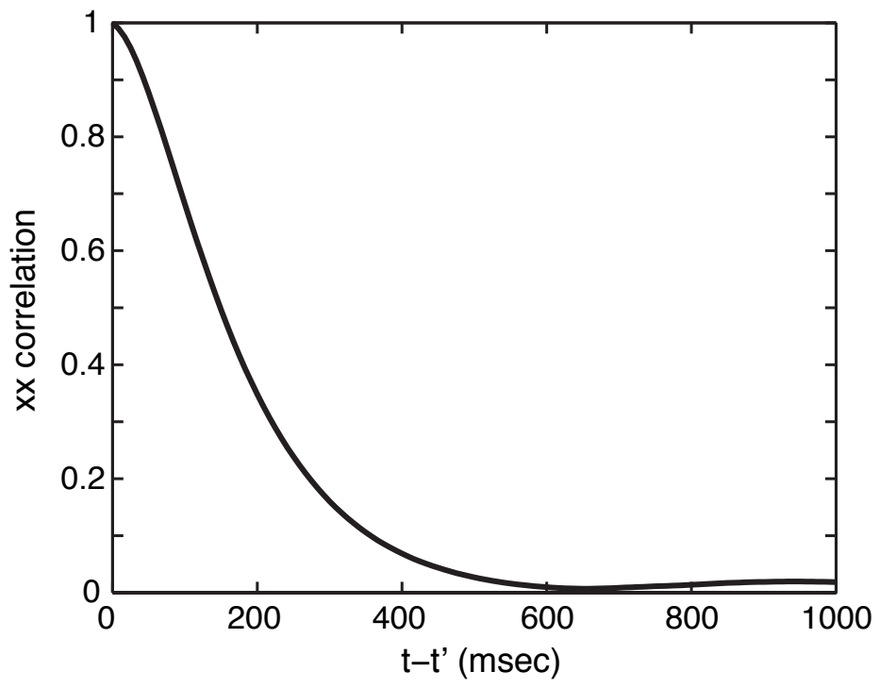

Recording from the retina:

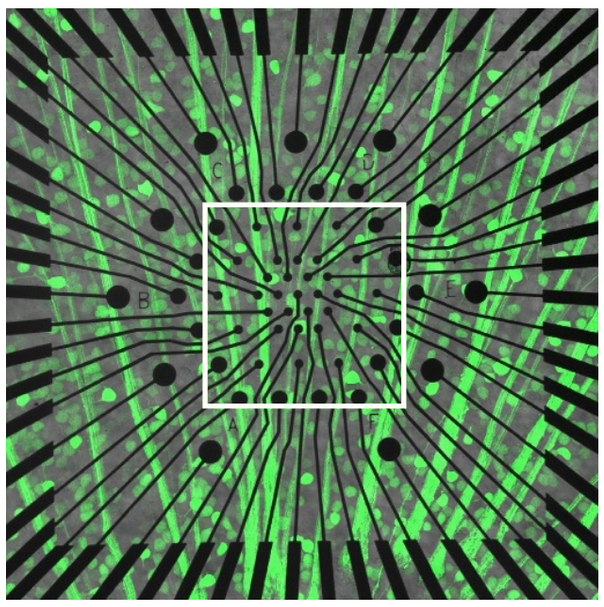

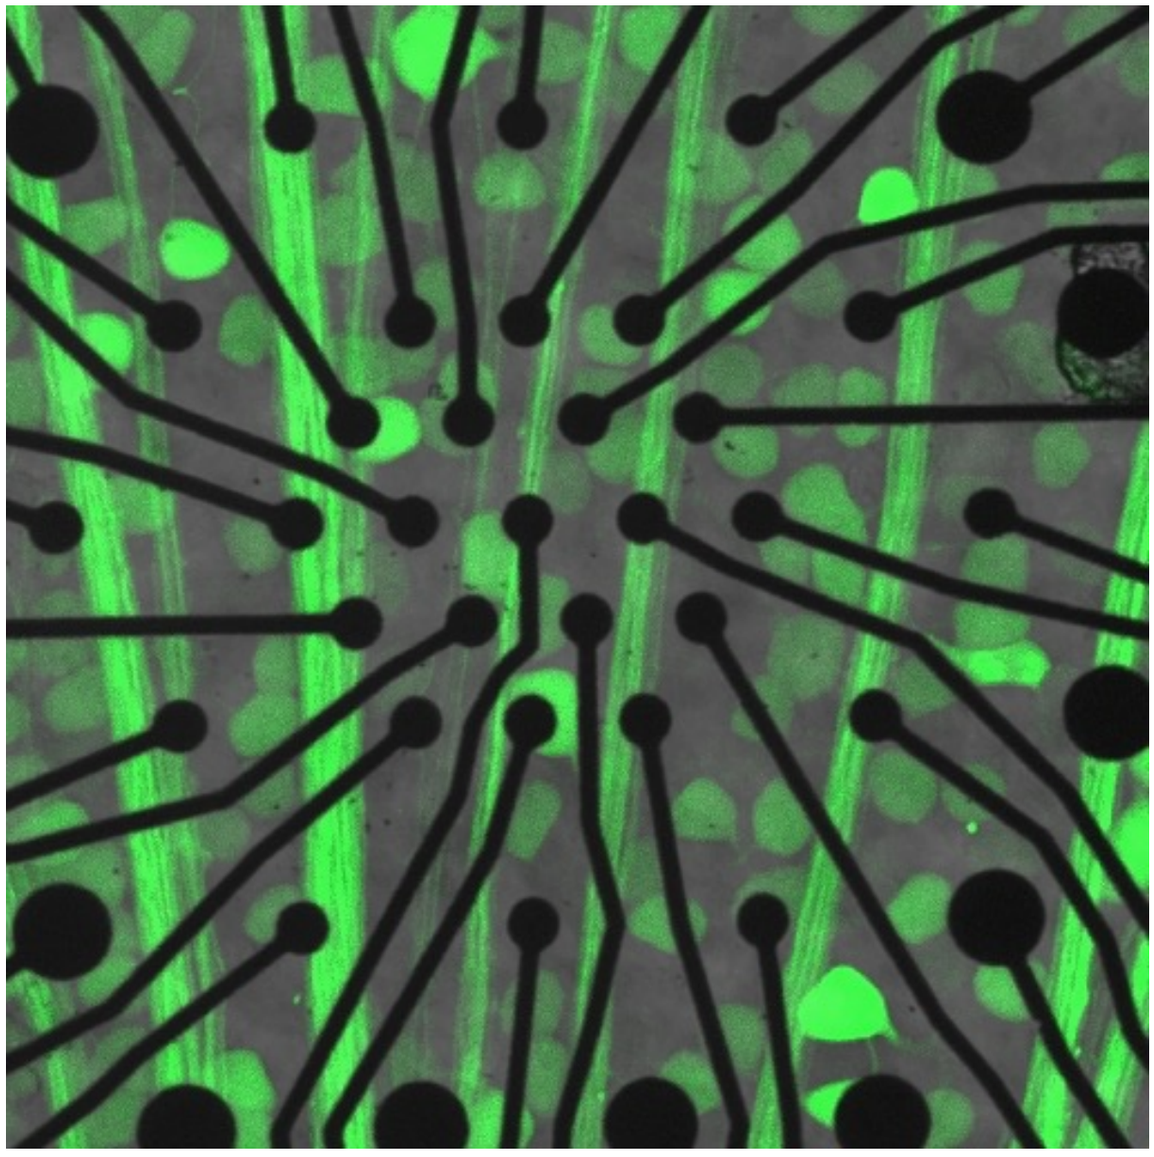

images from the Berry Lab, Ronen Segev

We bin time to create binary spike 'words':

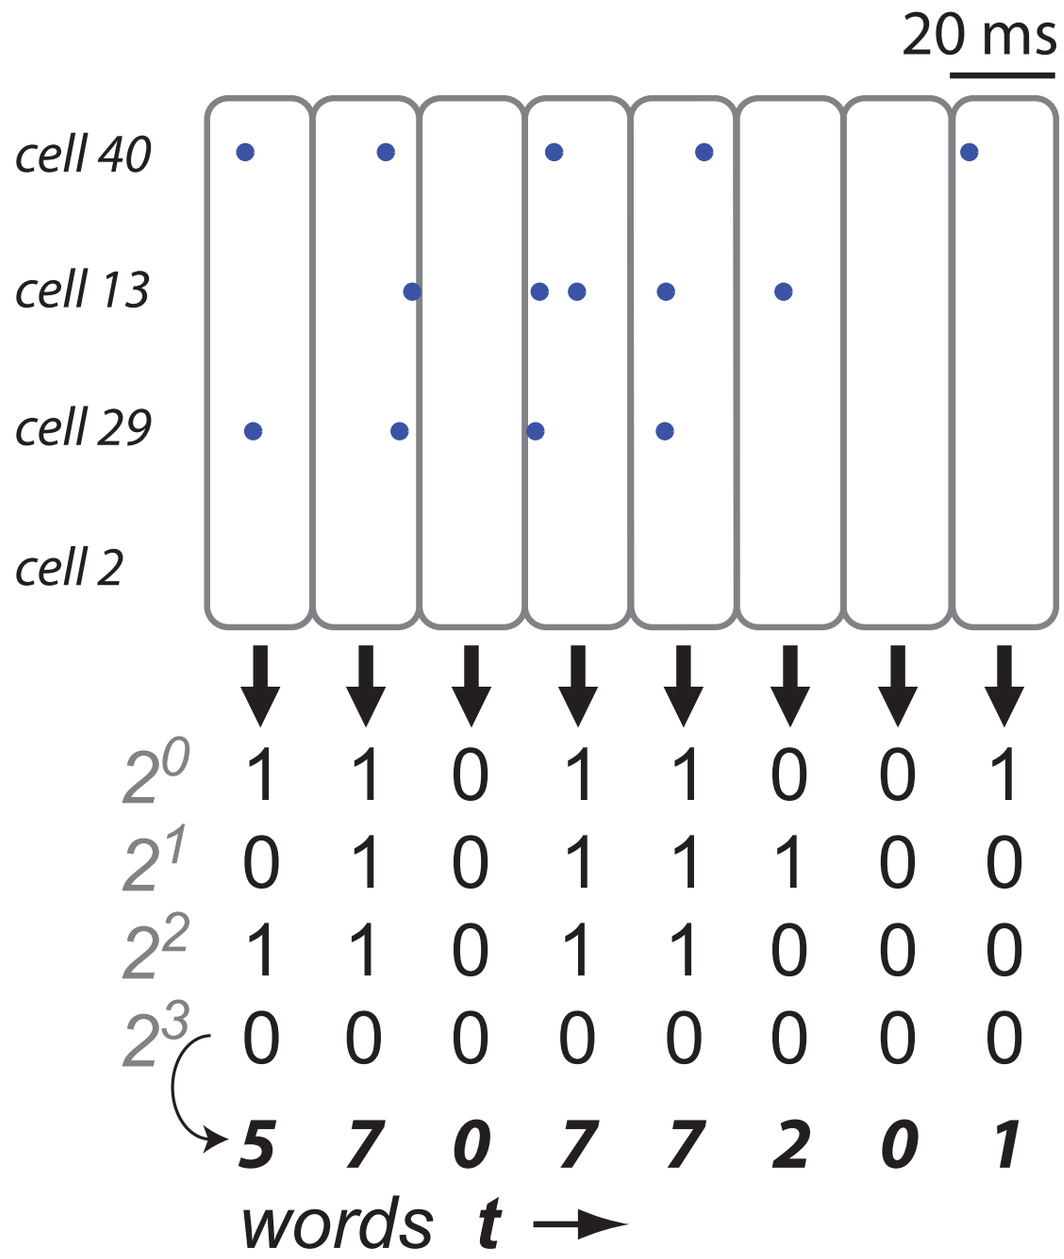

Schematic of our calculations:

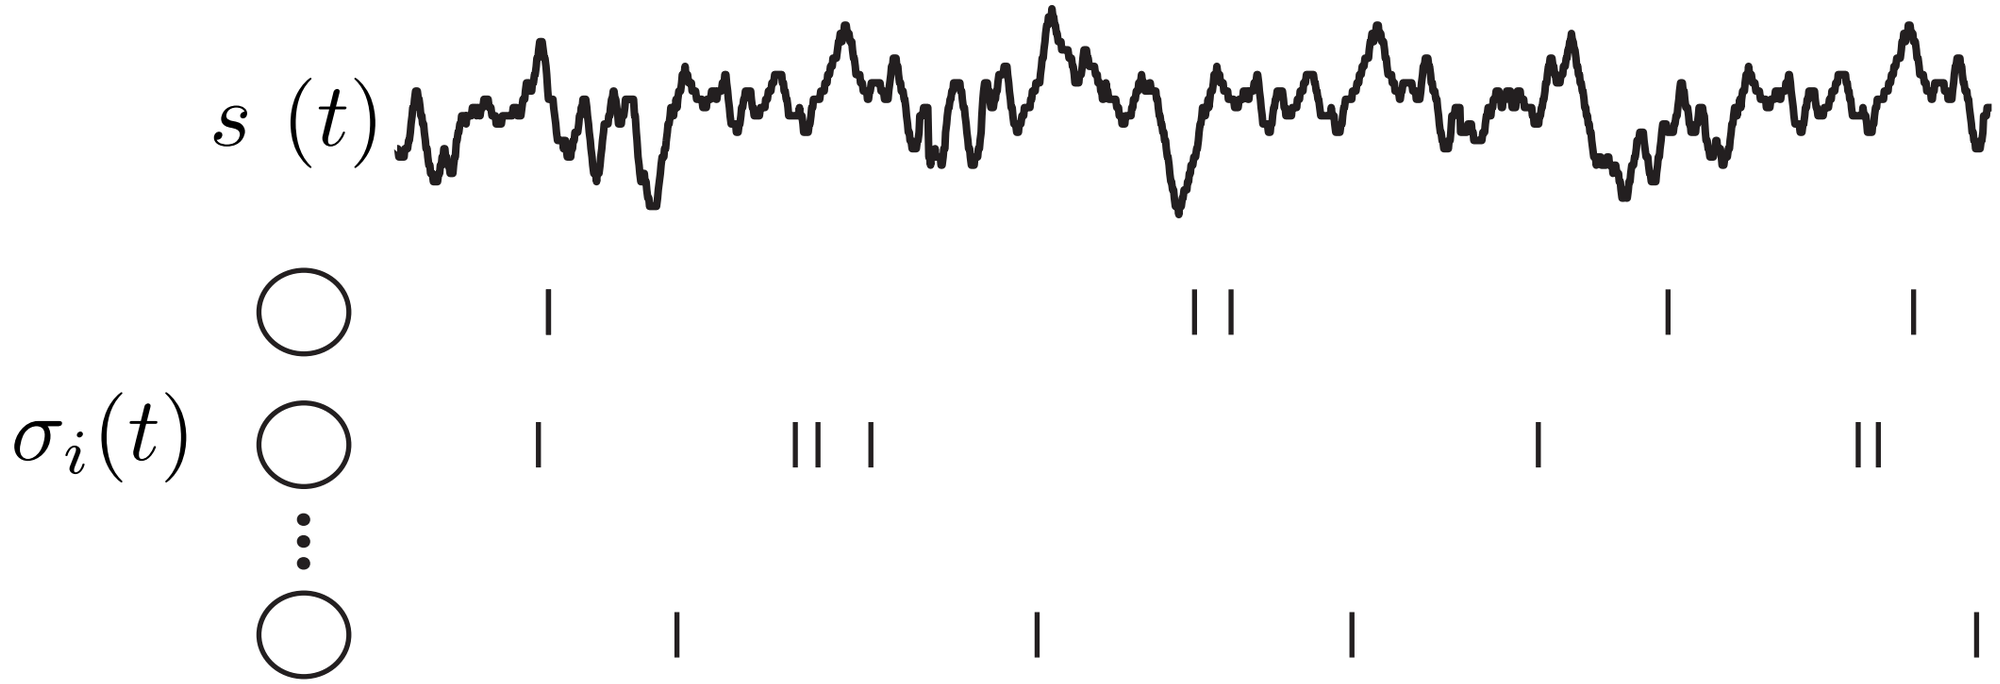

Schematic of our calculations:

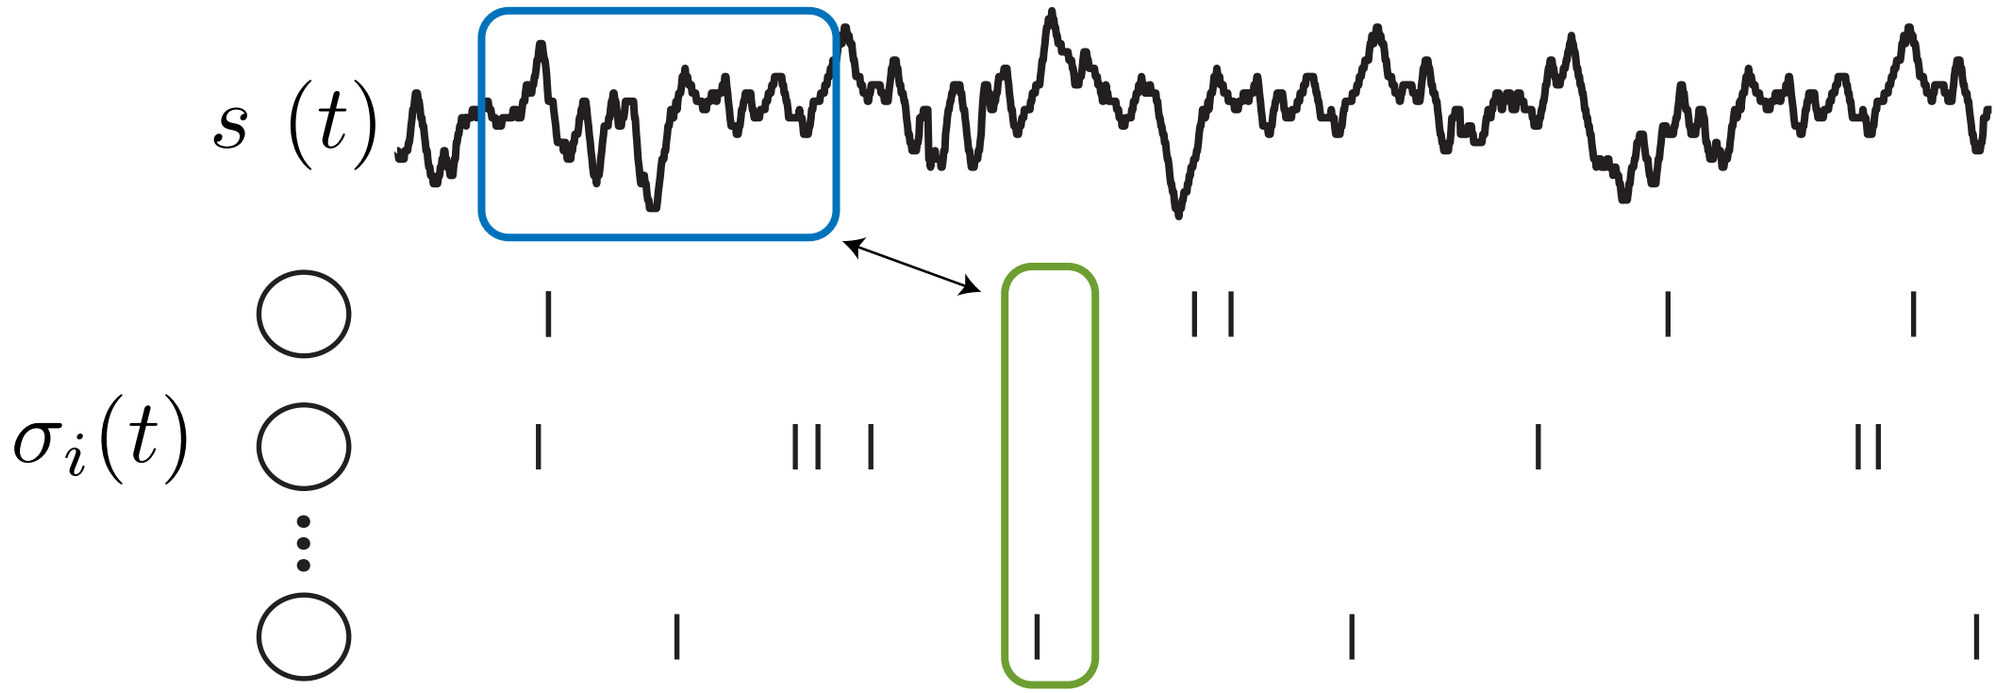

Schematic of our calculations:

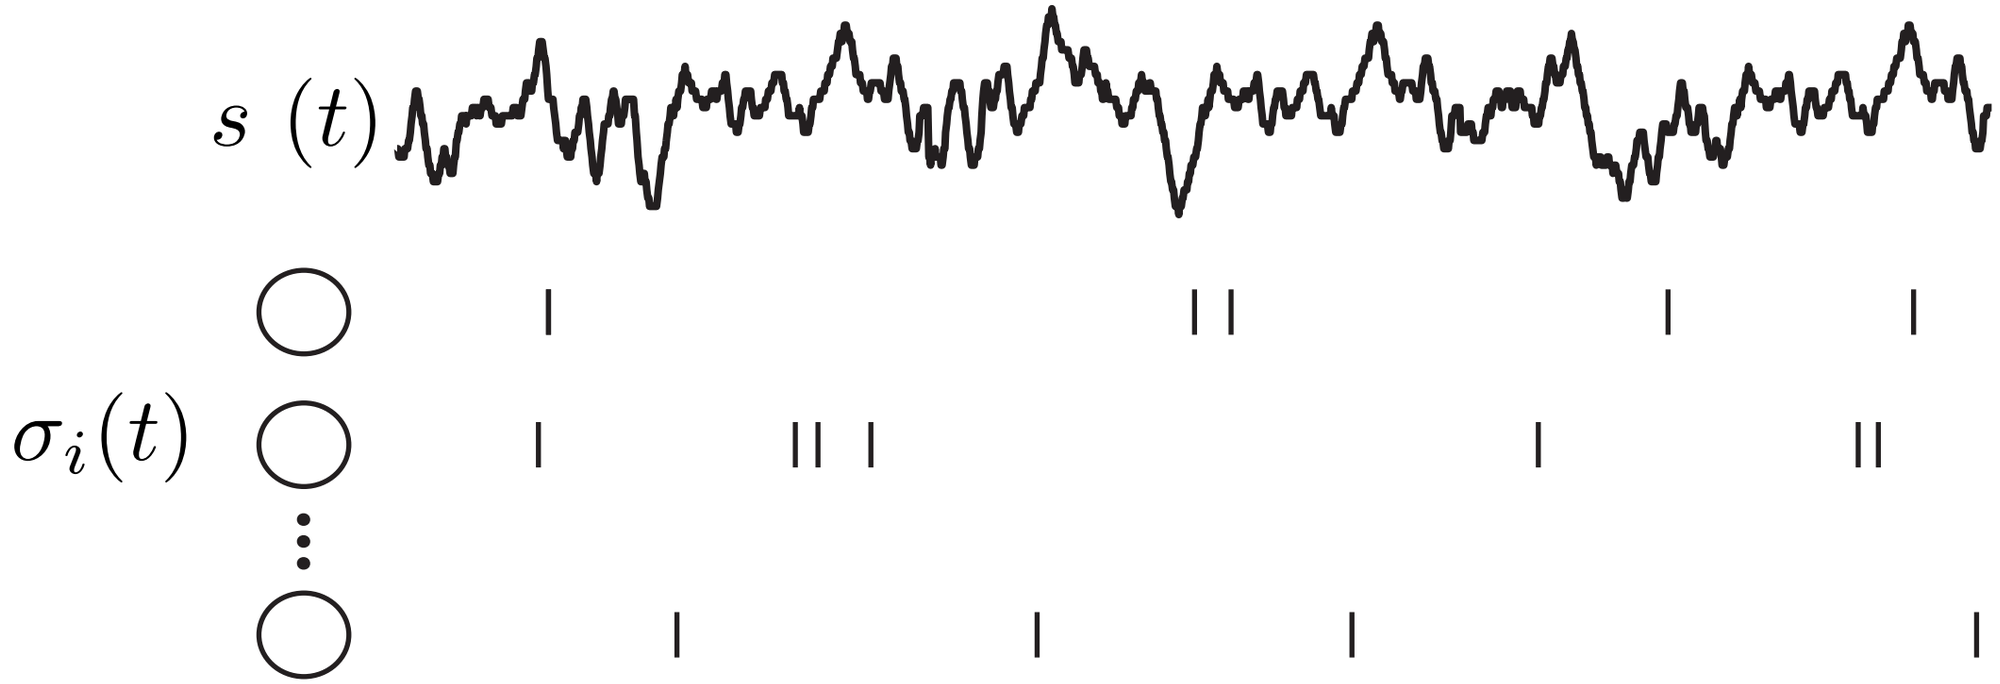

Schematic of our calculations:

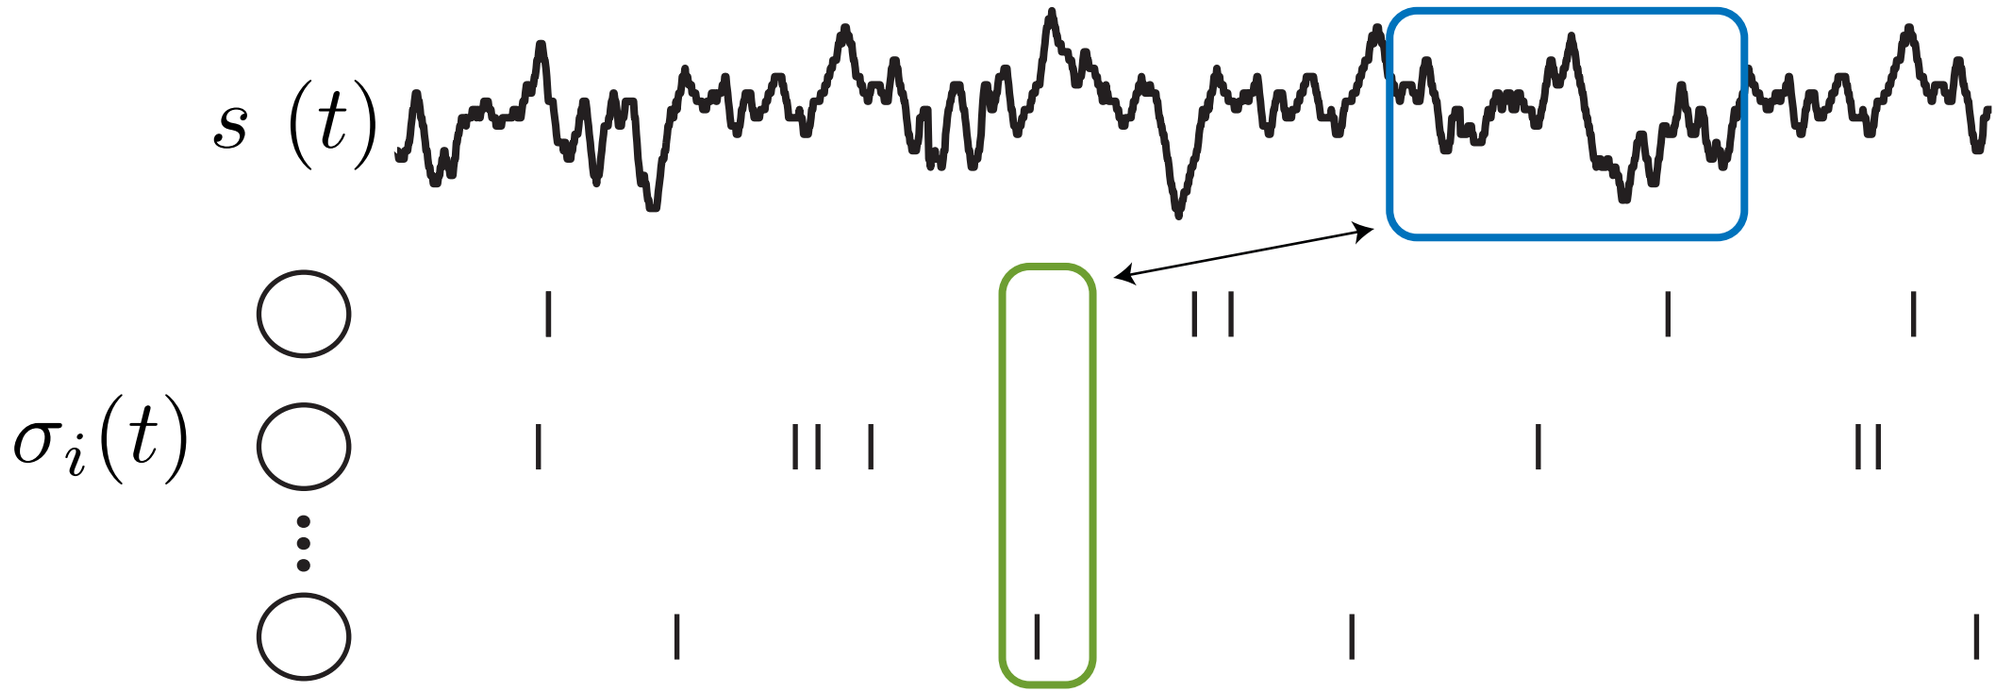

Schematic of our calculations:

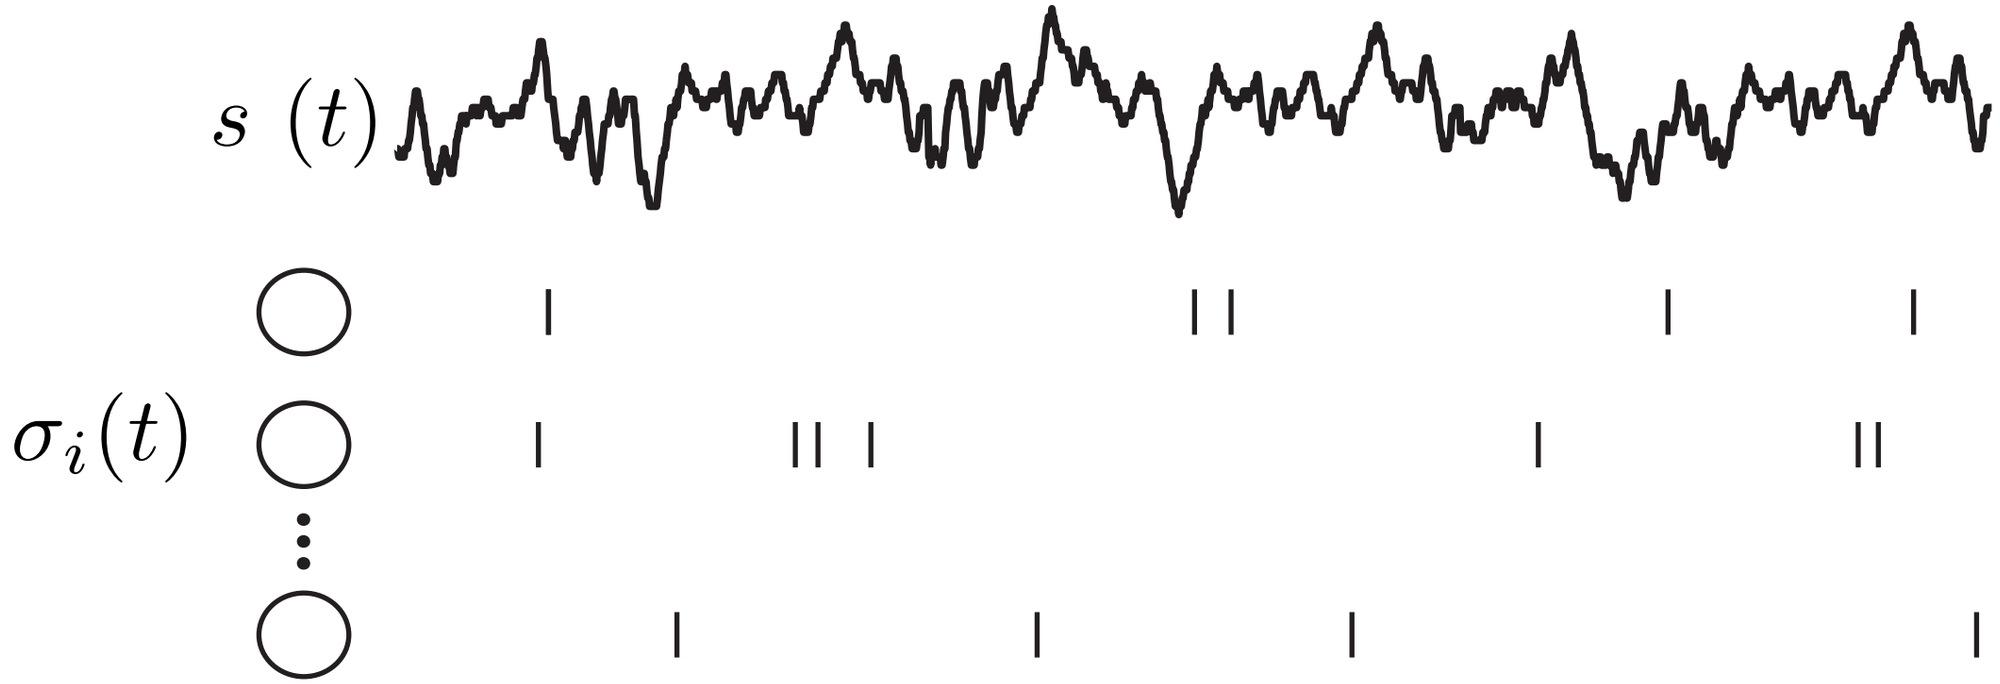

Schematic of our calculations:

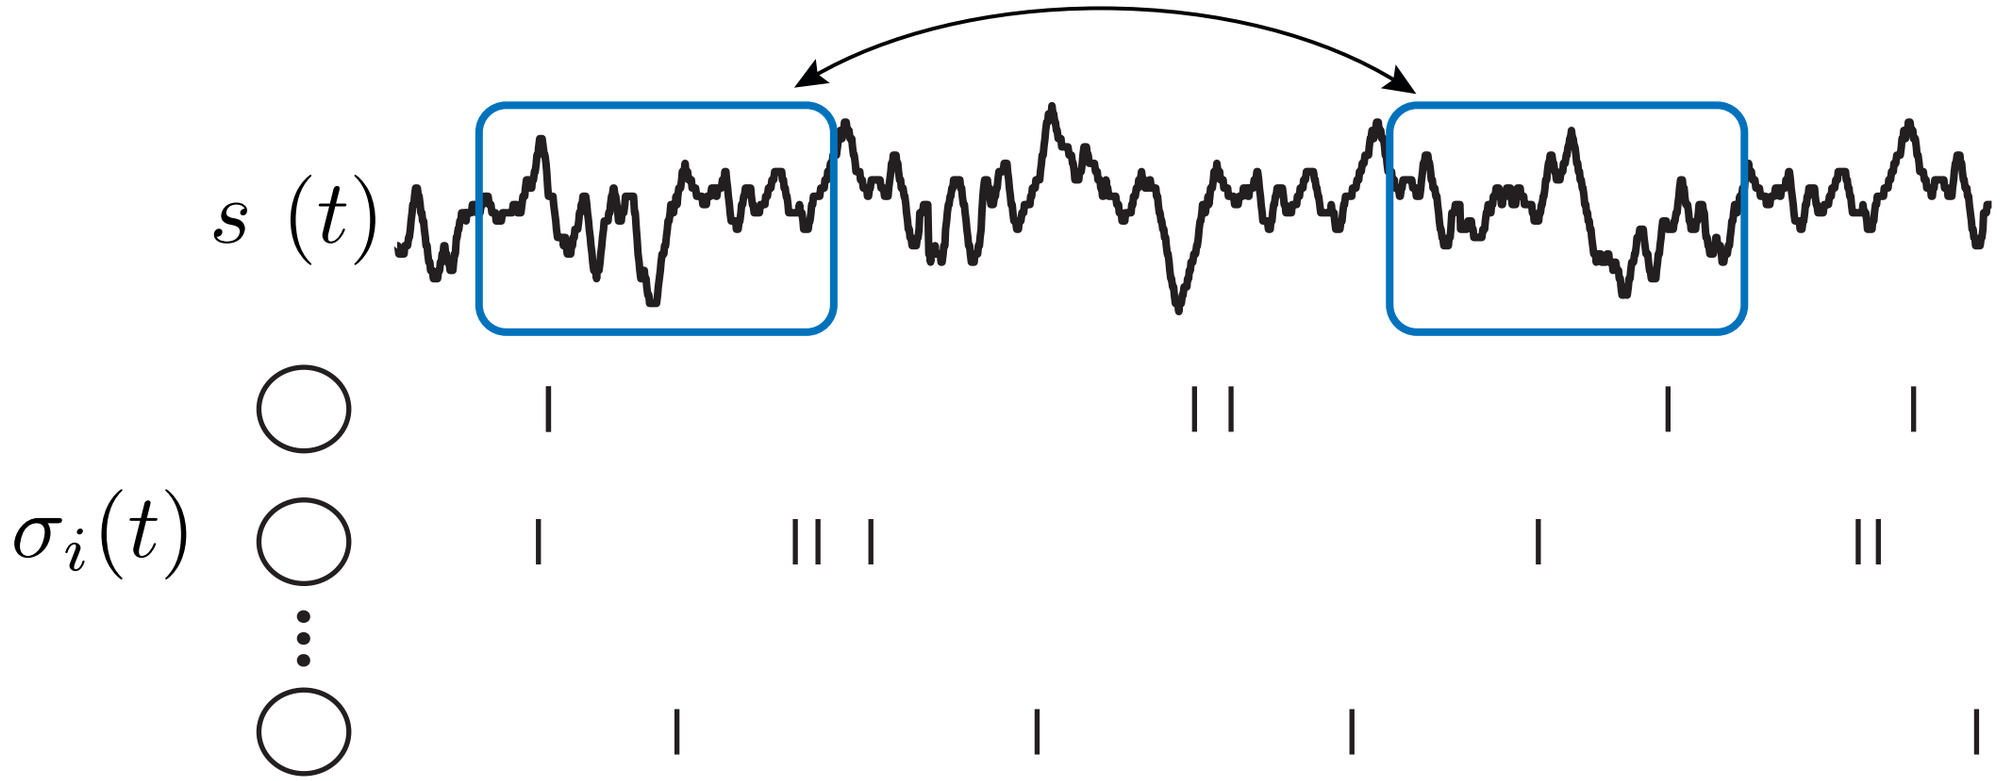

Schematic of our calculations:

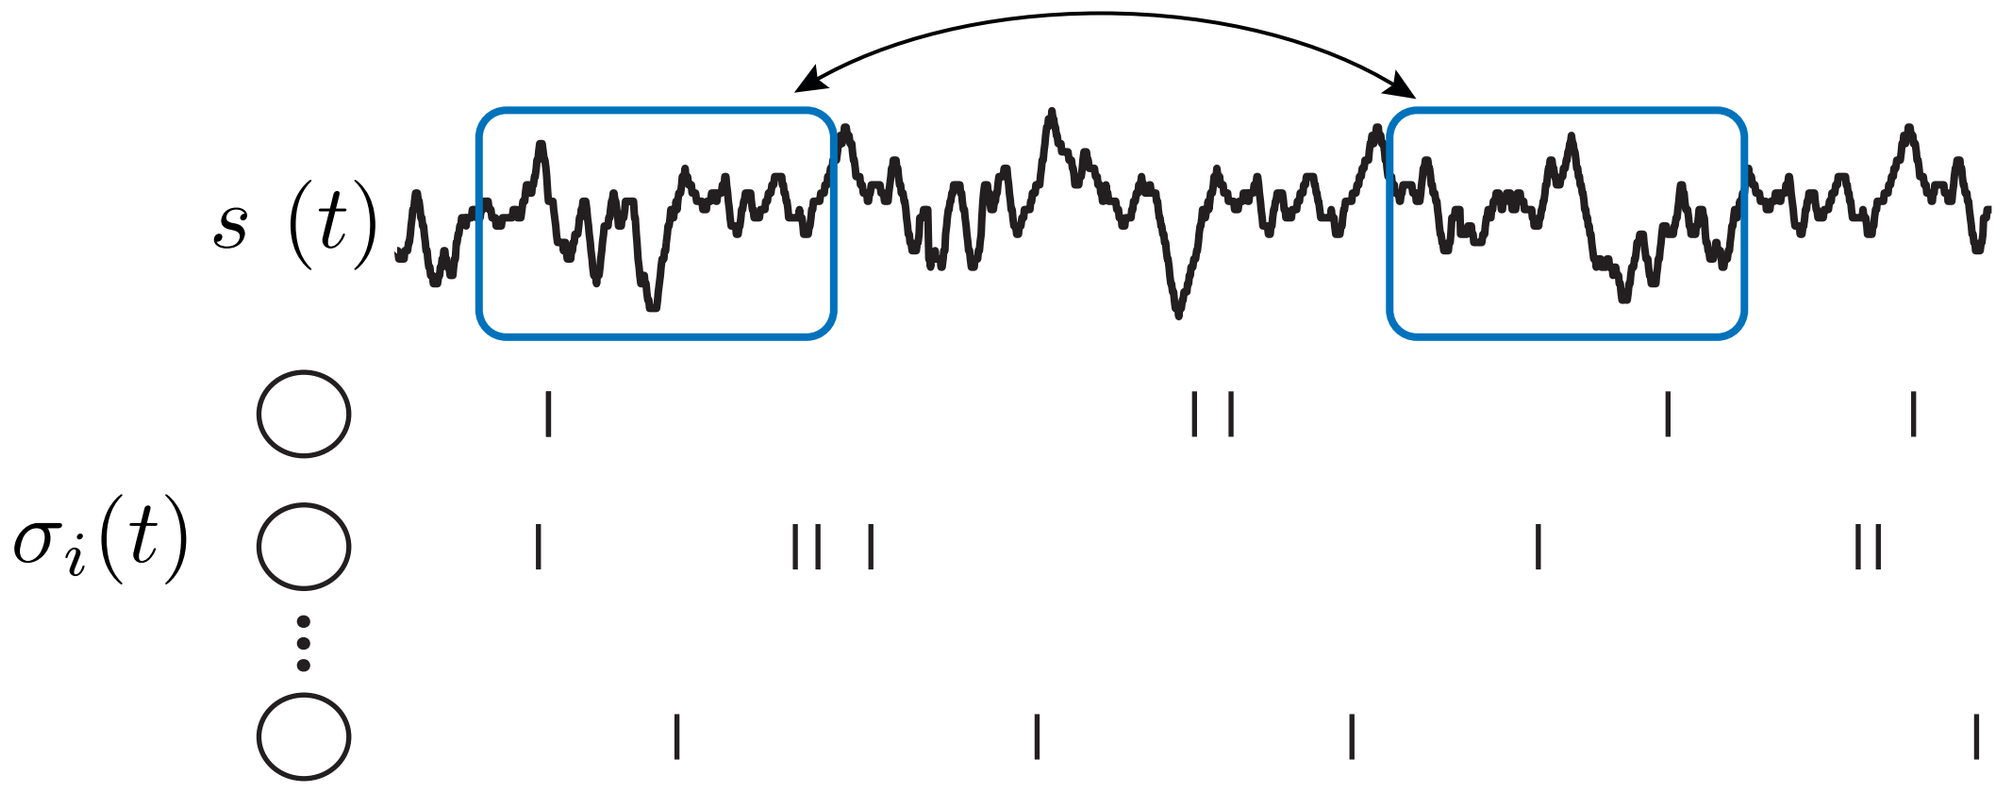

$$I(\text{past};\text{future}) = S(\text{future}) - S(\text{future}|\text{past})$$

Retina populations carry info about the future:

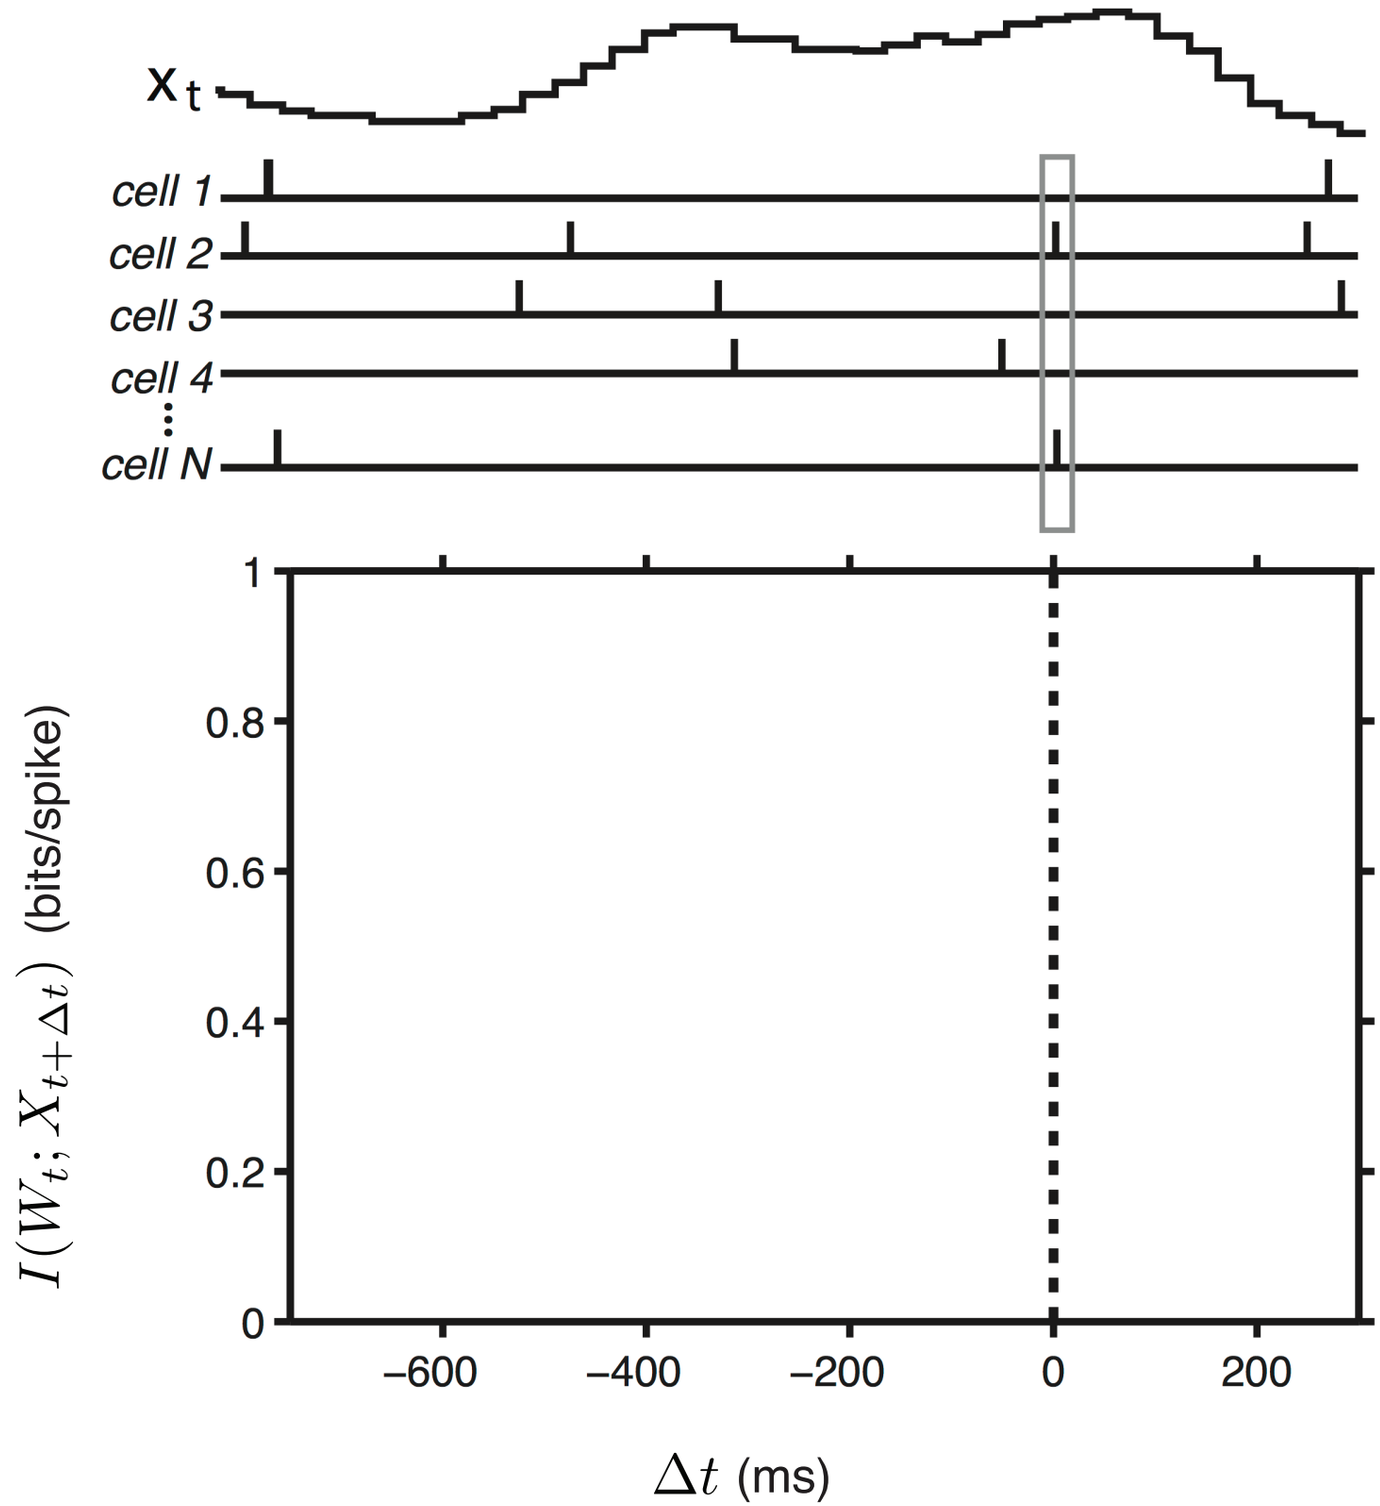

Retina populations carry info about the future:

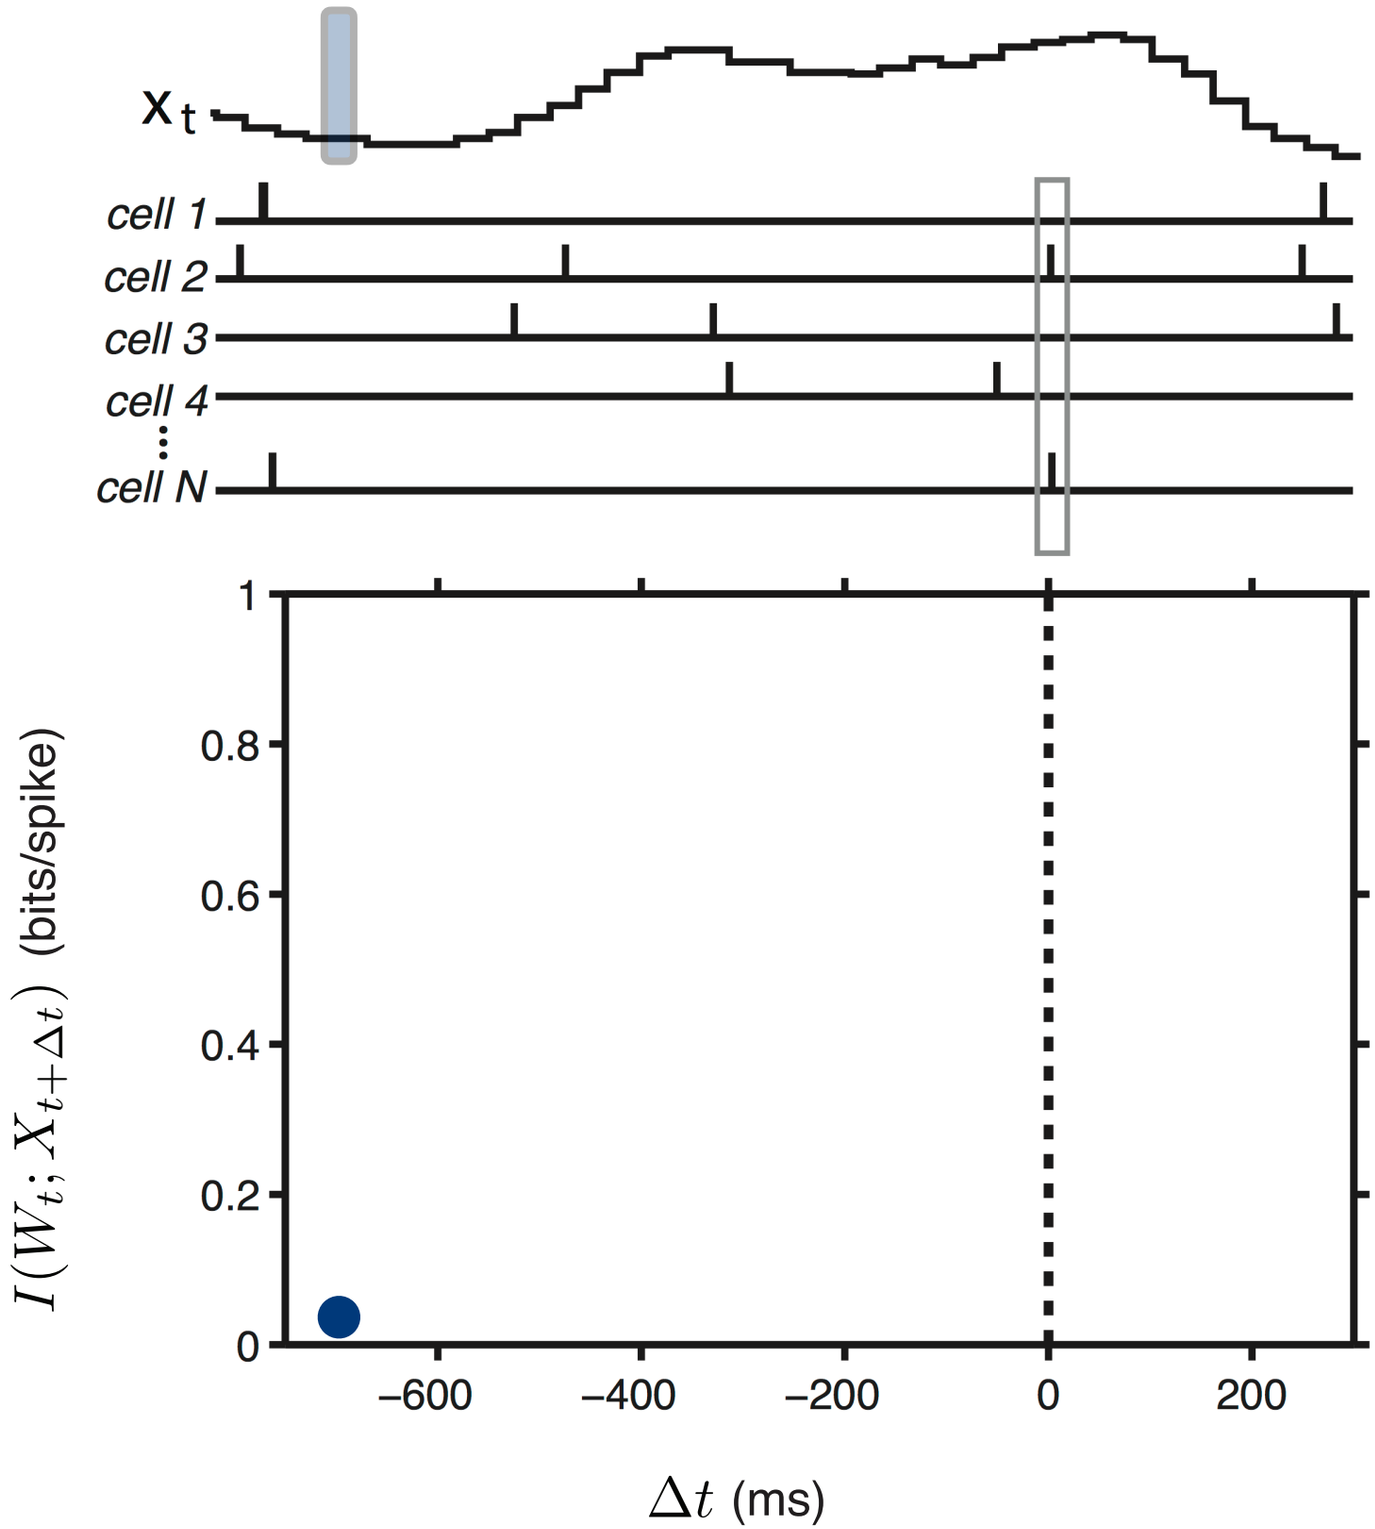

Retina populations carry info about the future:

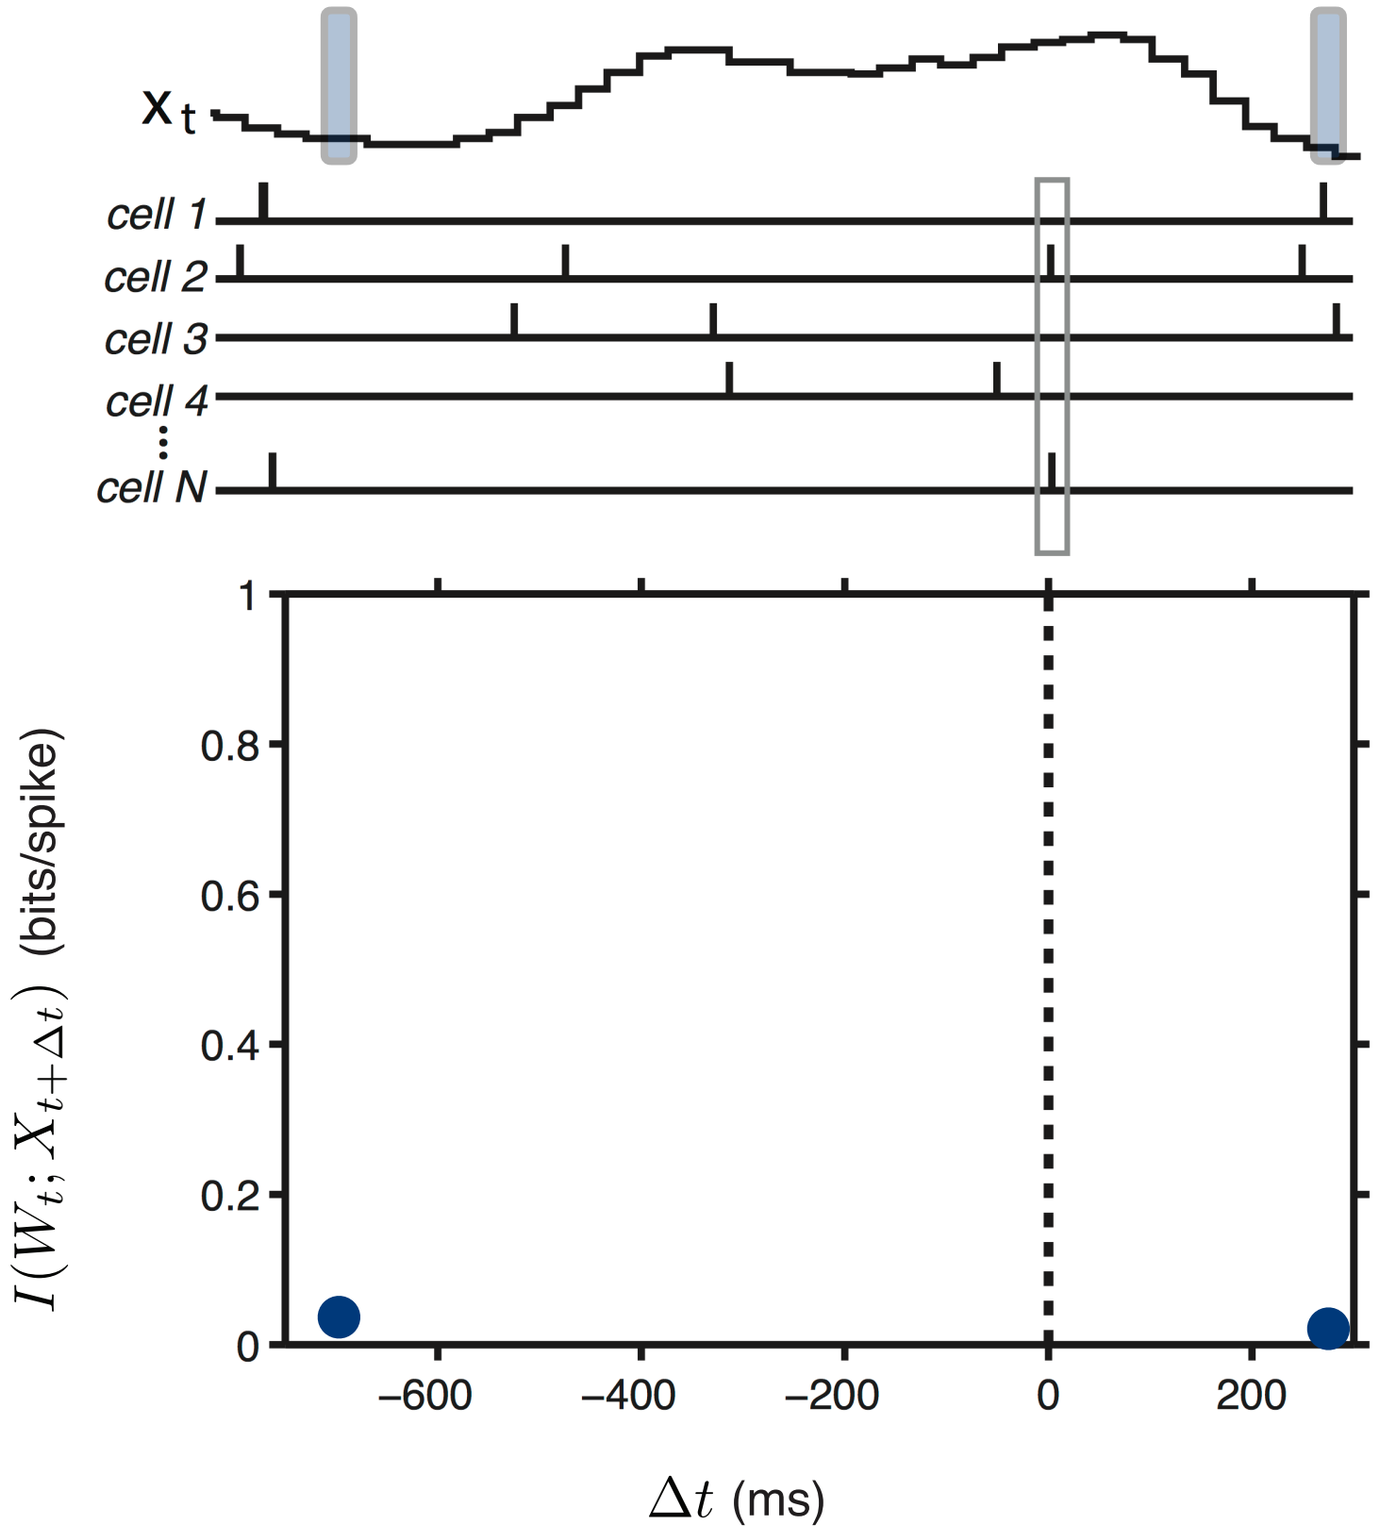

Retina populations carry info about the future:

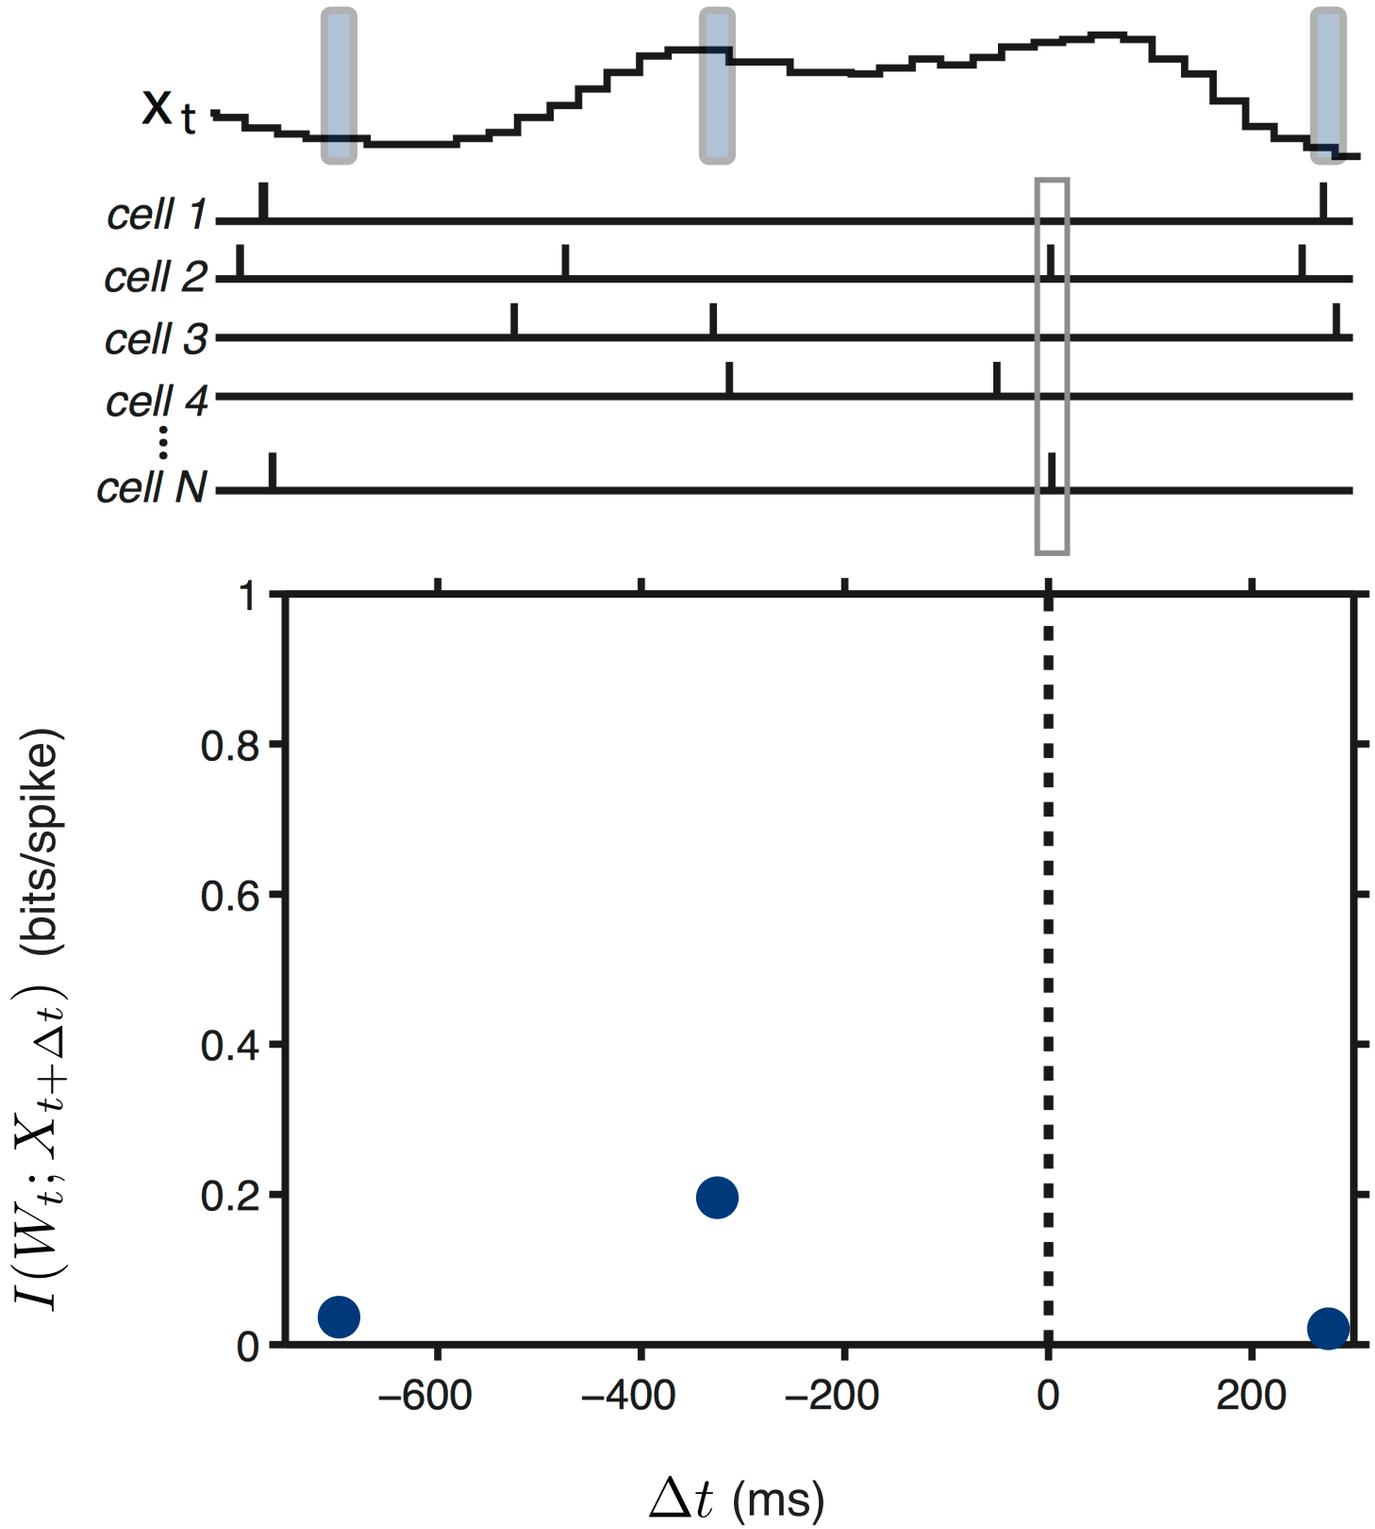

Retina populations carry info about the future:

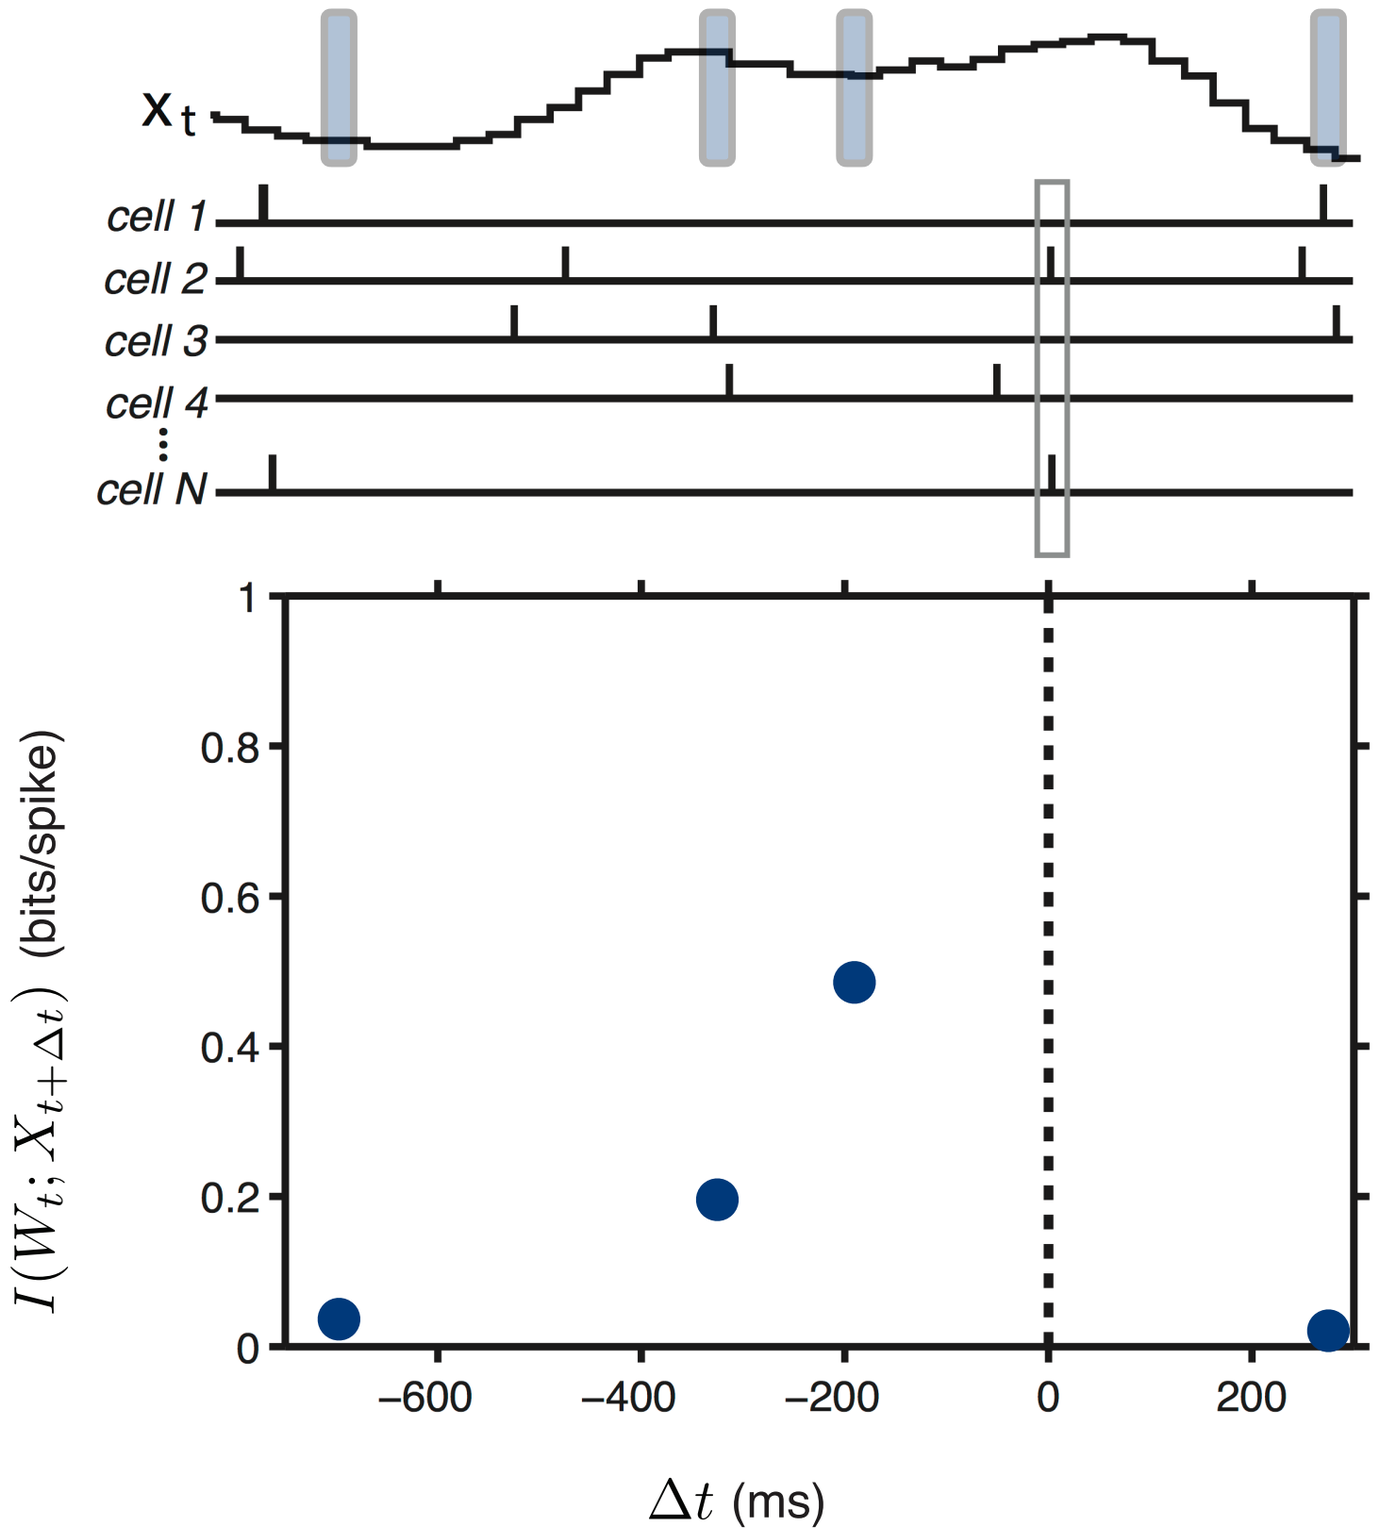

Retina populations carry info about the future:

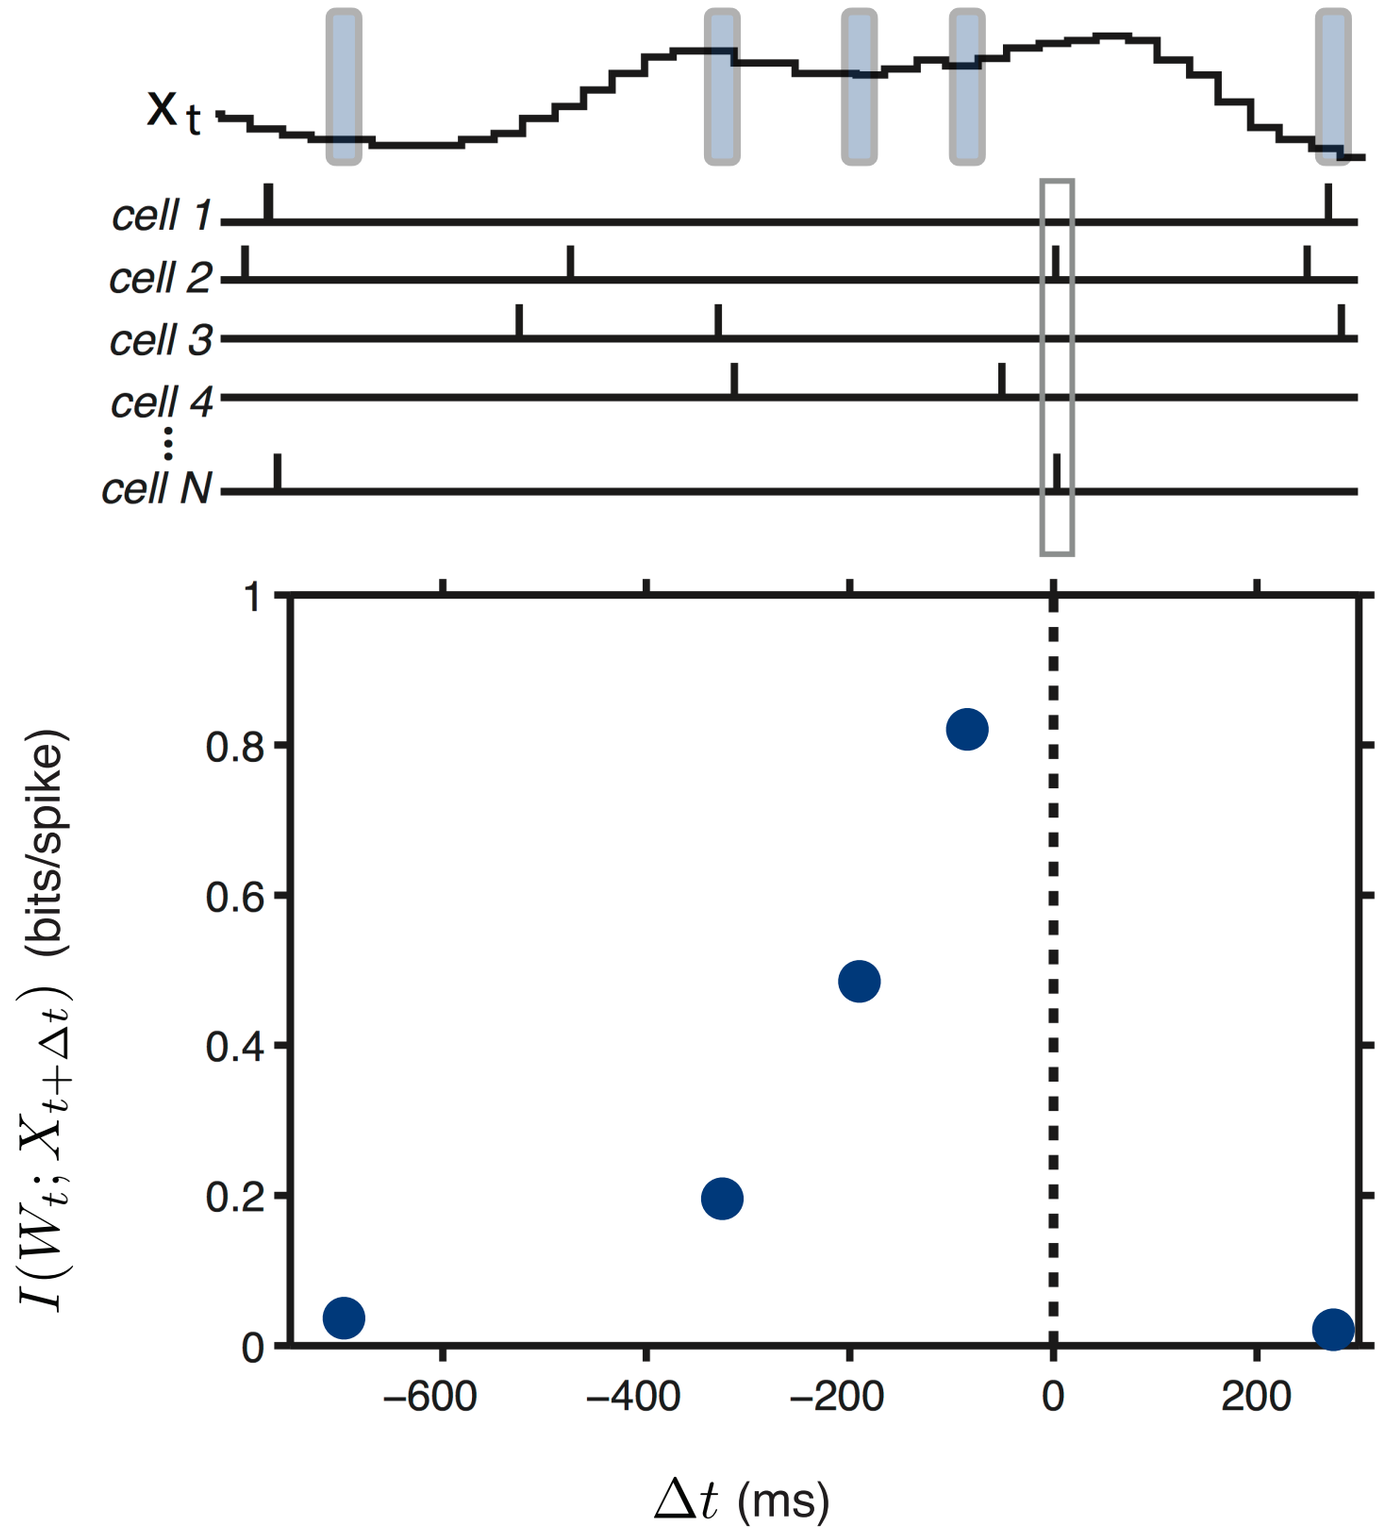

Retina populations carry info about the future:

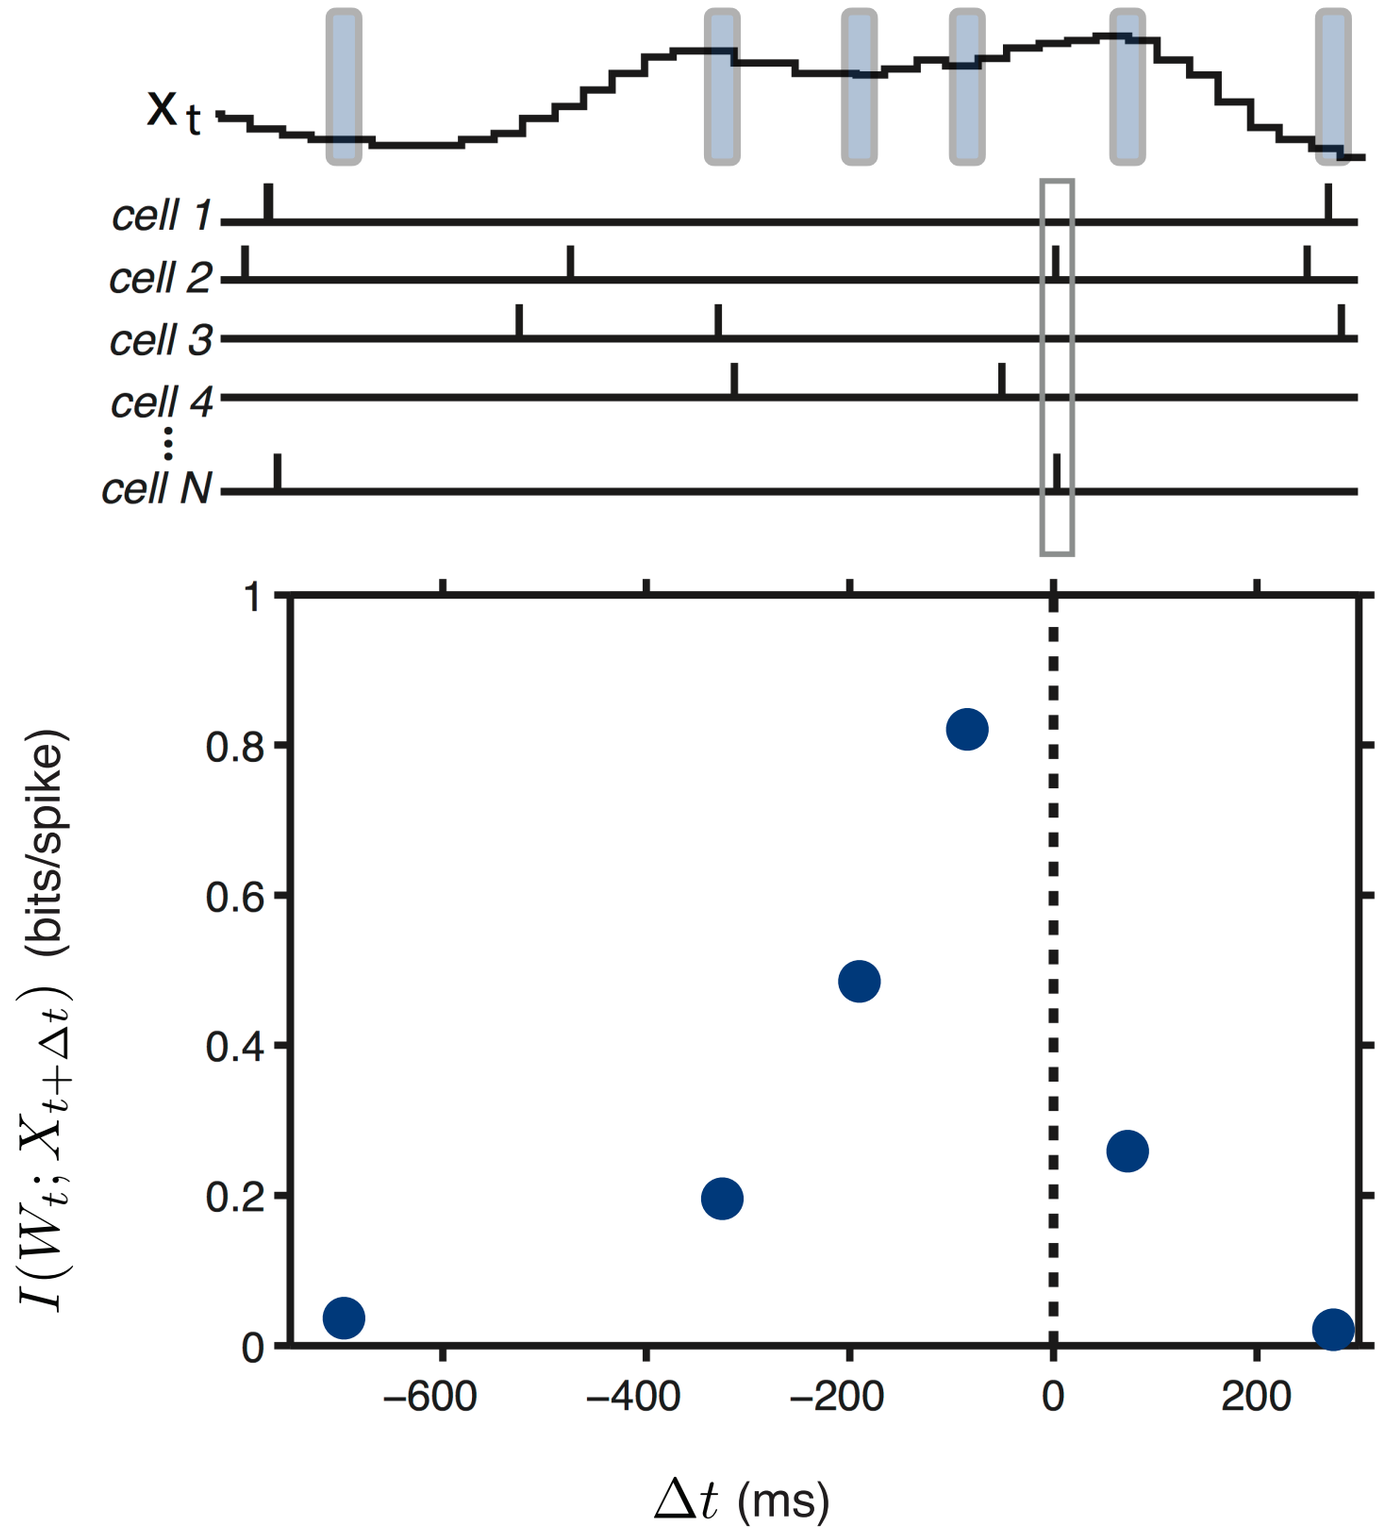

Retina populations carry info about the future:

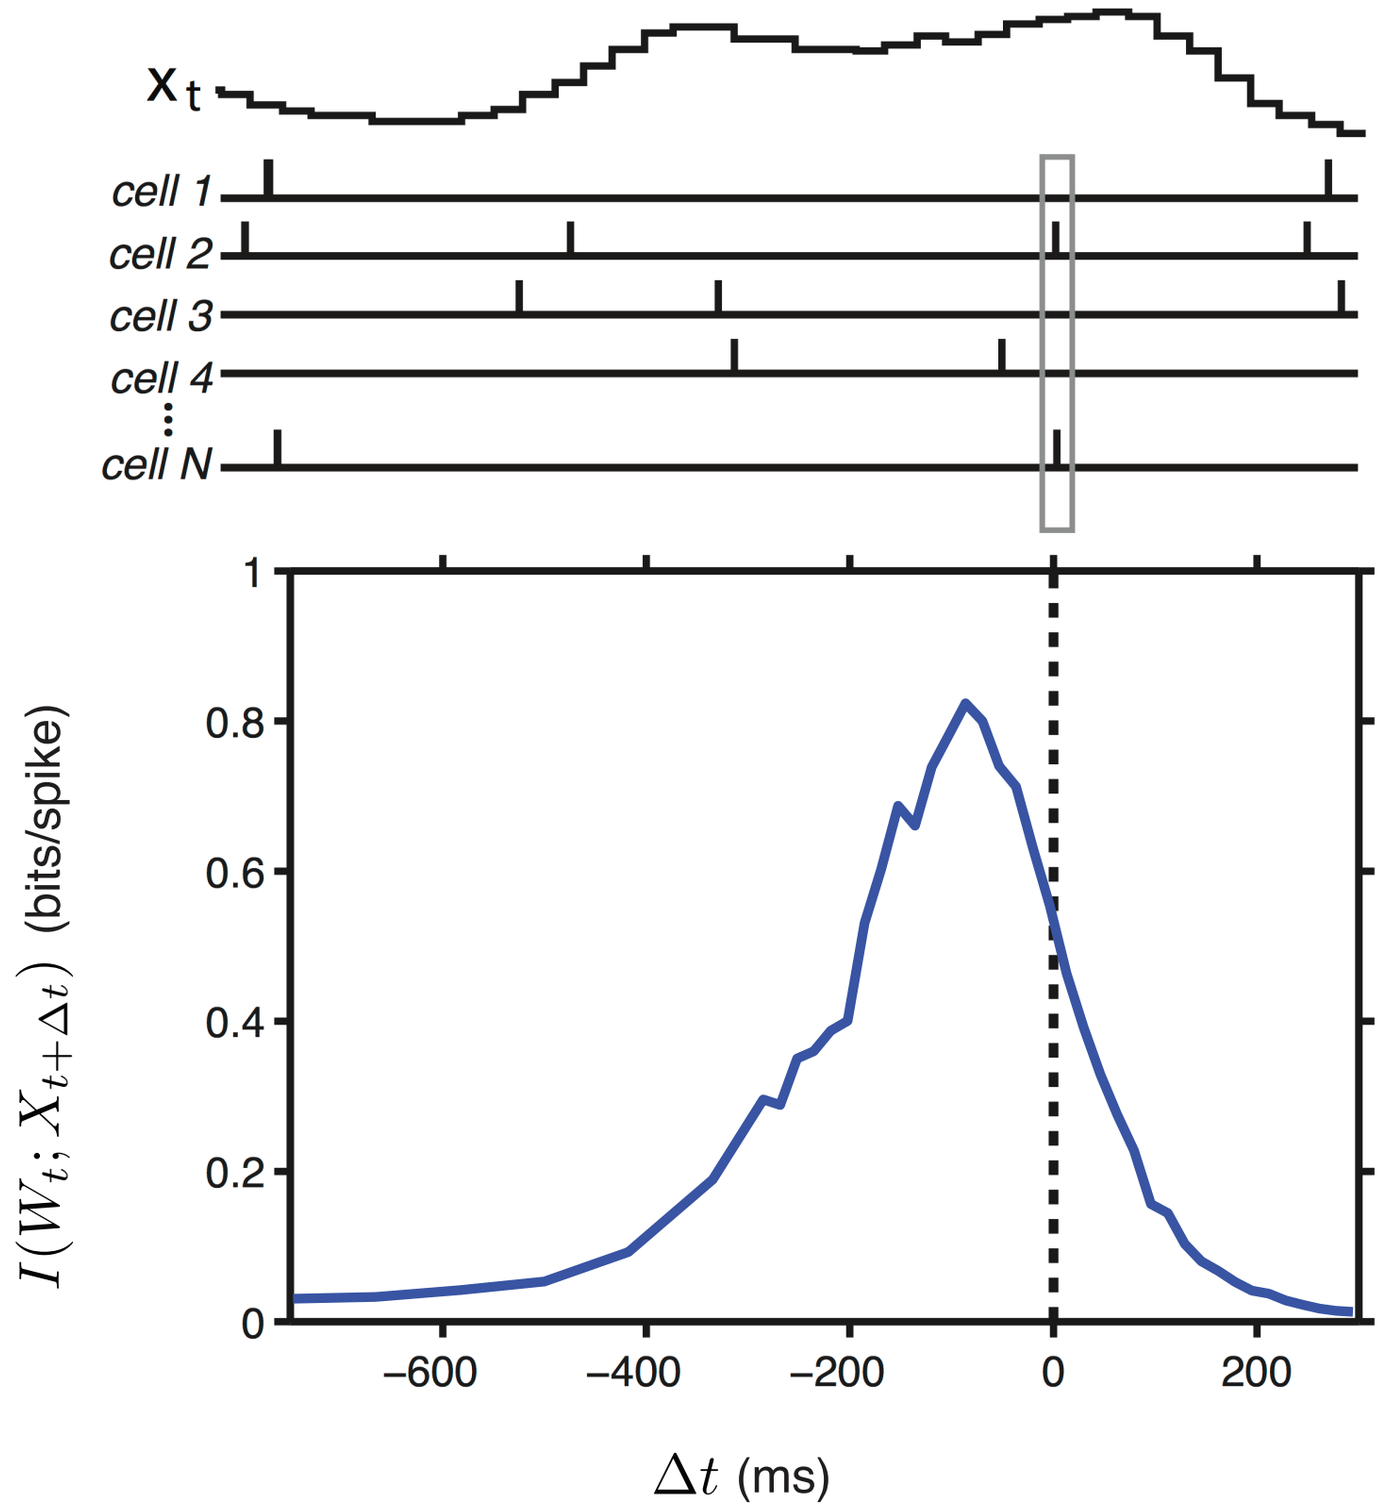

Retina populations carry info about the future:

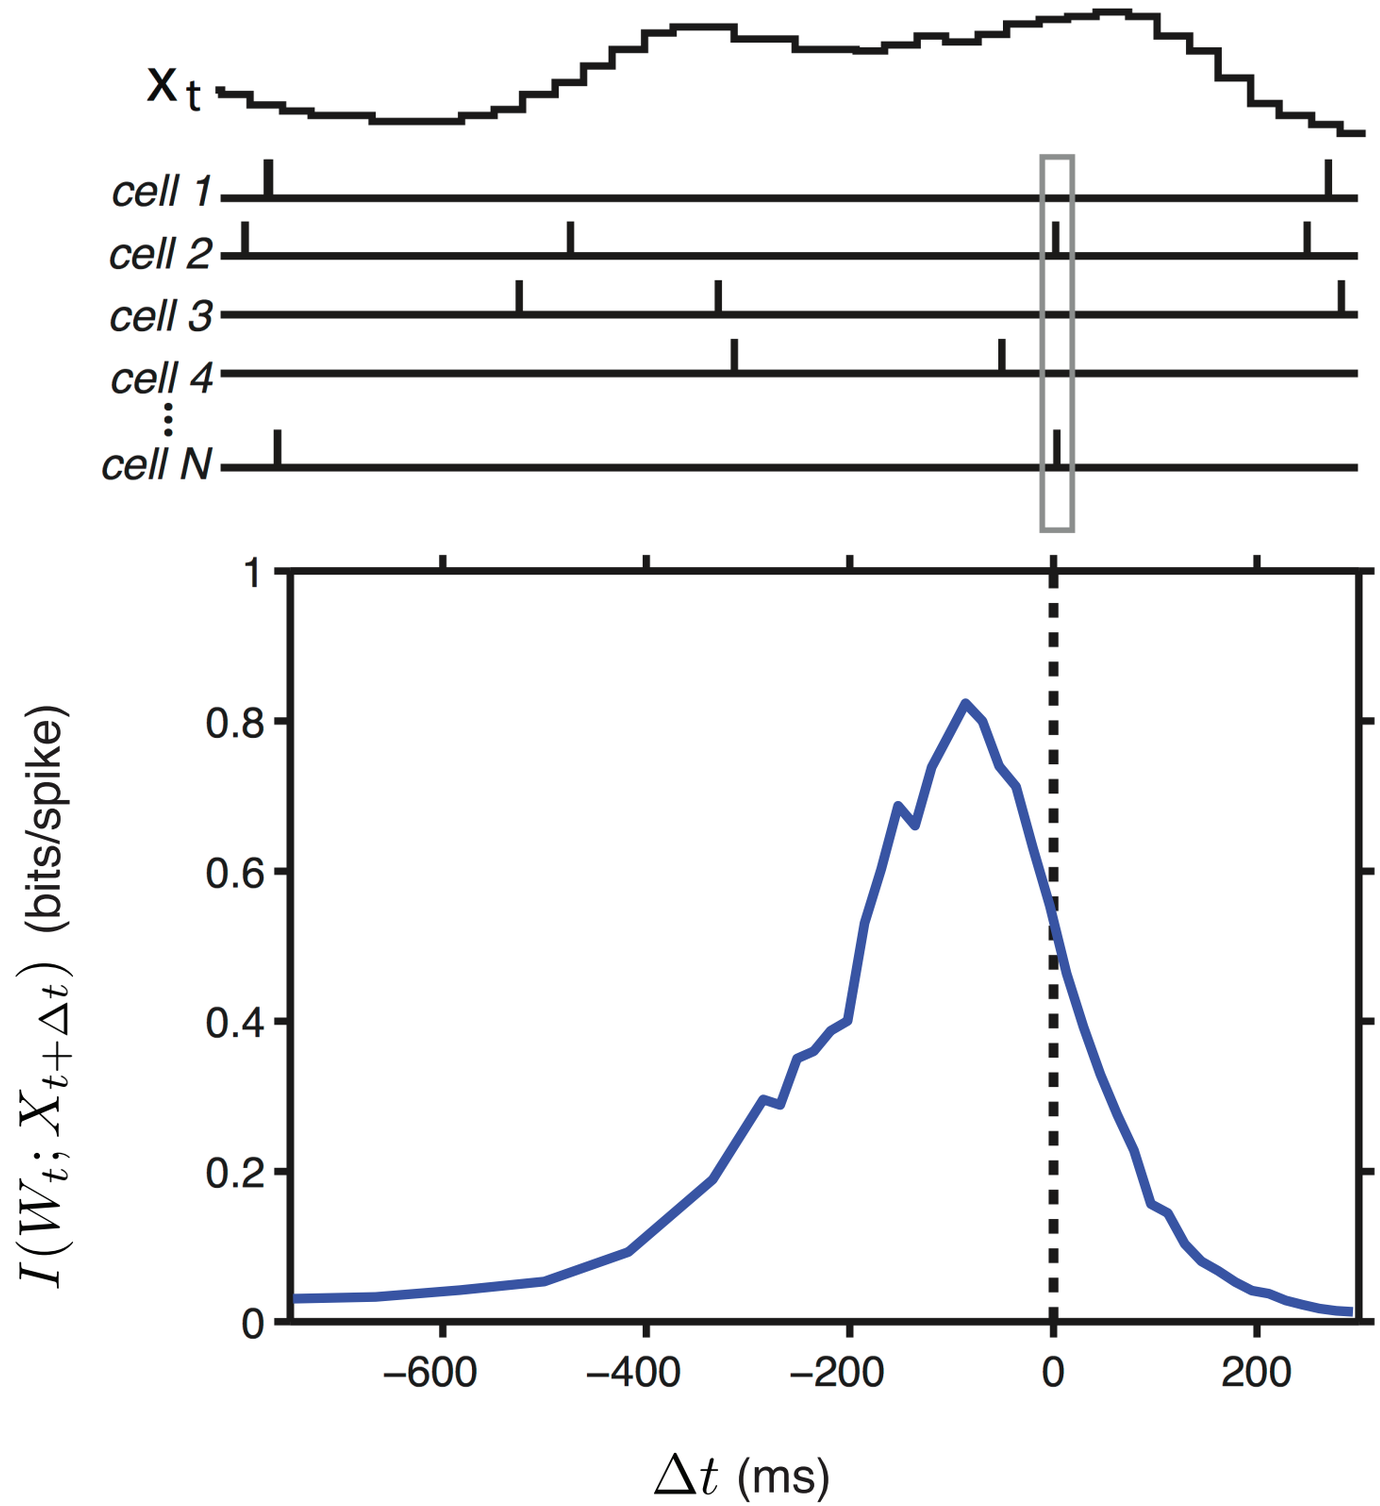

Palmer, Marre, Berry, Bialek PNAS (2015)

Retina populations carry info about the future:

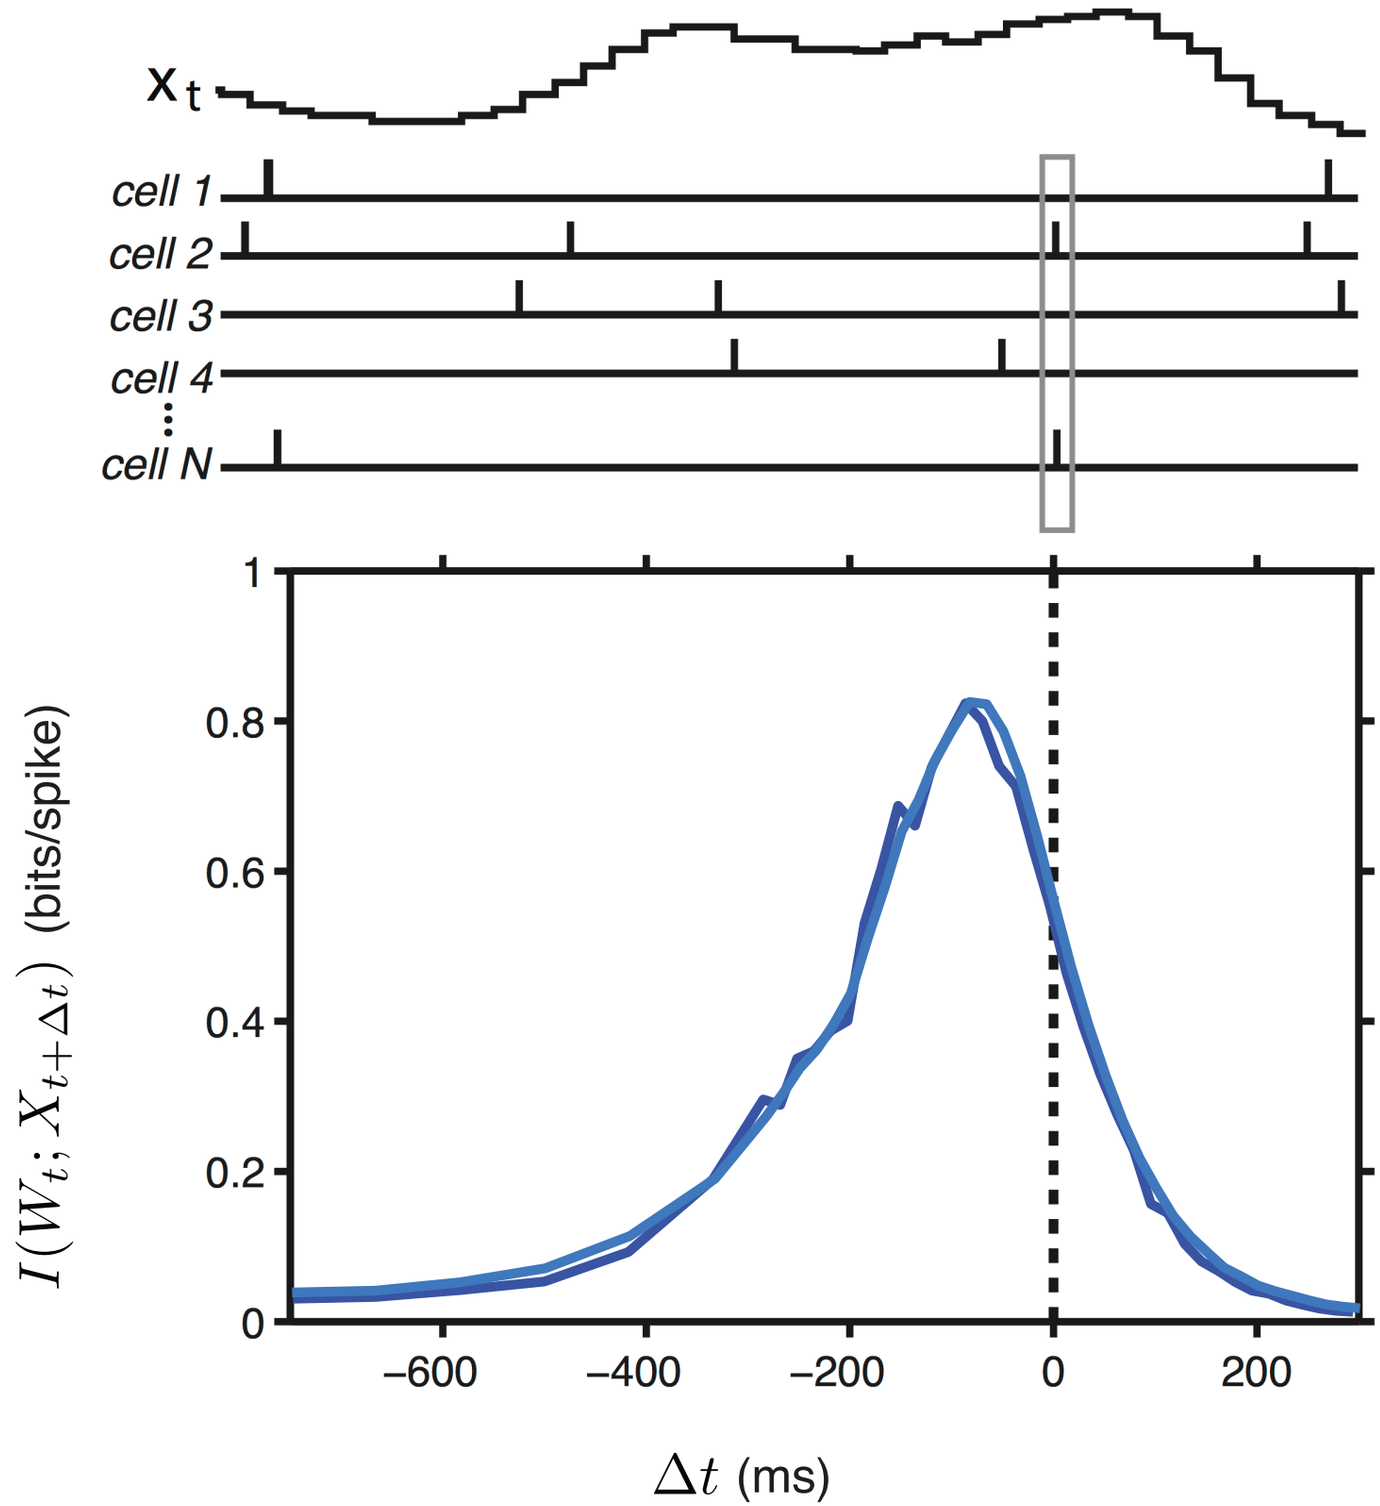

Palmer, Marre, Berry, Bialek PNAS (2015)

Retina populations carry info about the future:

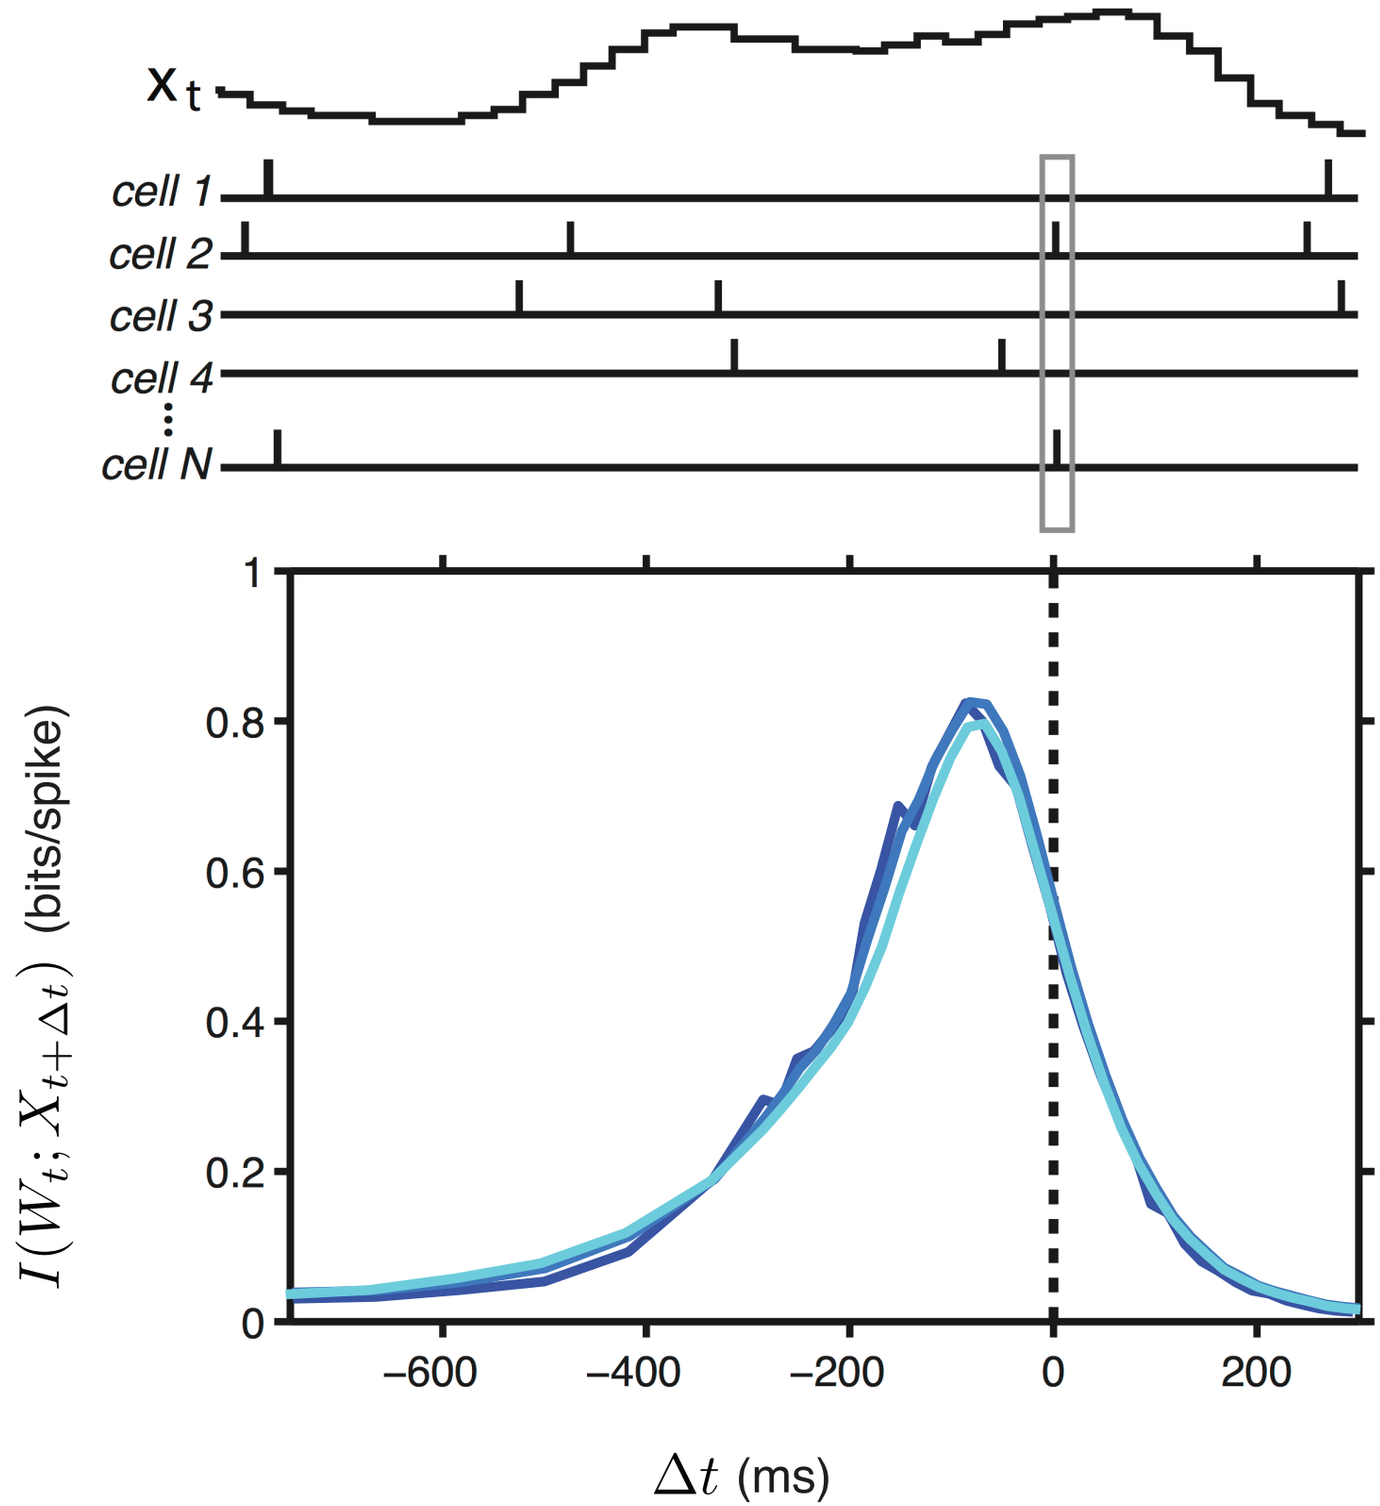

Palmer, Marre, Berry, Bialek PNAS (2015)

Retina populations carry info about the future:

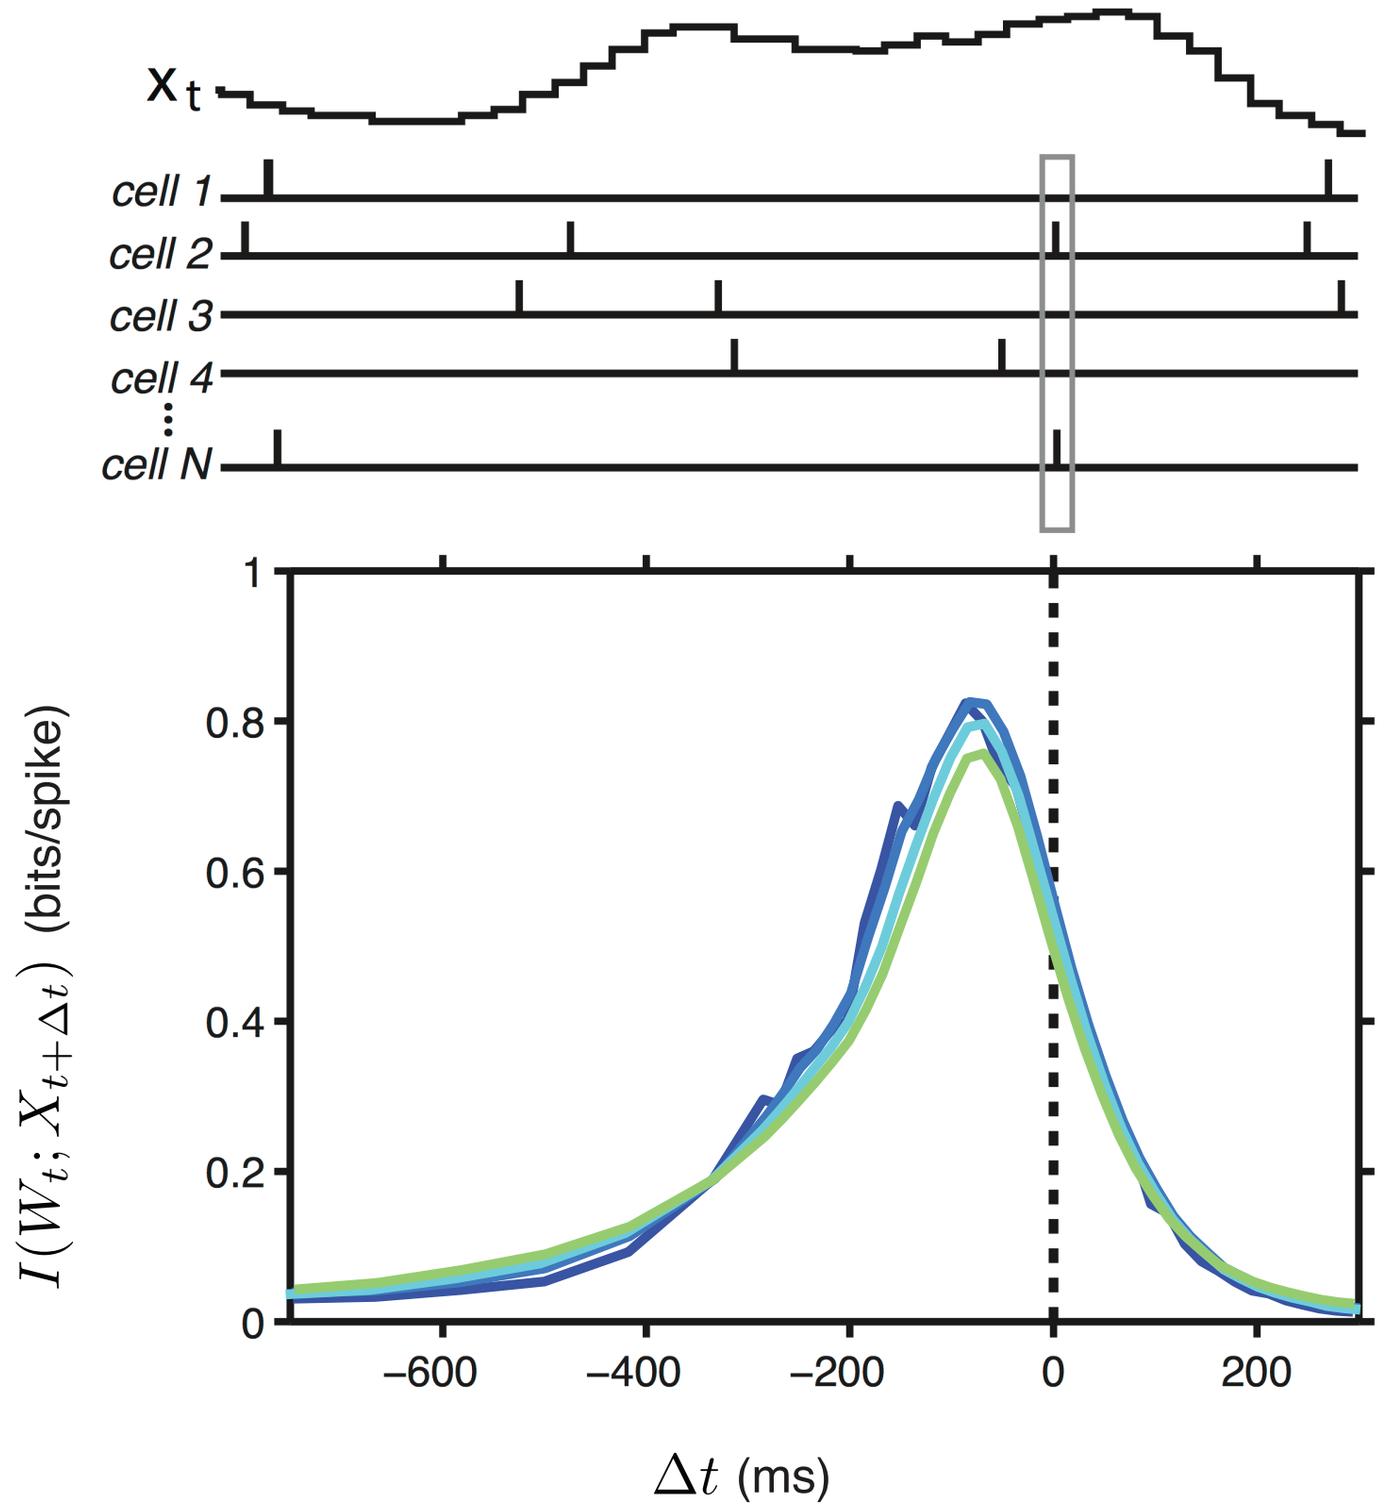

Palmer, Marre, Berry, Bialek PNAS (2015)

Retina populations carry info about the future:

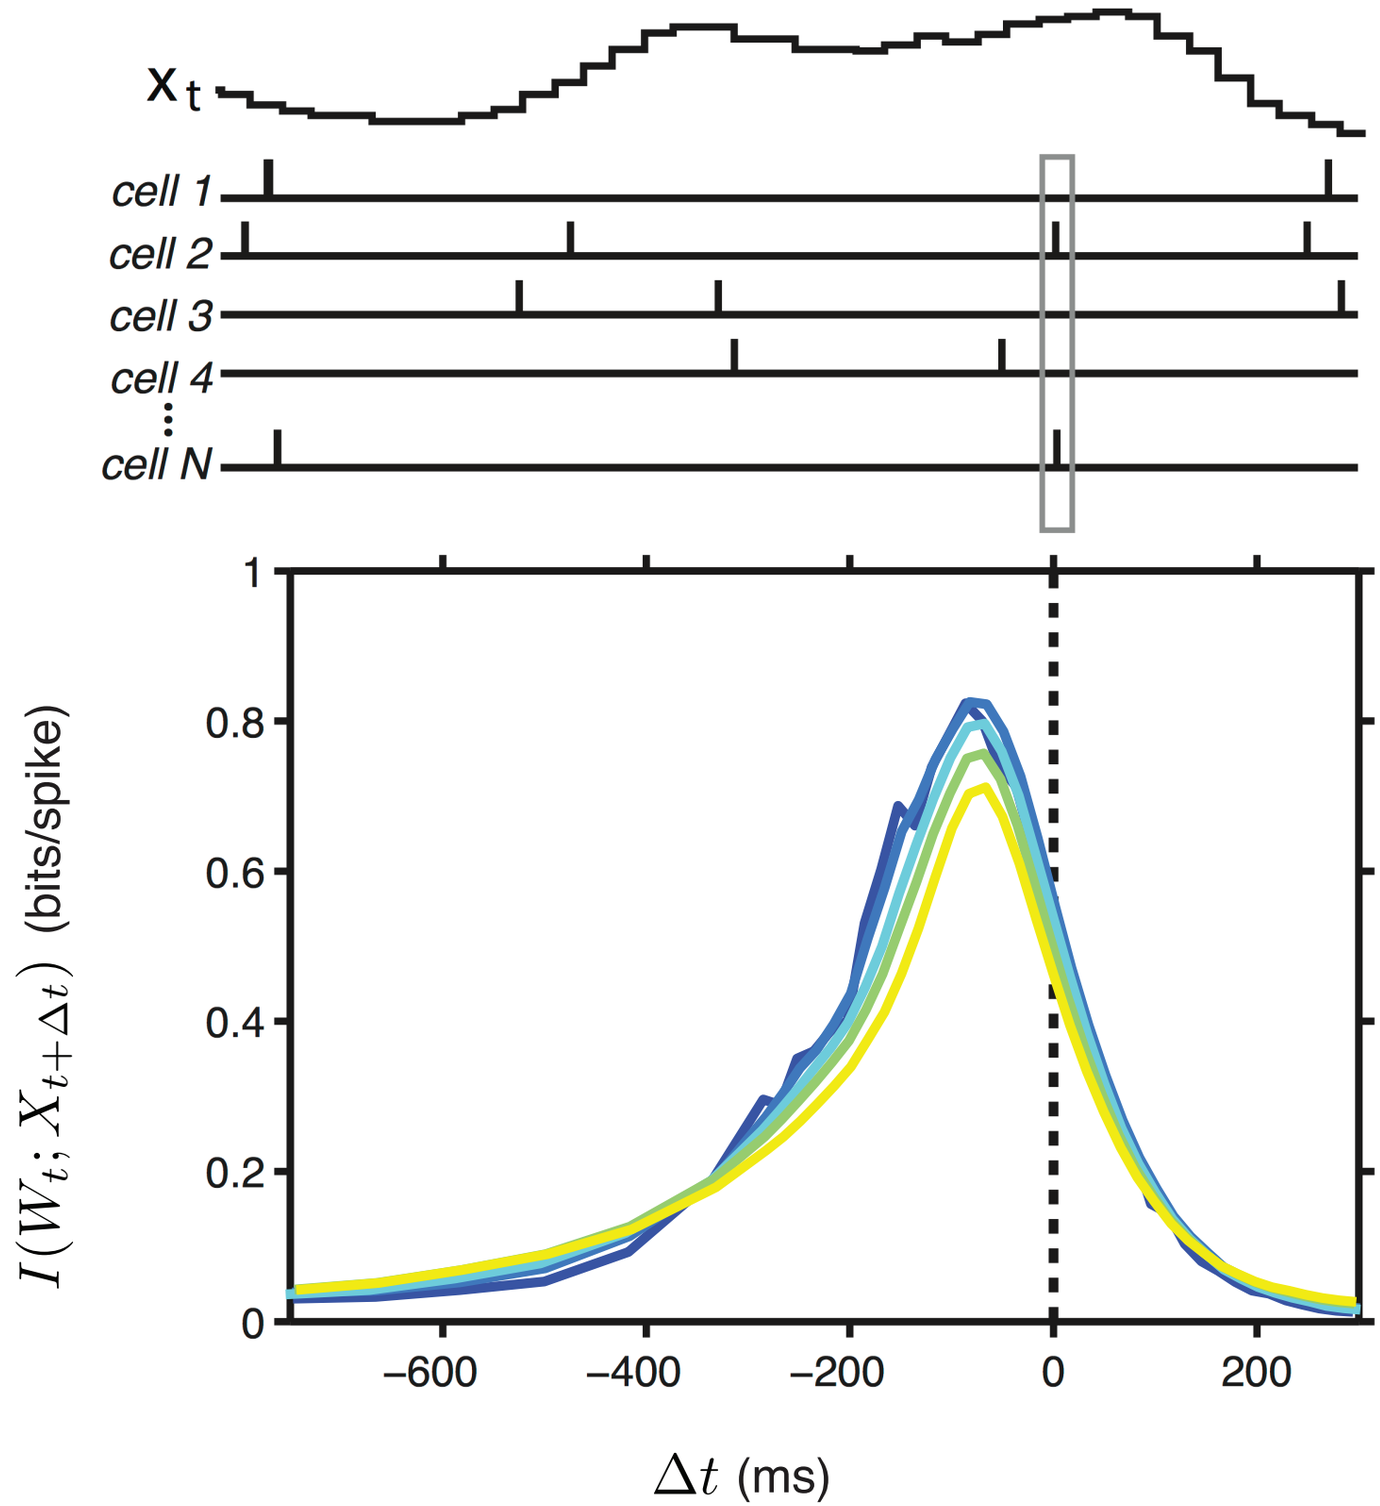

Palmer, Marre, Berry, Bialek PNAS (2015)

Retina populations carry info about the future:

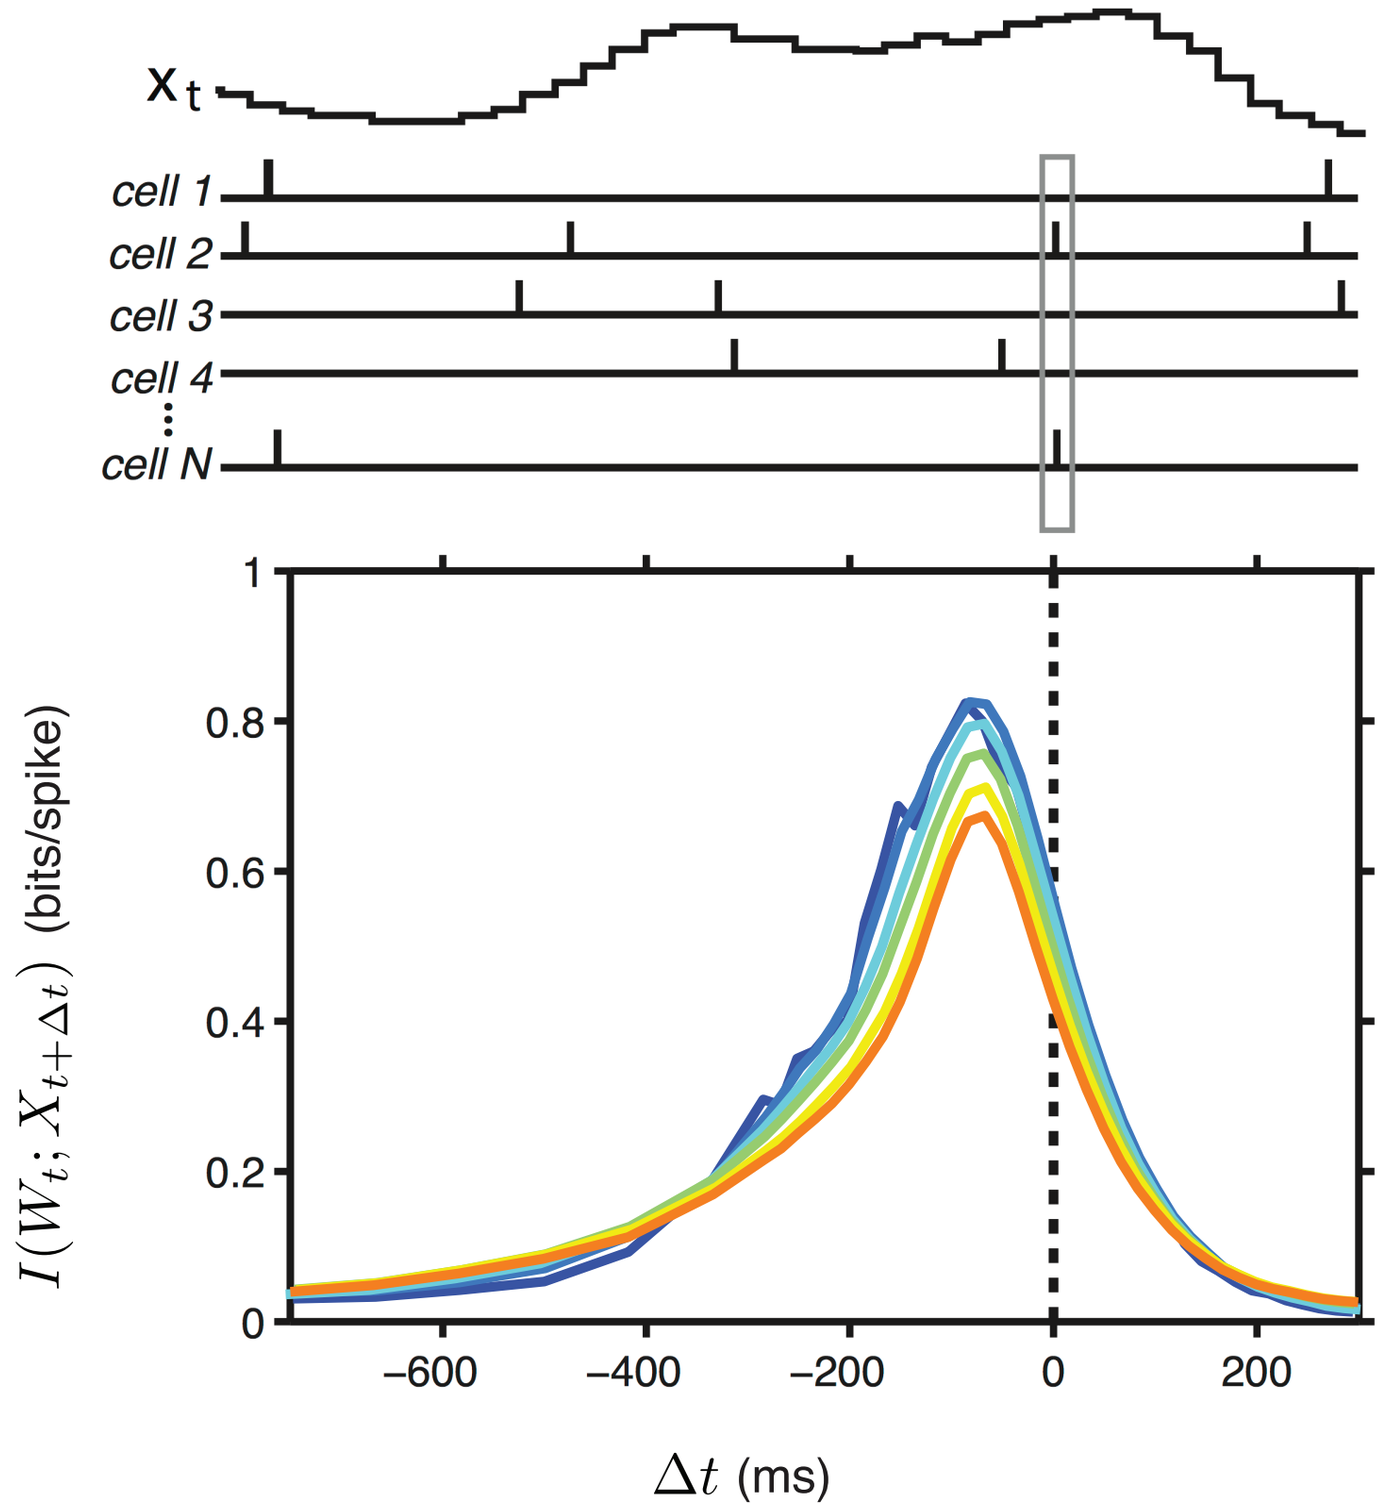

Palmer, Marre, Berry, Bialek PNAS (2015)

Retina populations carry info about the future:

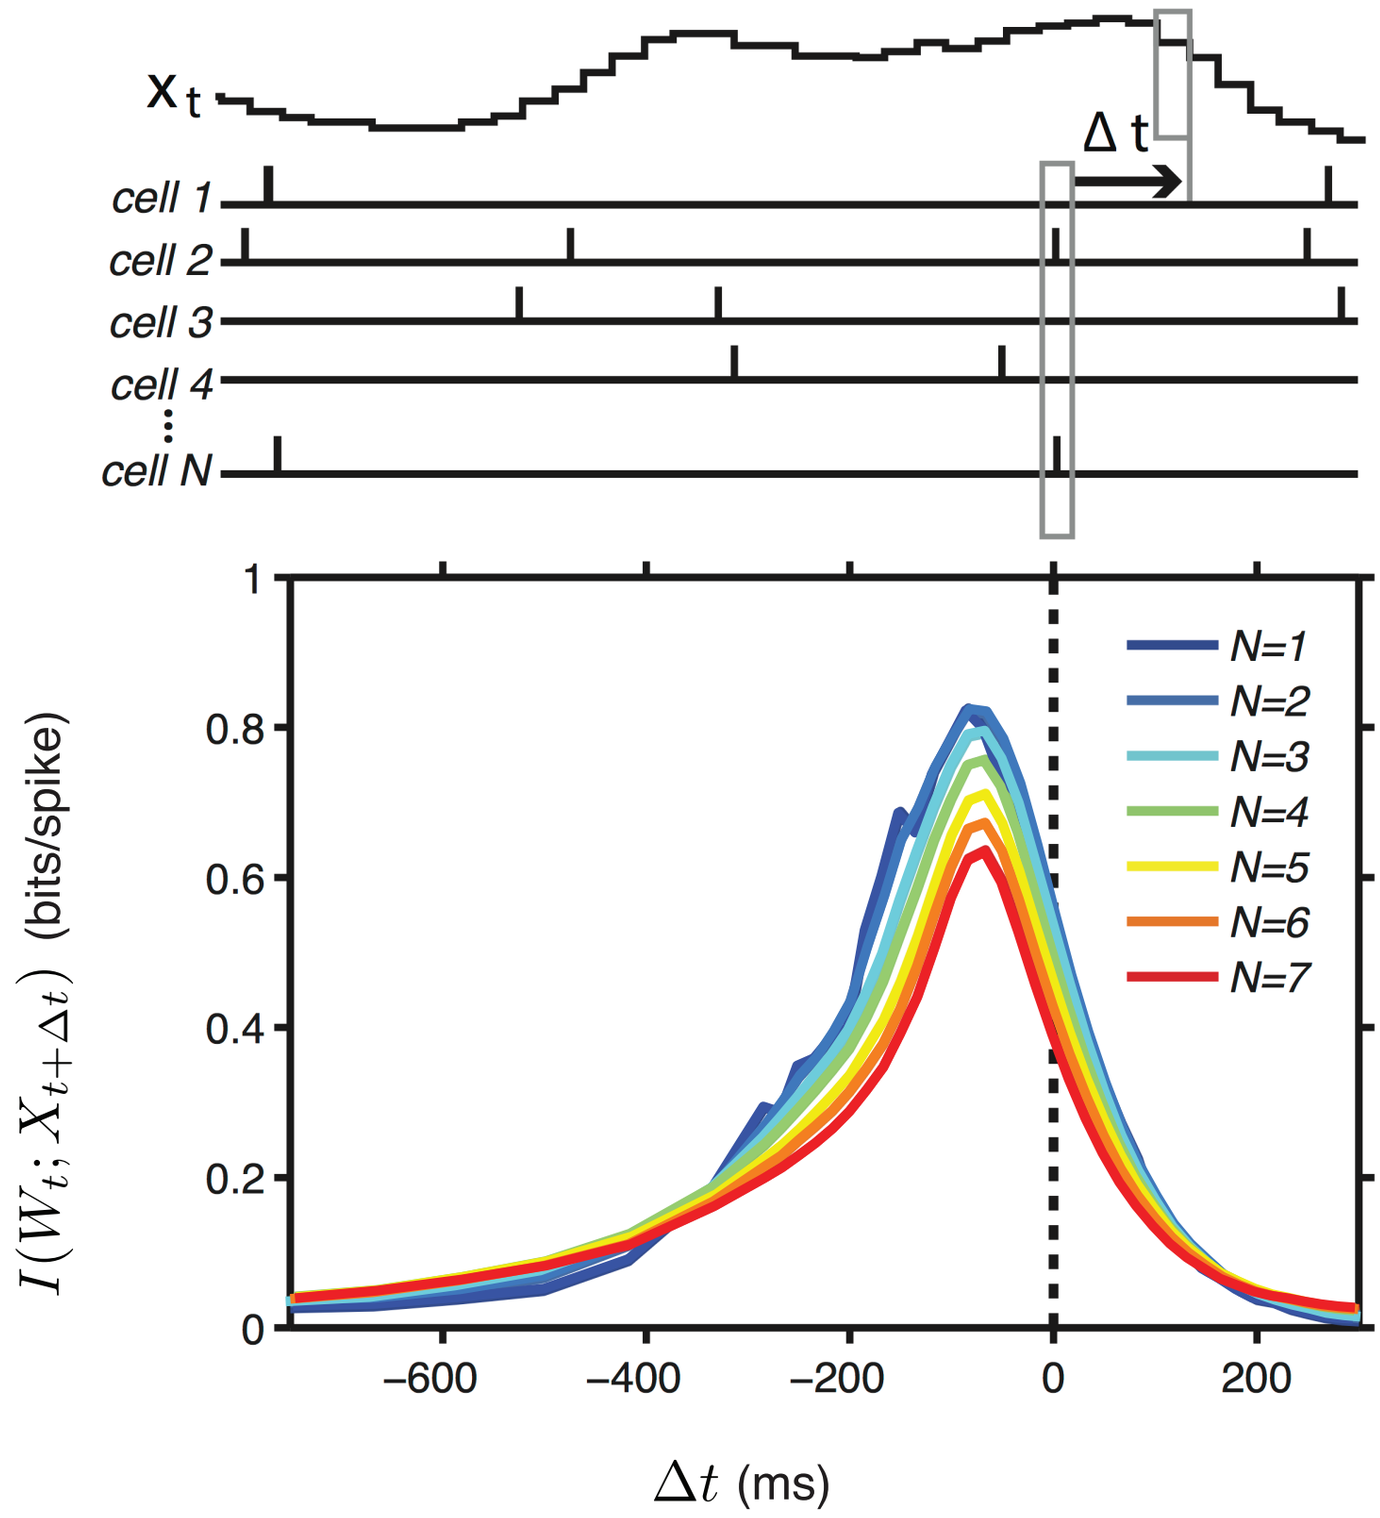

Palmer, Marre, Berry, Bialek PNAS (2015)

We compute the optimal code,
given the stimulus correlations:

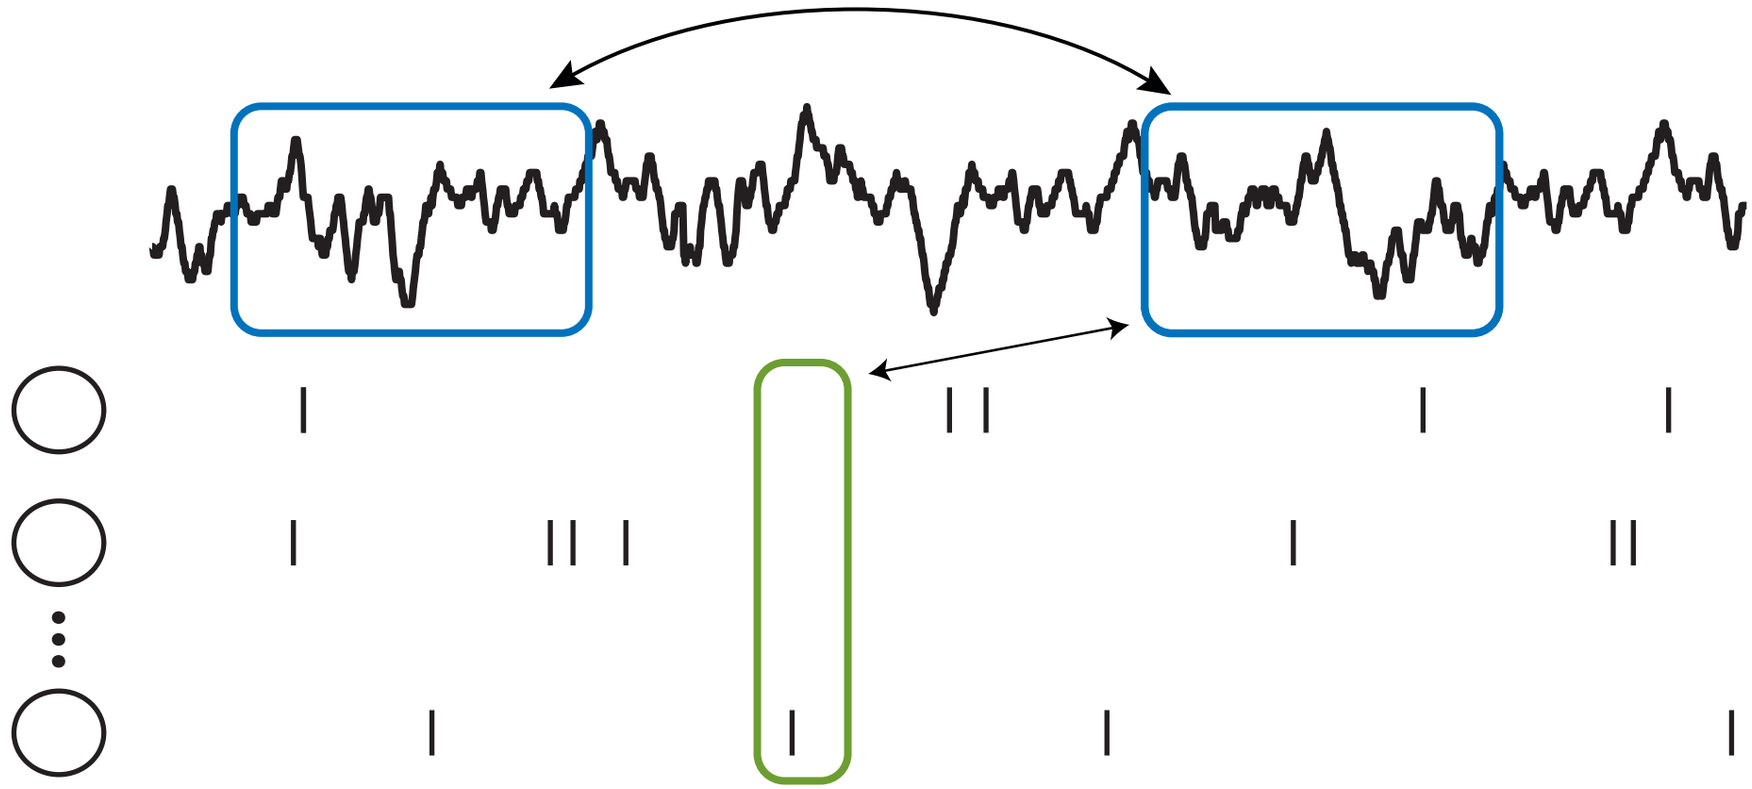

The bottleneck problem:

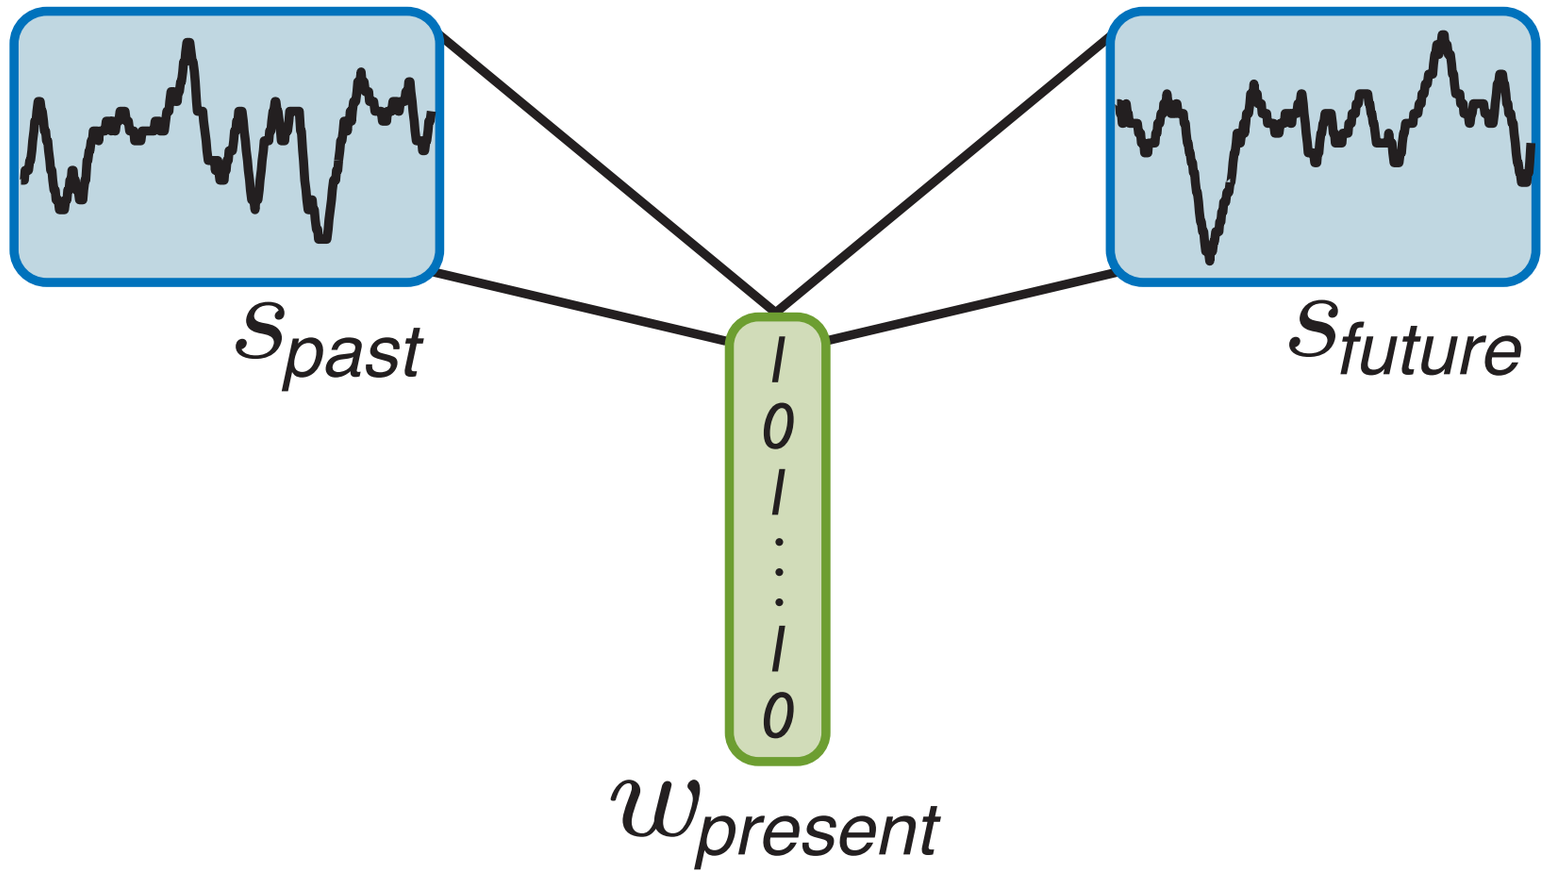

The bottleneck problem:

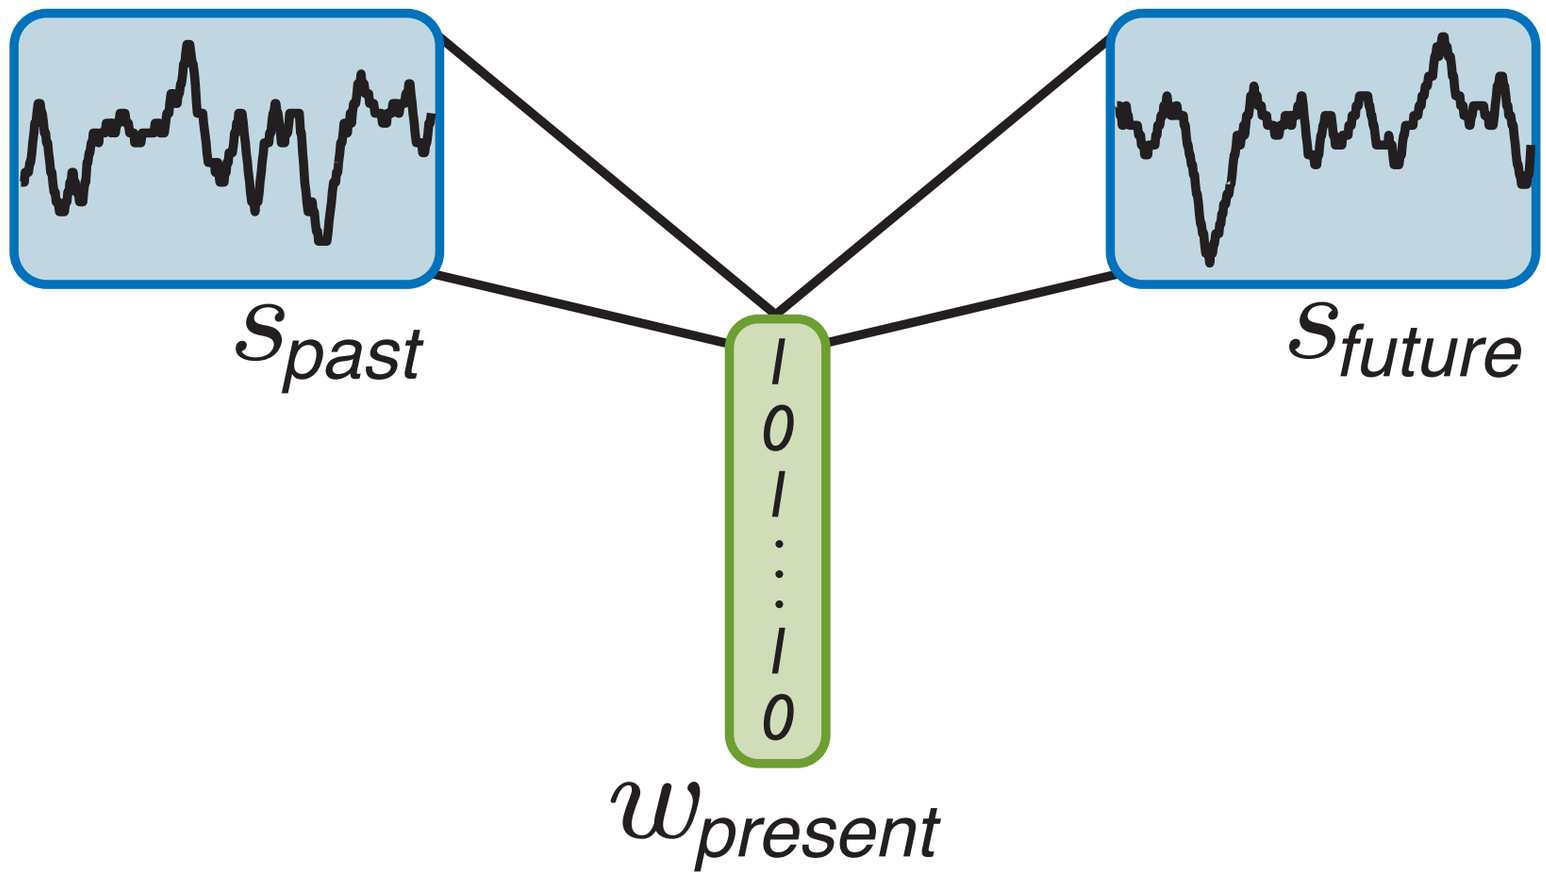

$$L = I_{past}(W_t; \vec{S}_{t-\Delta t}) - \beta I_{future}(W_t; \vec{S}_{t+\Delta t})$$

Optimal compression:

$$L = I_{past}(W_t; \vec{S}_{t-\Delta t}) - \beta I_{future}(W_t; \vec{S}_{t+\Delta t})$$

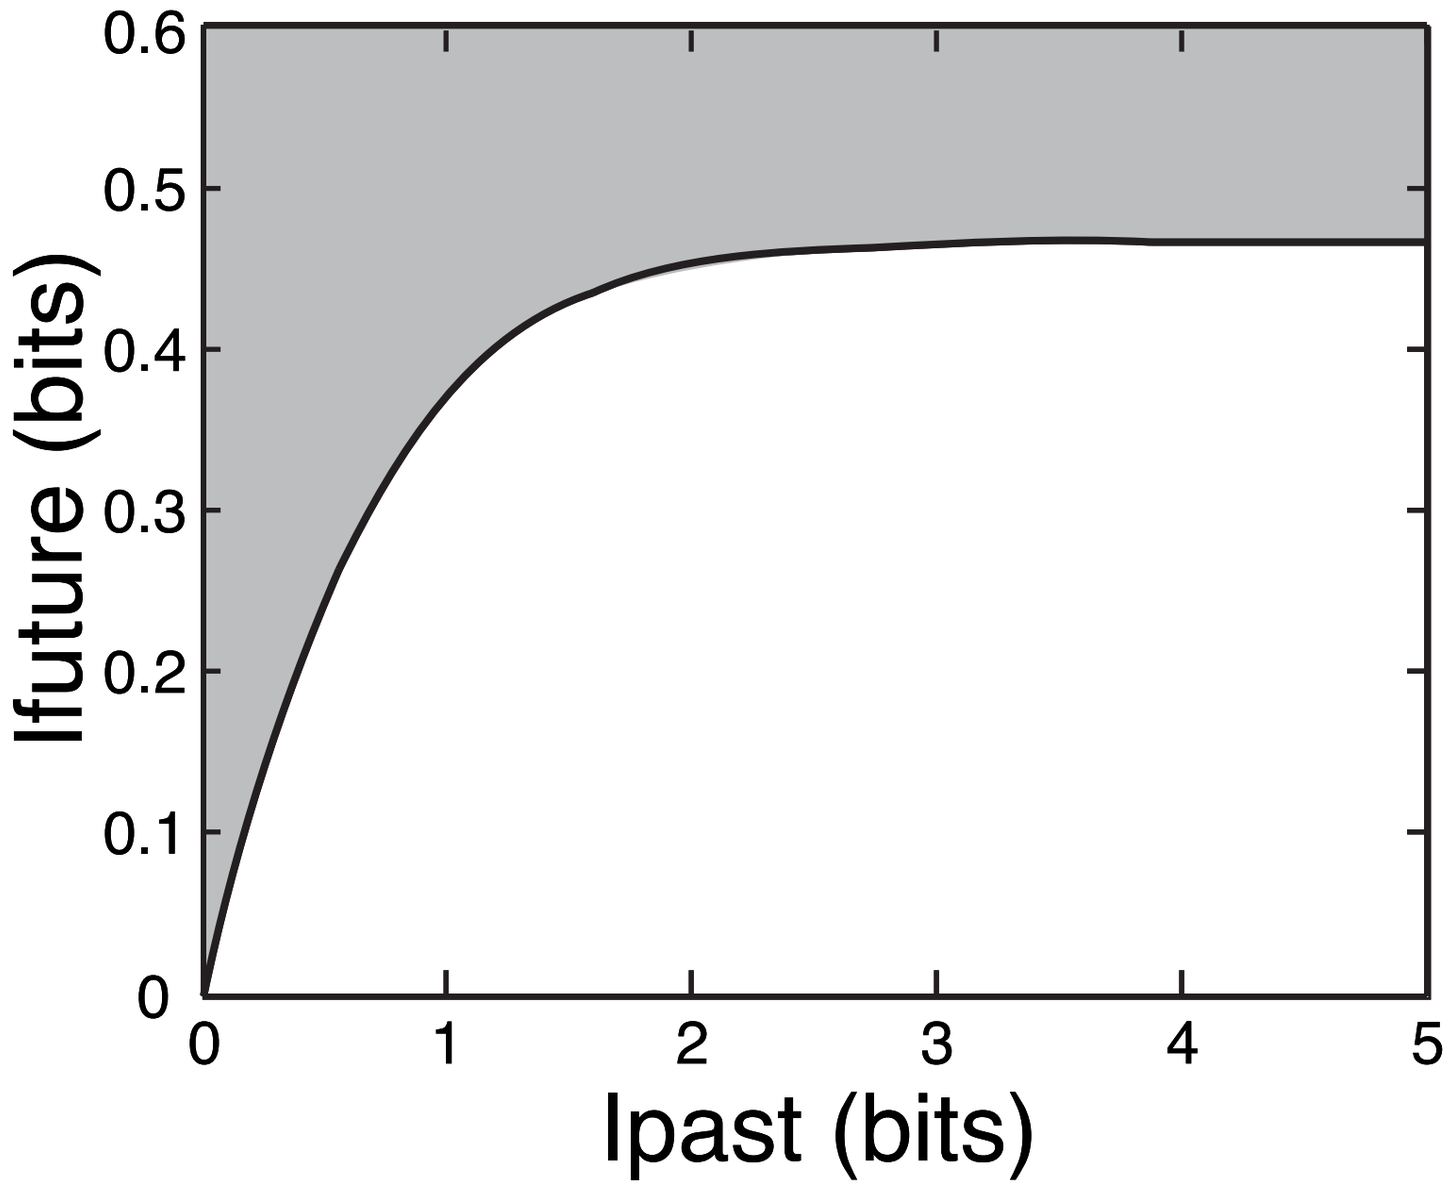

Tishby, Pereira, Bialek (1999)
Bialek, Nemenman, Tishby (2001)
Chechik, Globerson, Tishby, Weiss (2005)

Optimal compression:

$$L = I_{past}(W_t; \vec{S}_{t-\Delta t}) - \beta I_{future}(W_t; \vec{S}_{t+\Delta t})$$

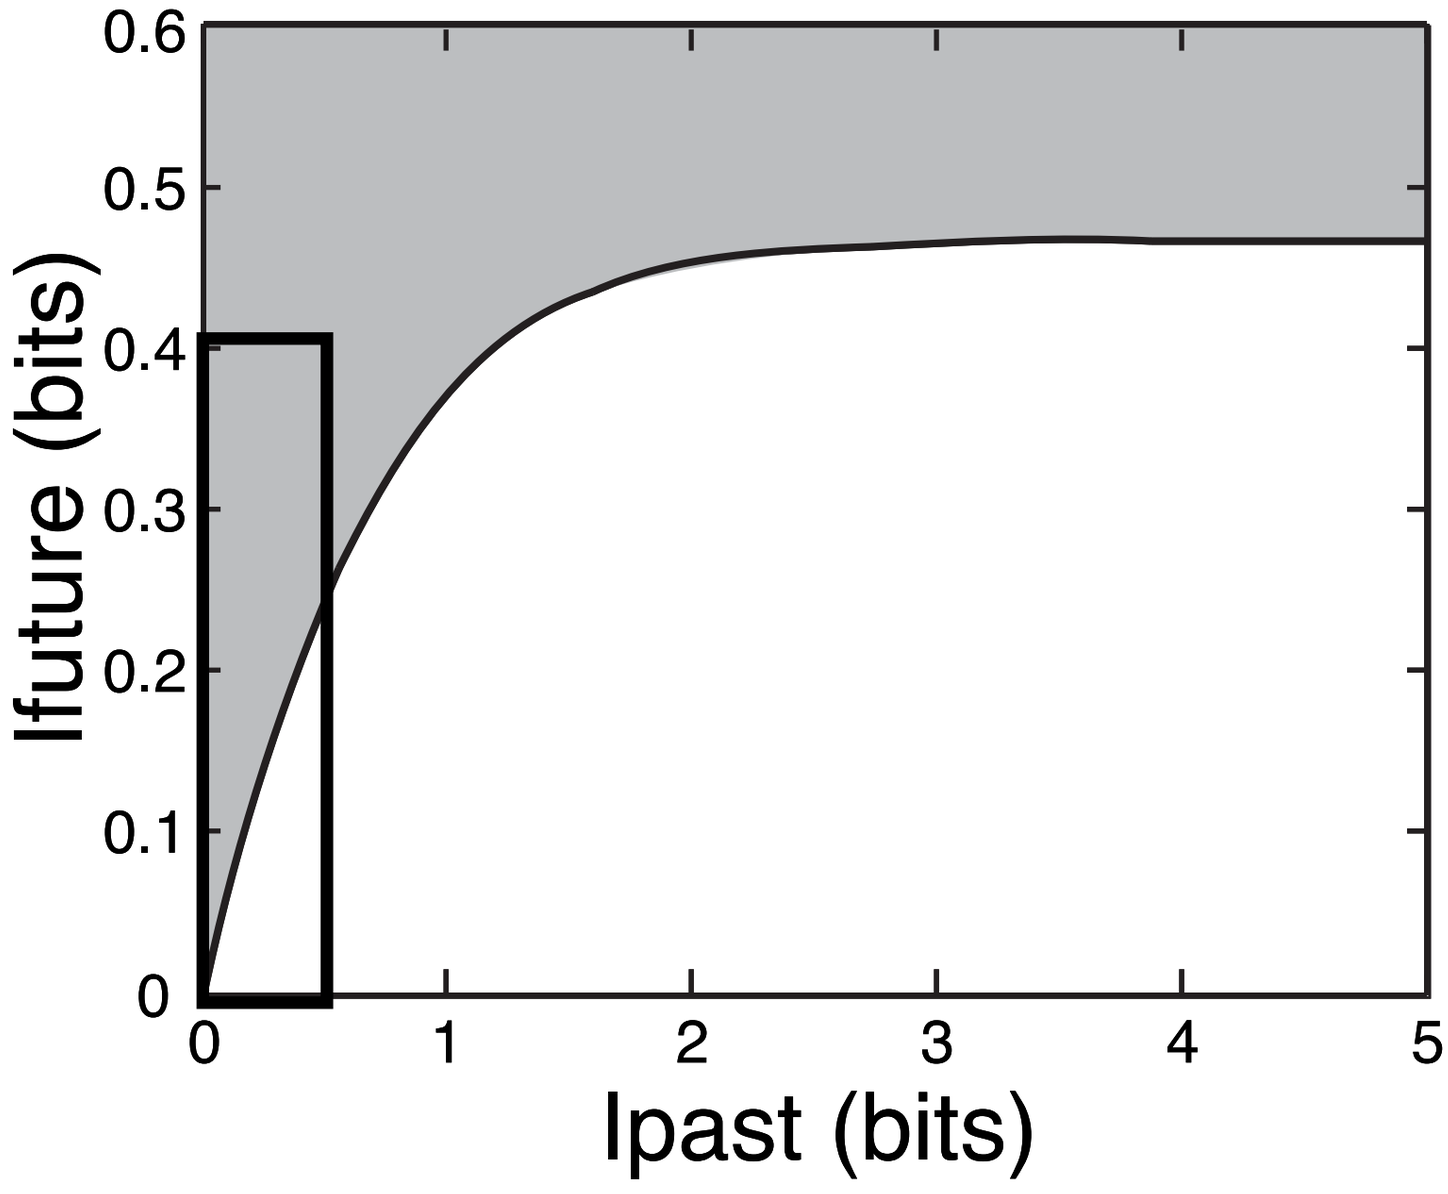

Tishby, Pereira, Bialek (1999)
Bialek, Nemenman, Tishby (2001)
Chechik, Globerson, Tishby, Weiss (2005)

Spiking patterns sit close to the bound:

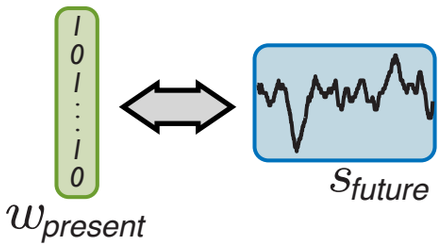

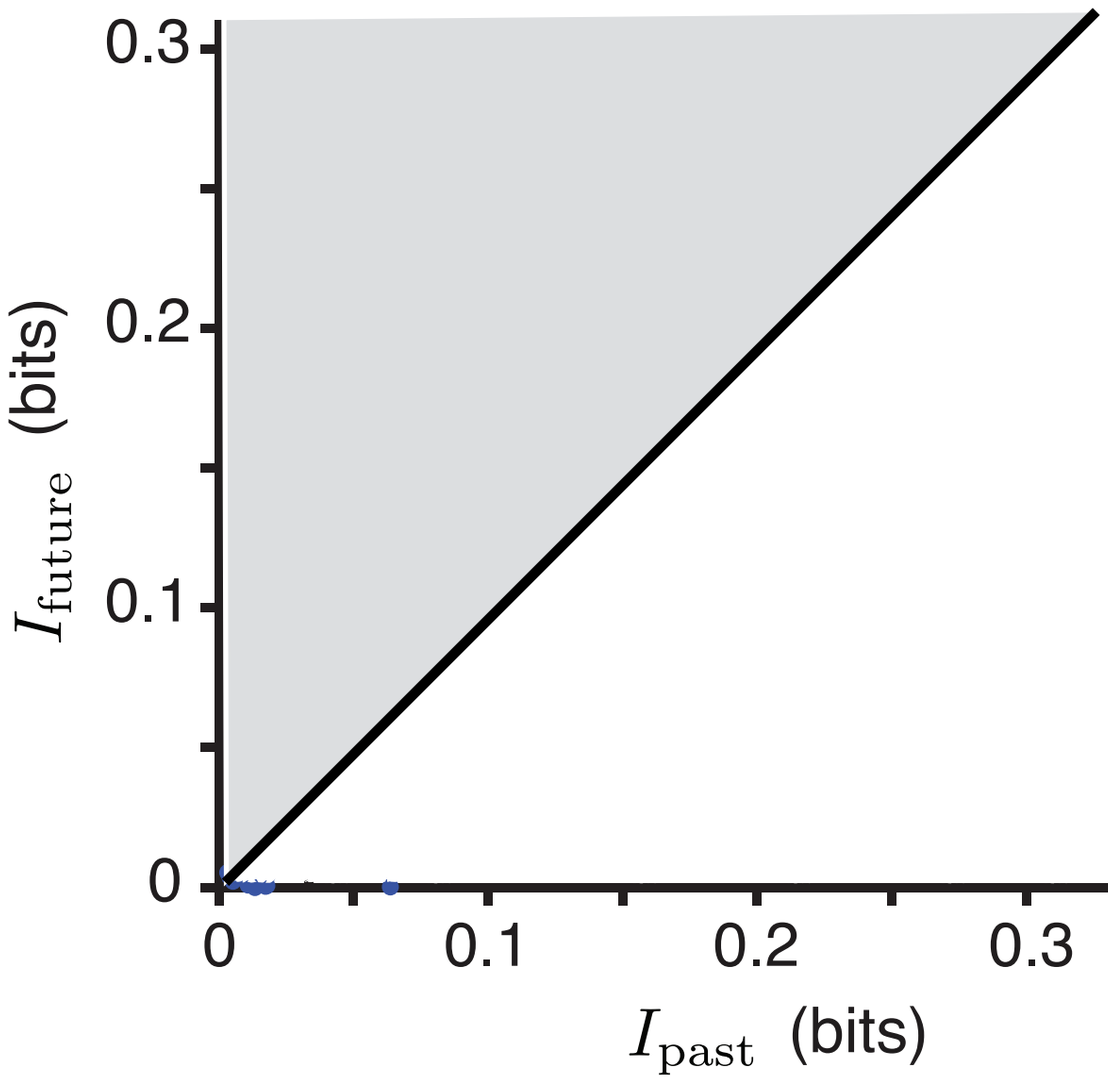

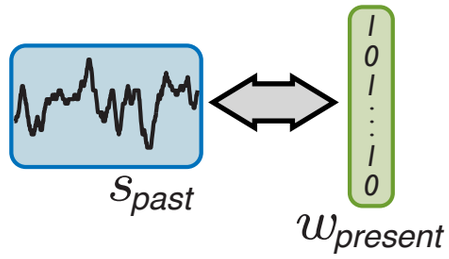

Spiking patterns sit close to the bound:

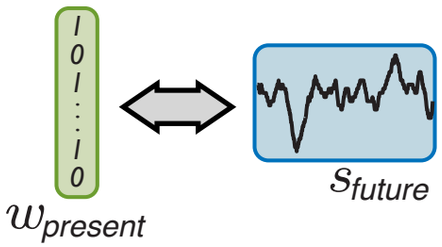

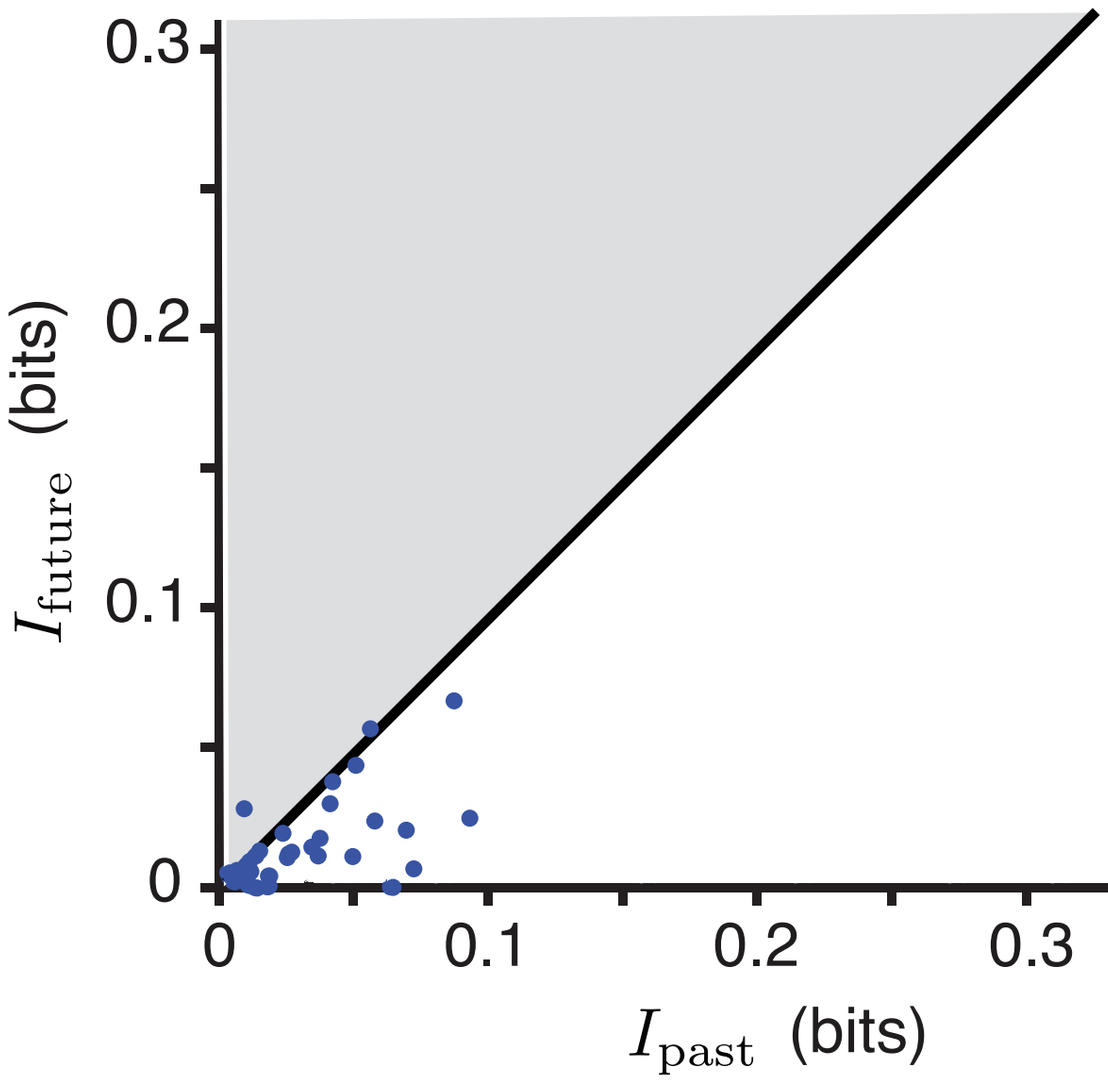

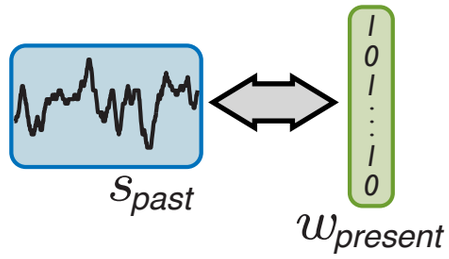

Spiking patterns sit close to the bound:

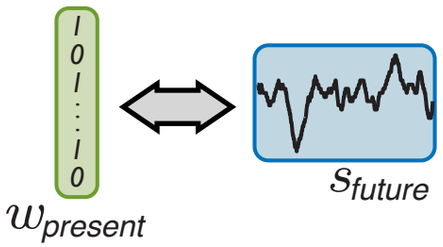

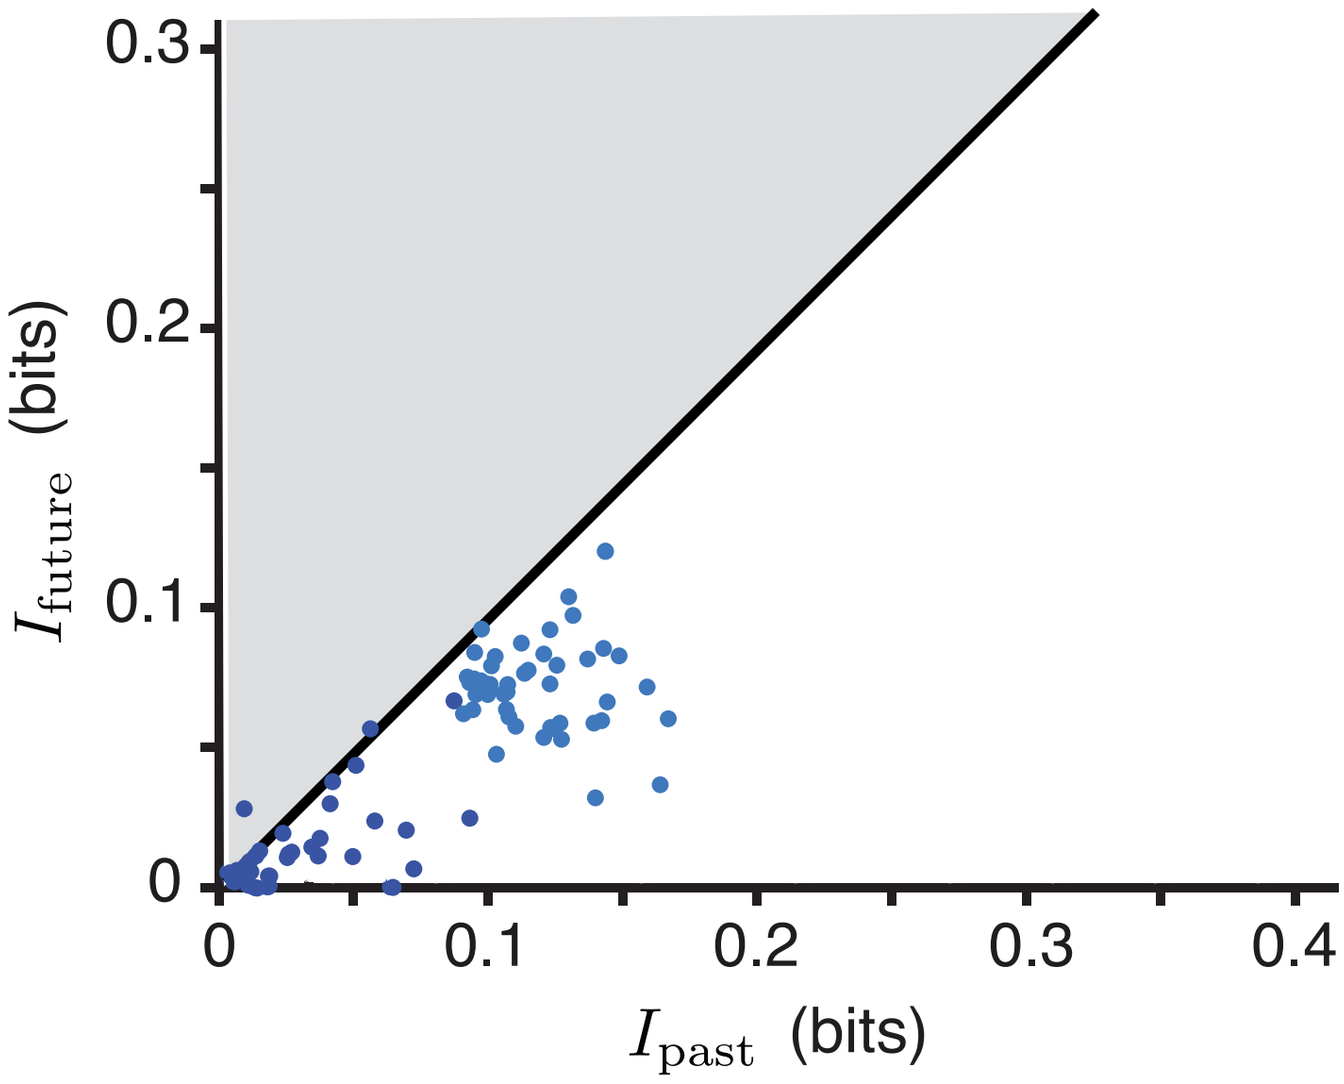

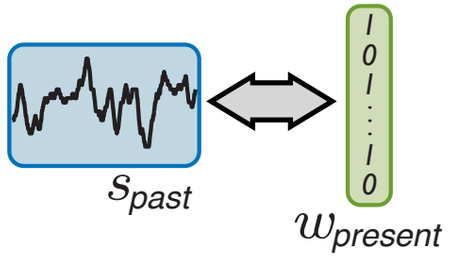

Spiking patterns sit close to the bound:

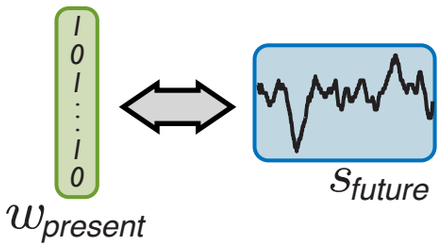

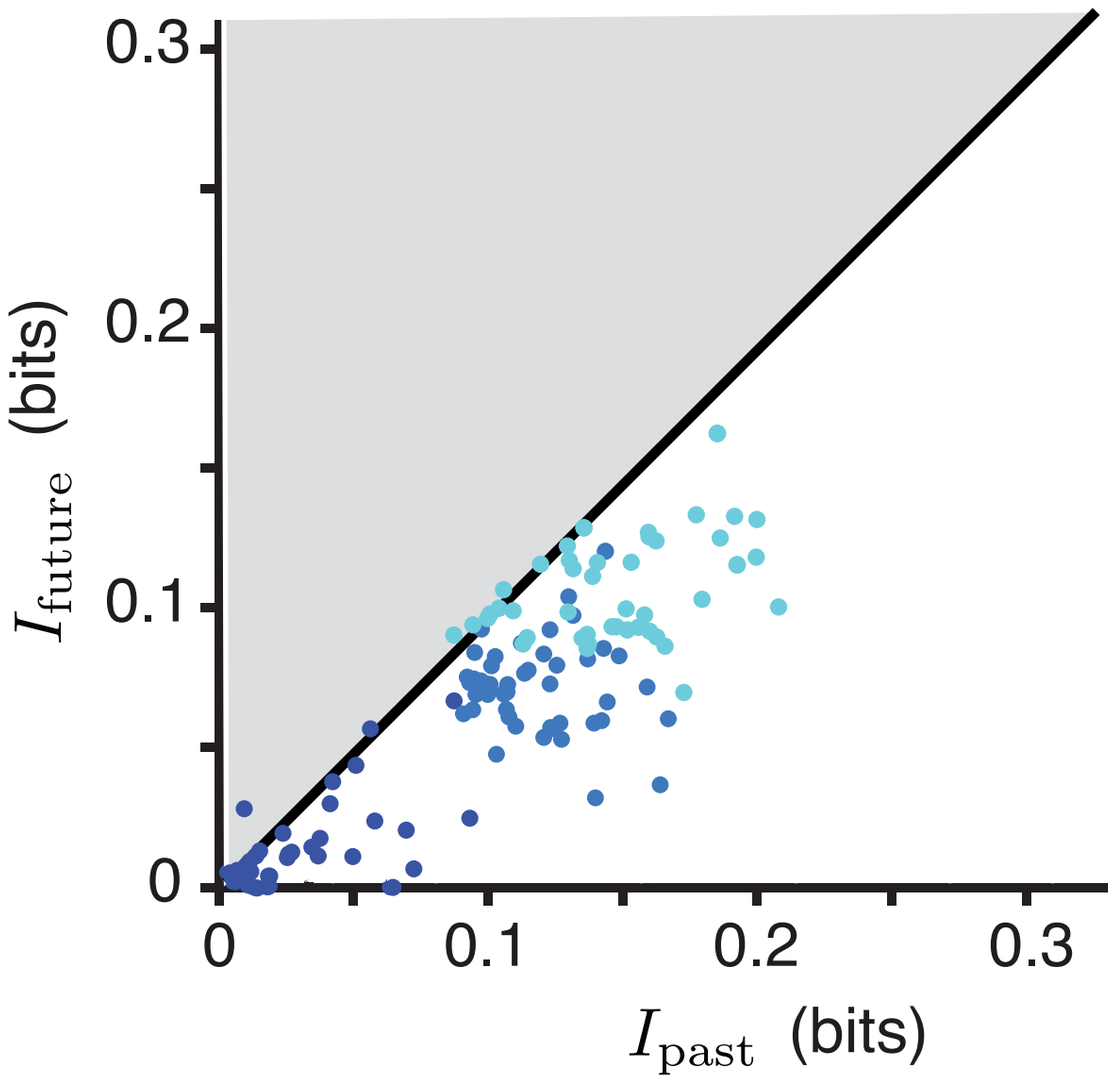

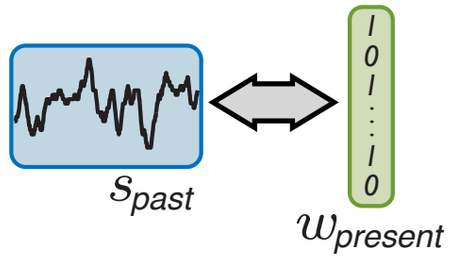

Spiking patterns sit close to the bound:

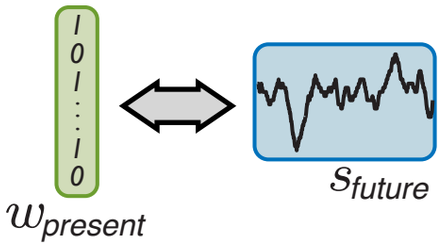

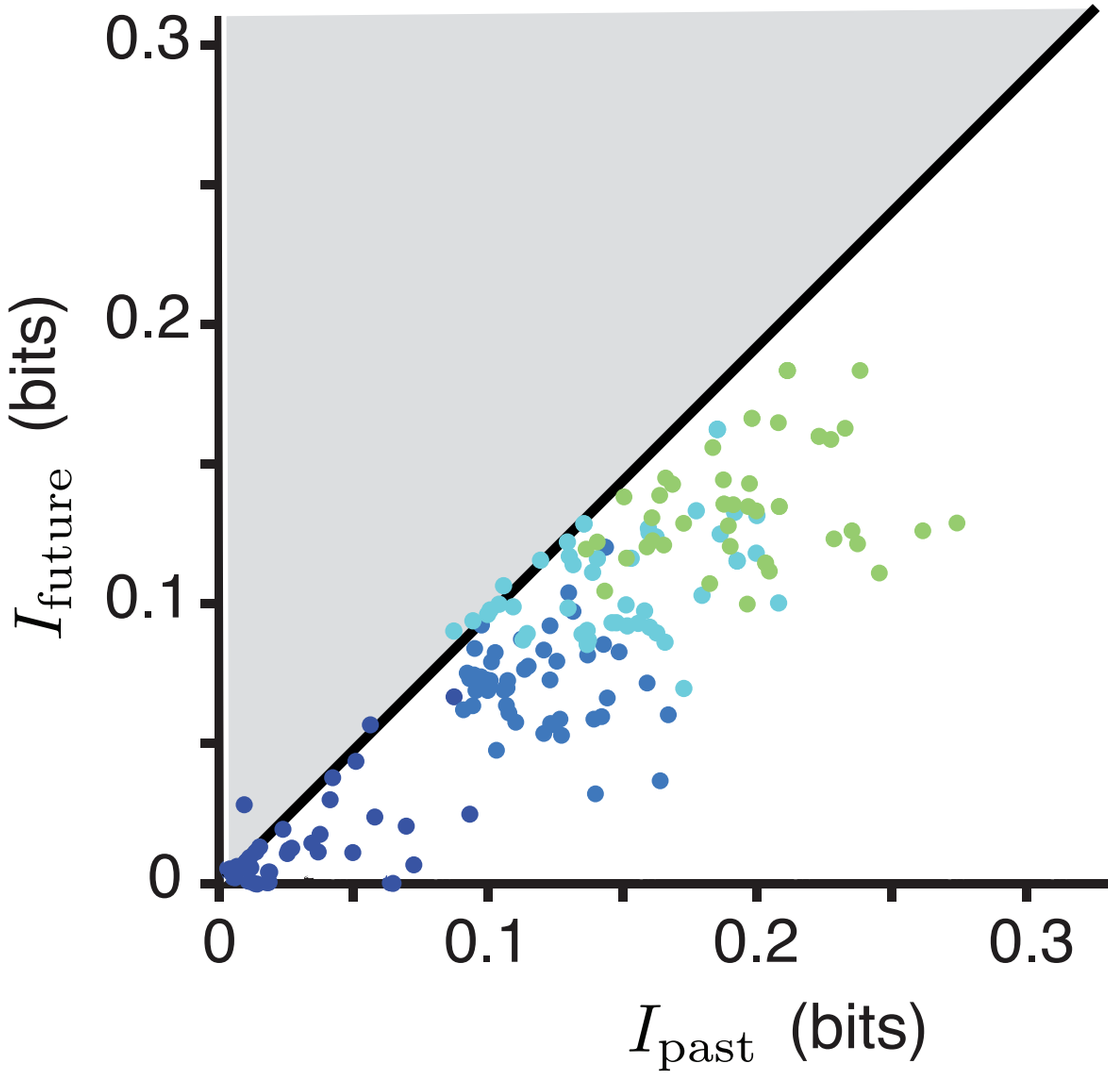

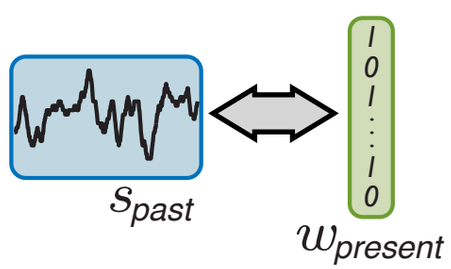

Spiking patterns sit close to the bound:

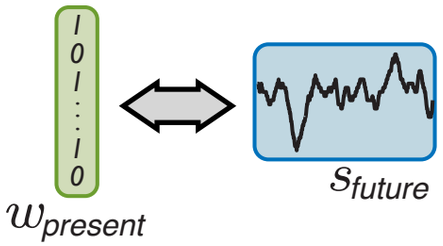

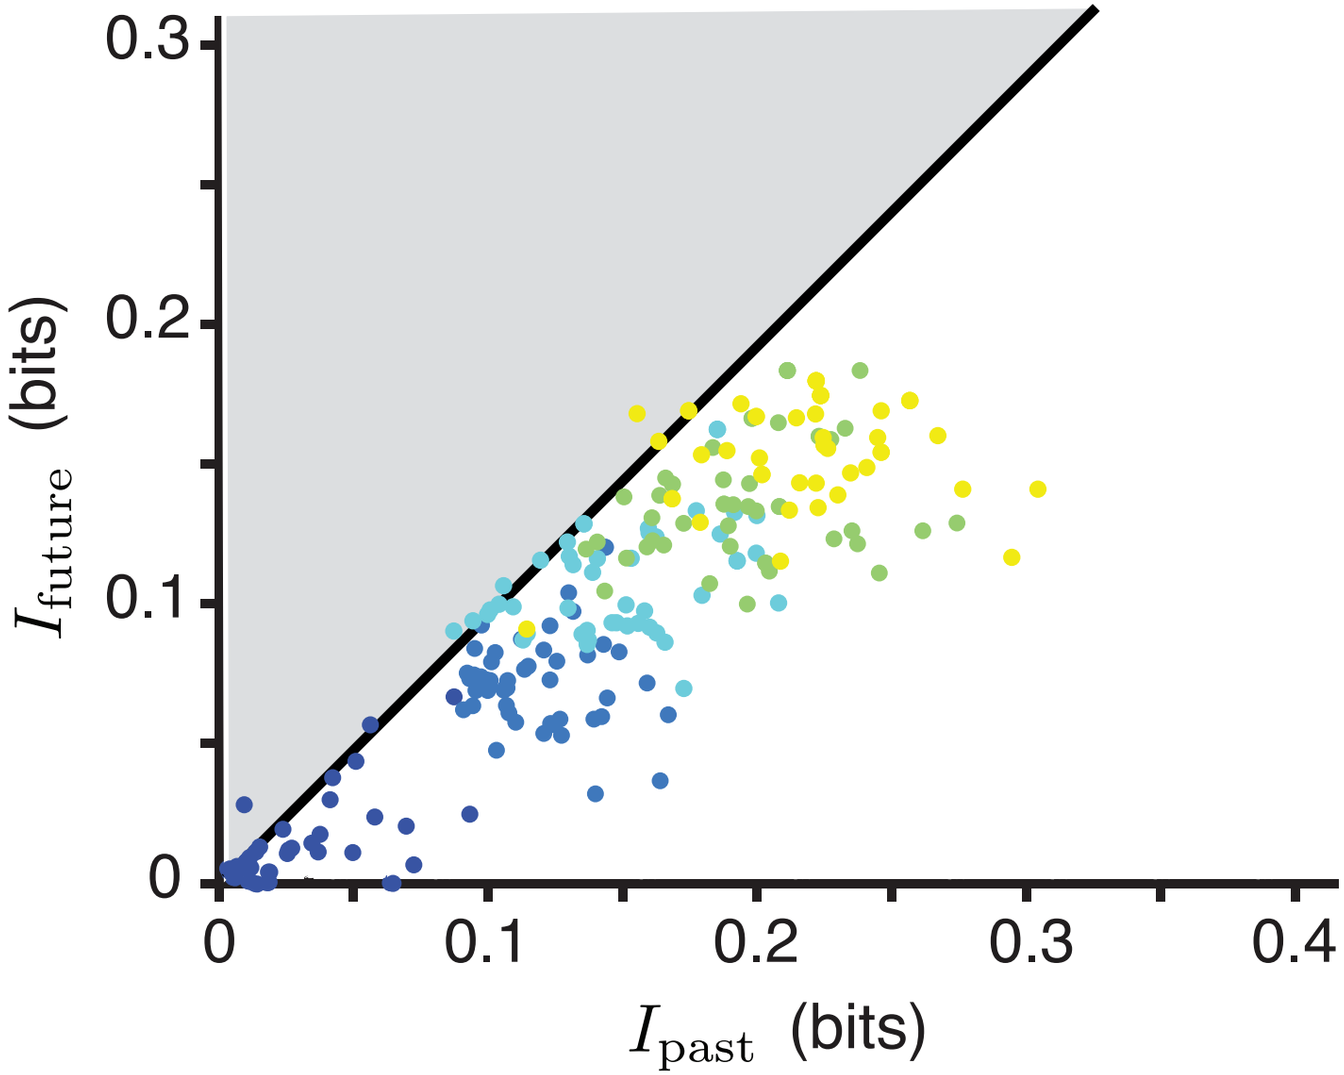

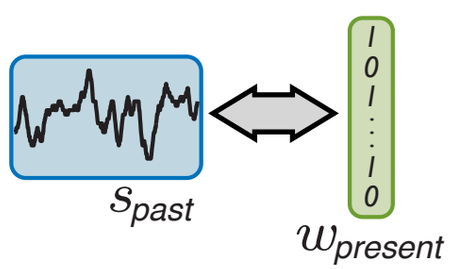

Spiking patterns sit close to the bound:

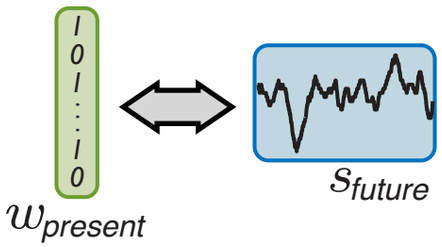

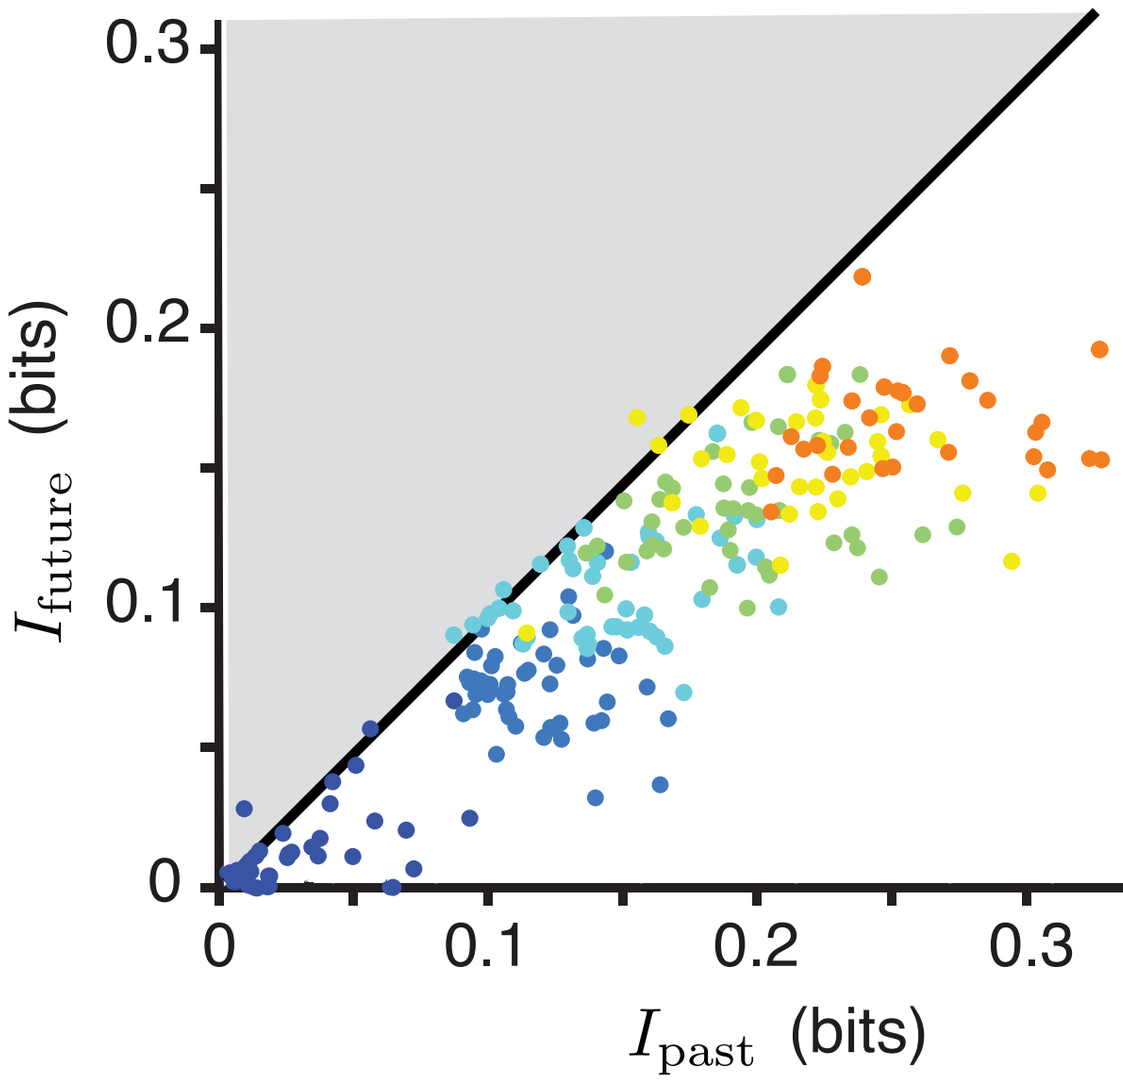

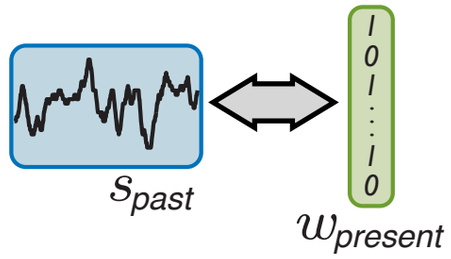

Spiking patterns sit close to the bound:

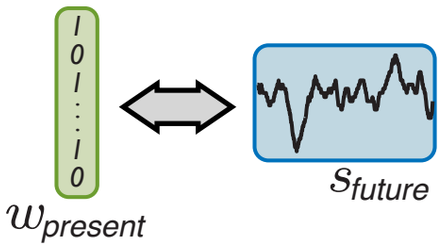

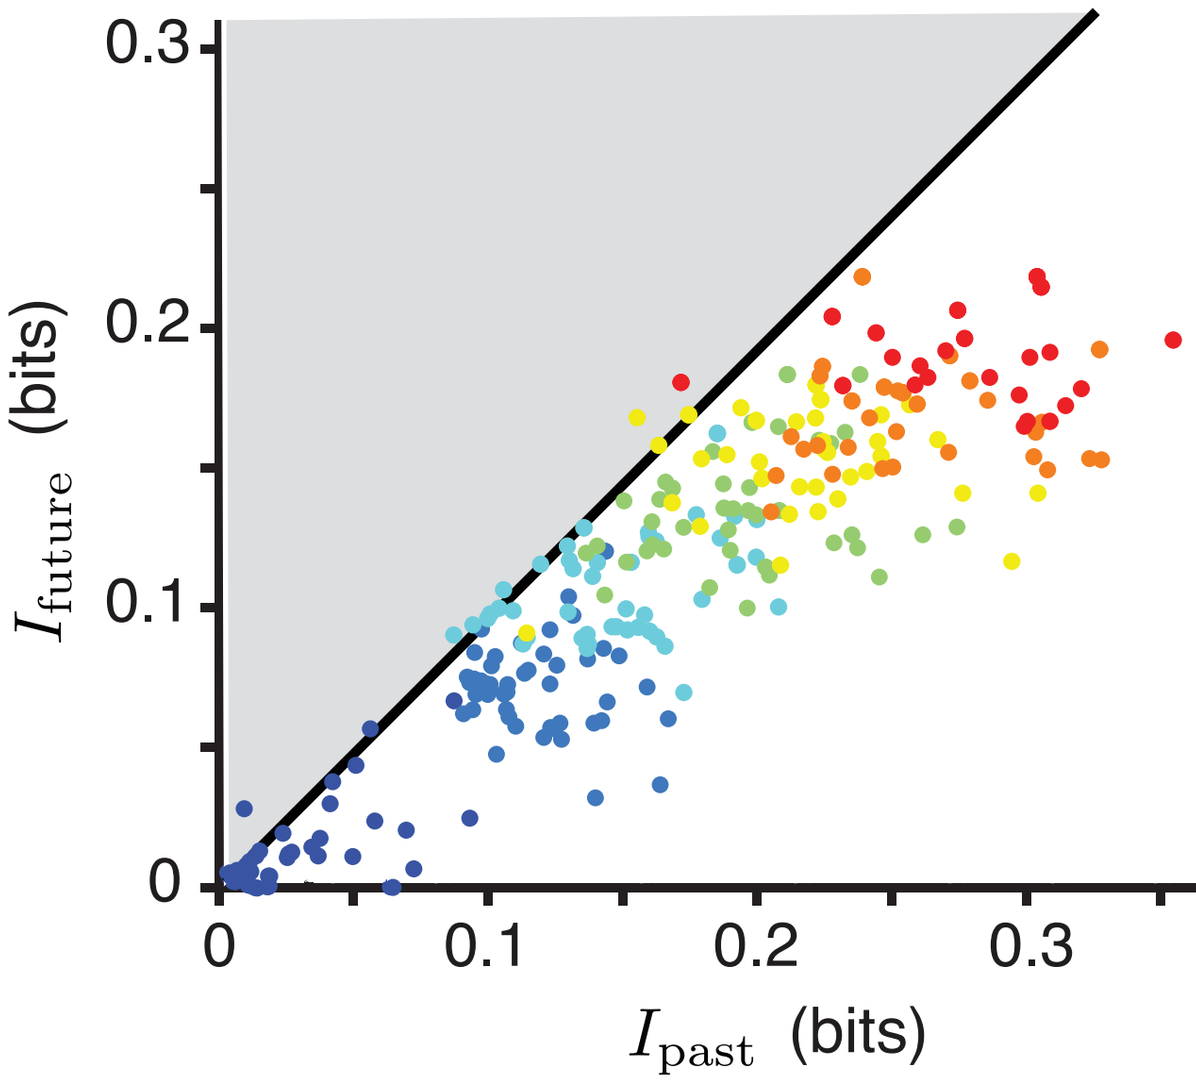

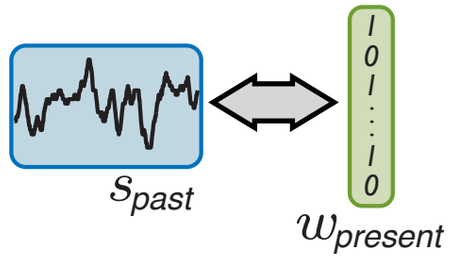

Retinal populations saturate the bound:

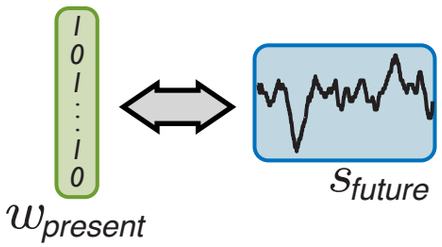

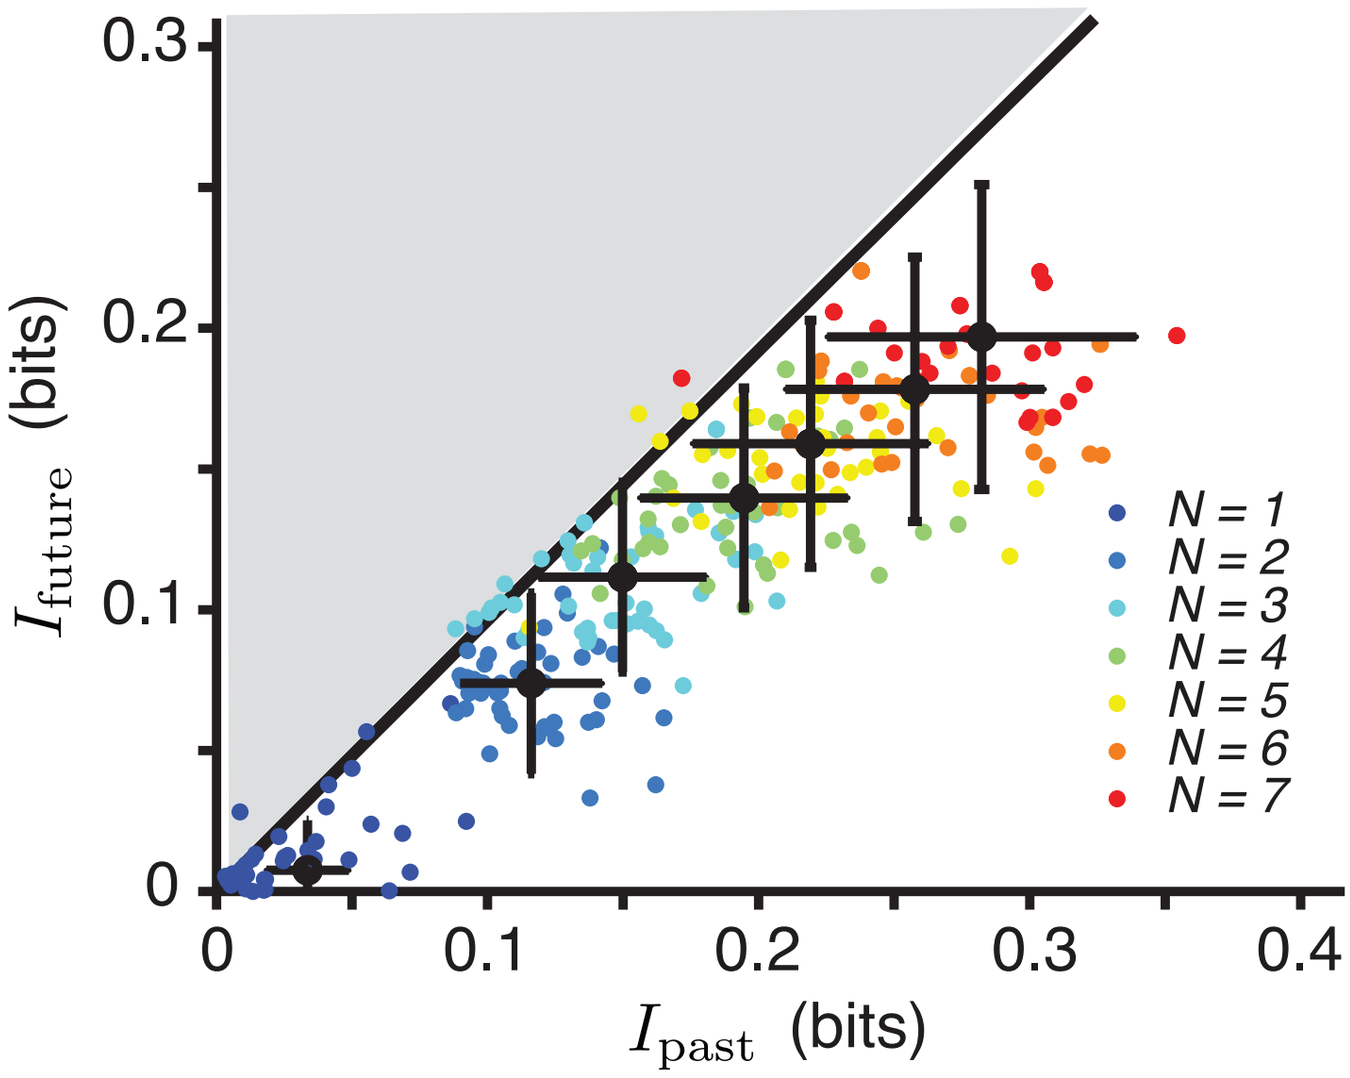

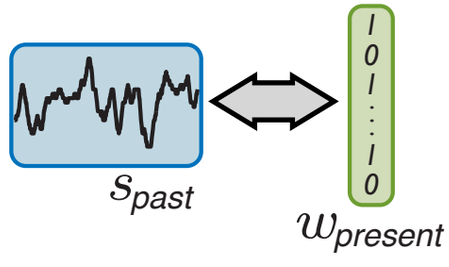

This doesn't work with just simple linear filters:

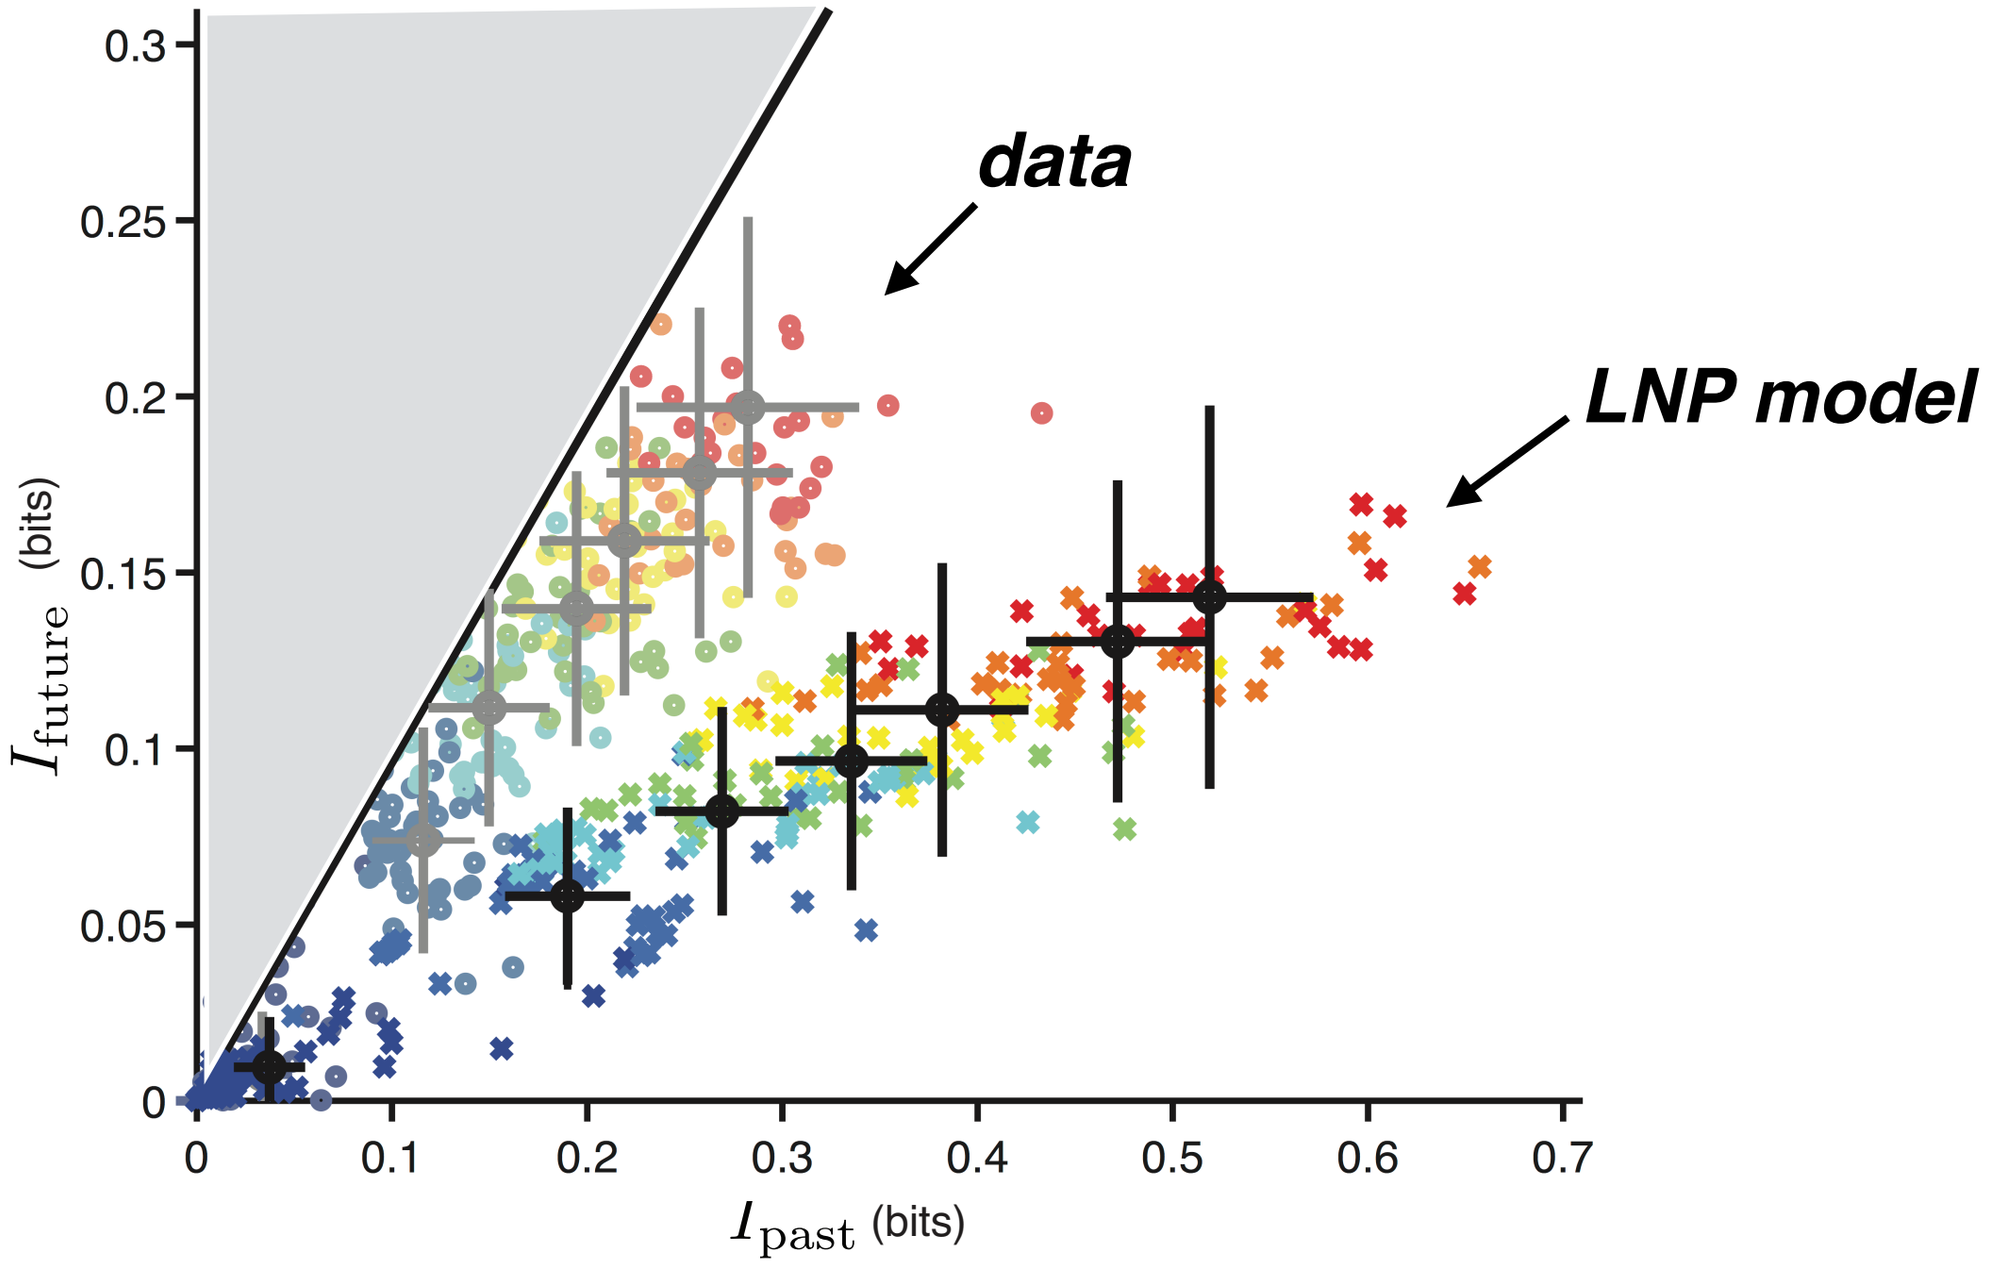

Palmer et al. PNAS (2015)

Recording from the rat retina:

wired.co.uk

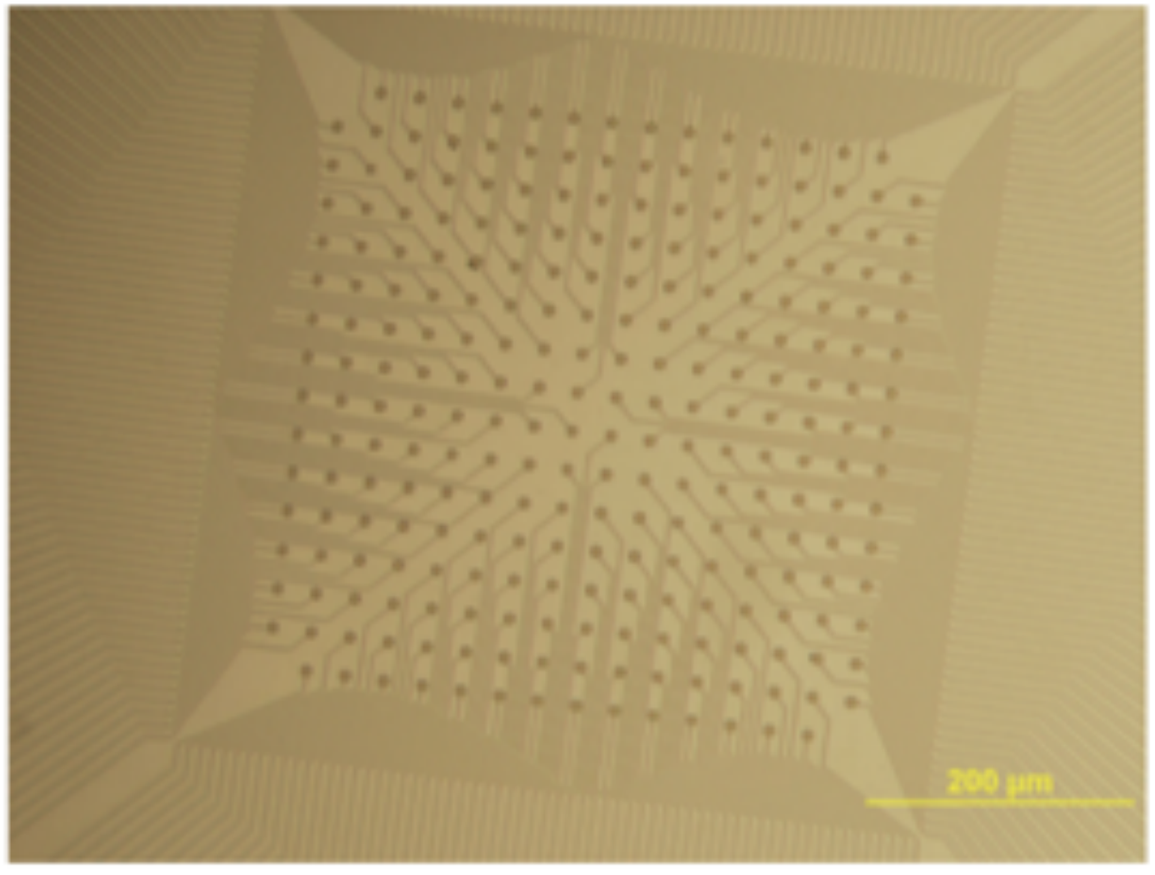

data from the Marre Lab, Institute of Vision, Paris VI University, France

Compute the bound on the predictive info:

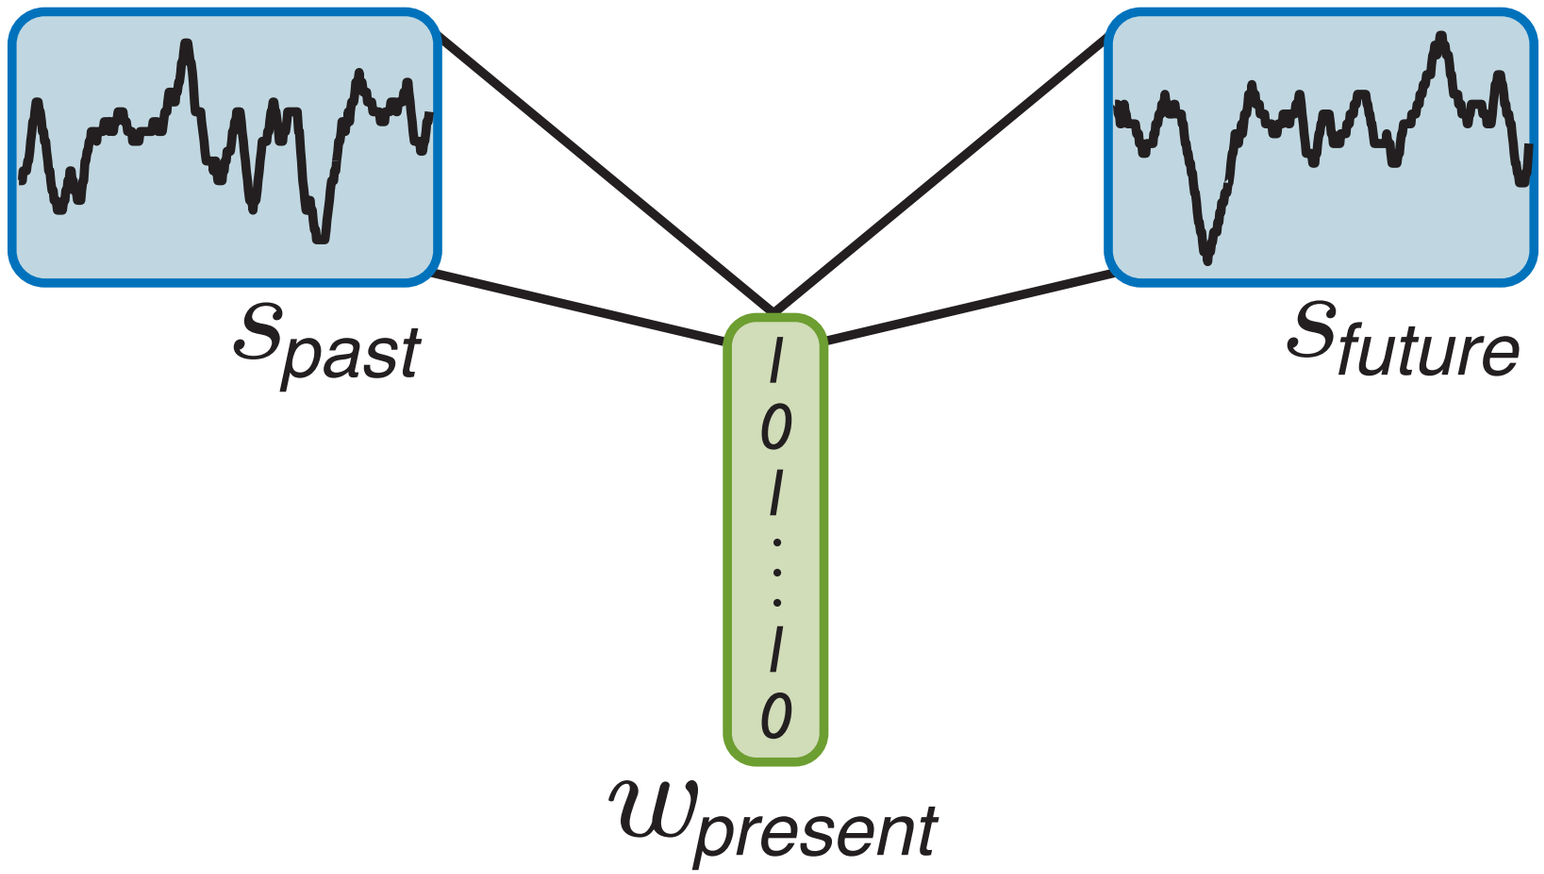

$$L = I_{past}(W_t; \vec{S}_{t-\Delta t}) - \beta I_{future}(W_t; \vec{S}_{t+\Delta t})$$

Performance relative to the bound:

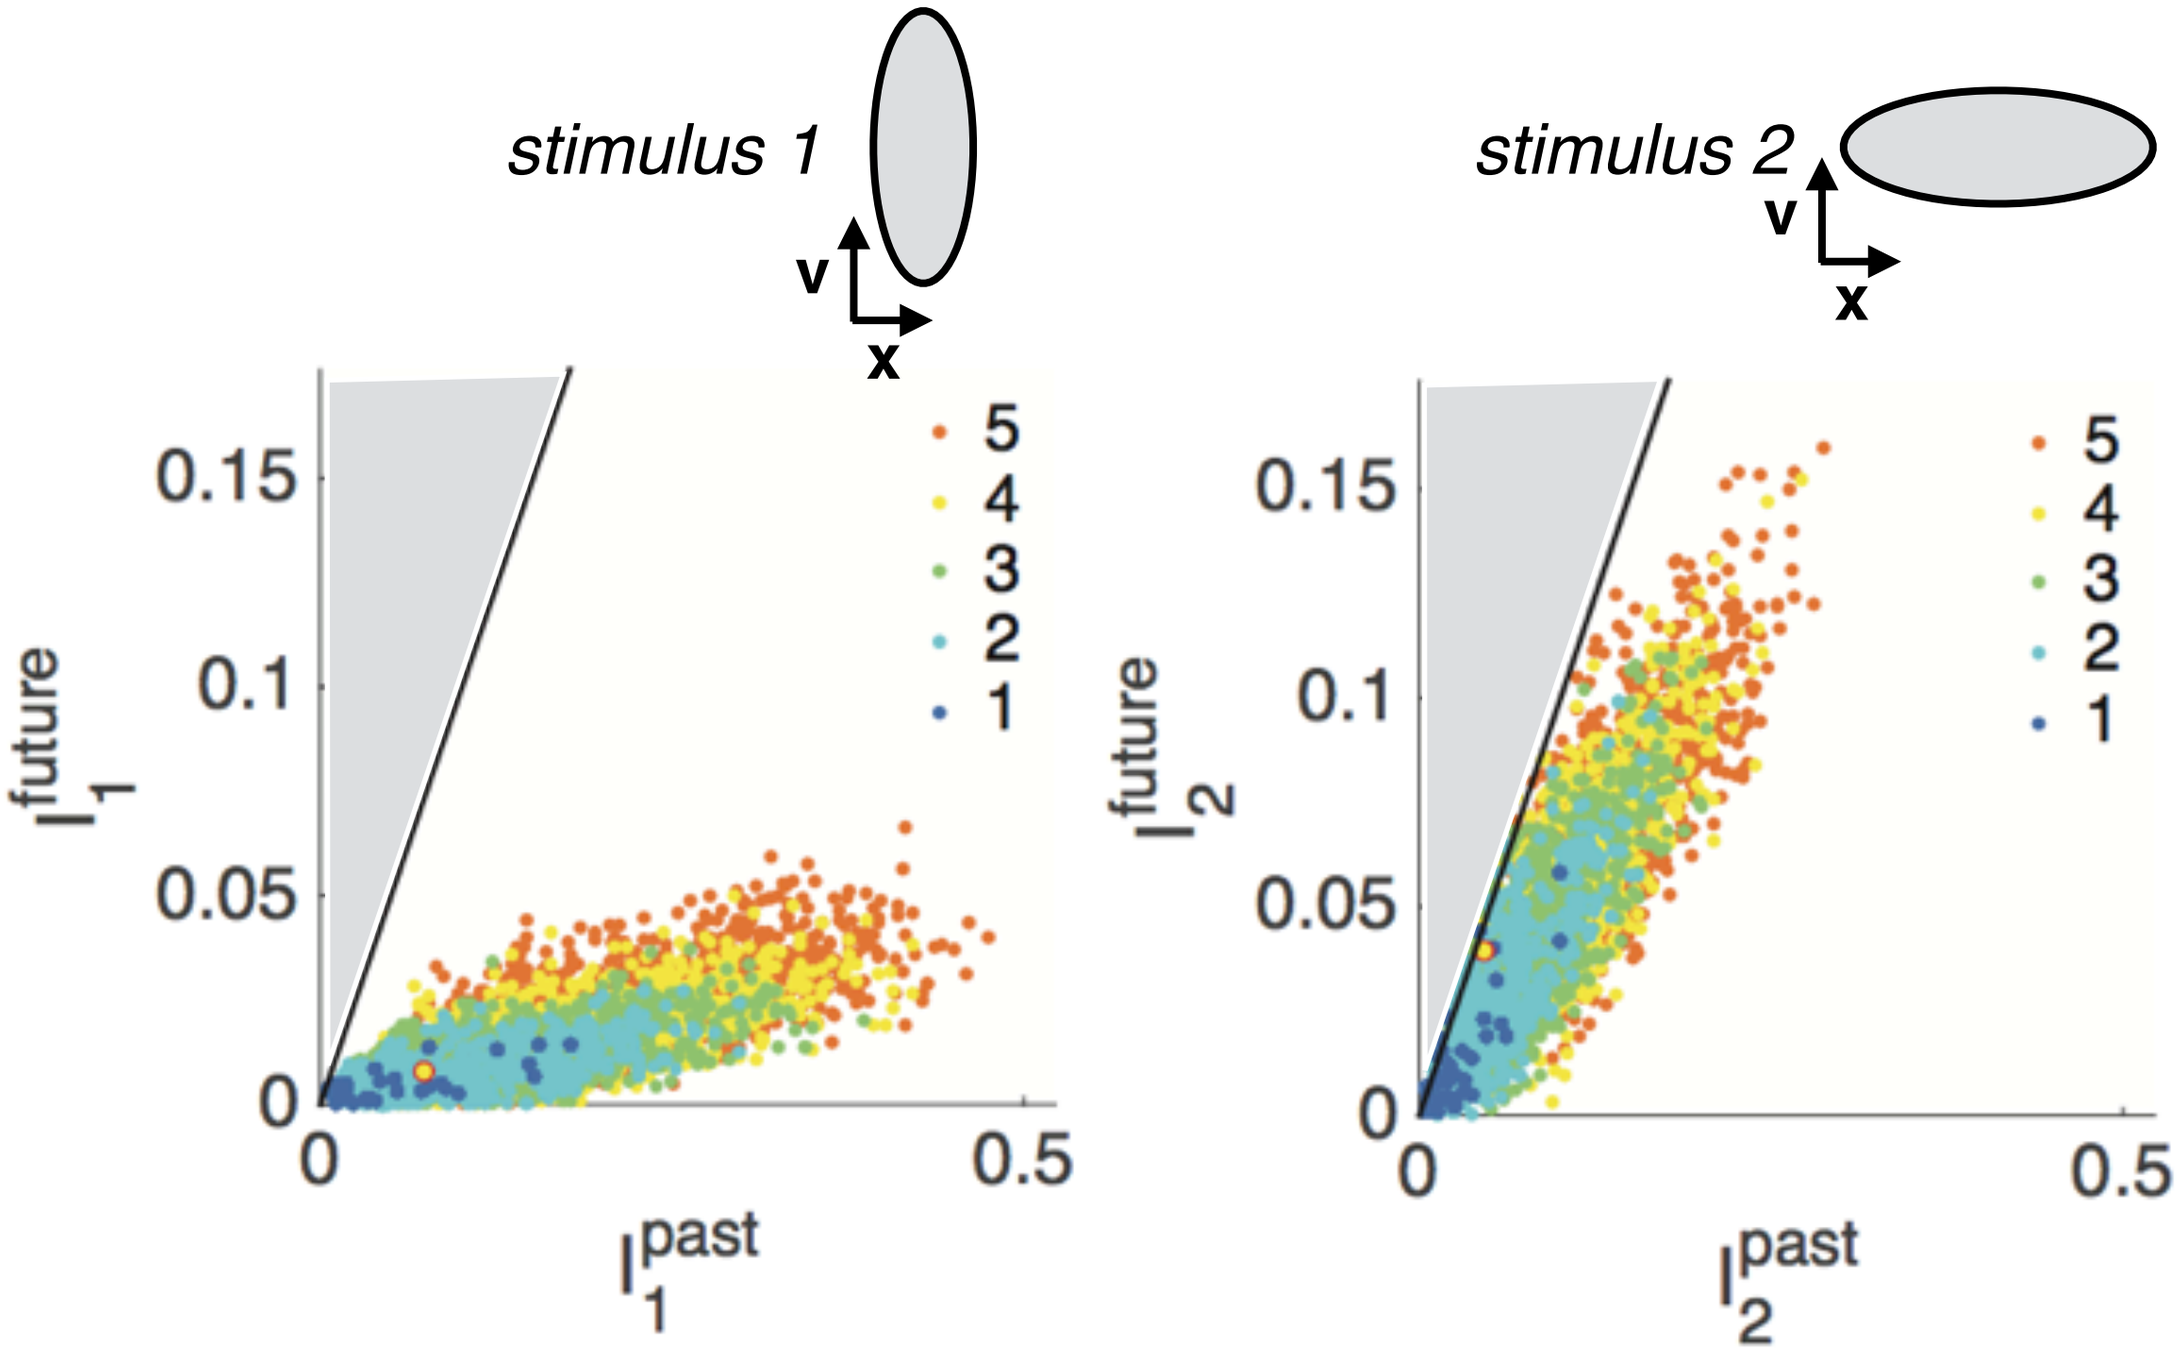

An example group of four cells:

$$I\left(W_{t-\Delta t}; F_t\right) = H\left(W_{t-\Delta t}\right) + H\left(F_t\right) - H\left(W_{t-\Delta t}, F_t\right)$$

N = 4

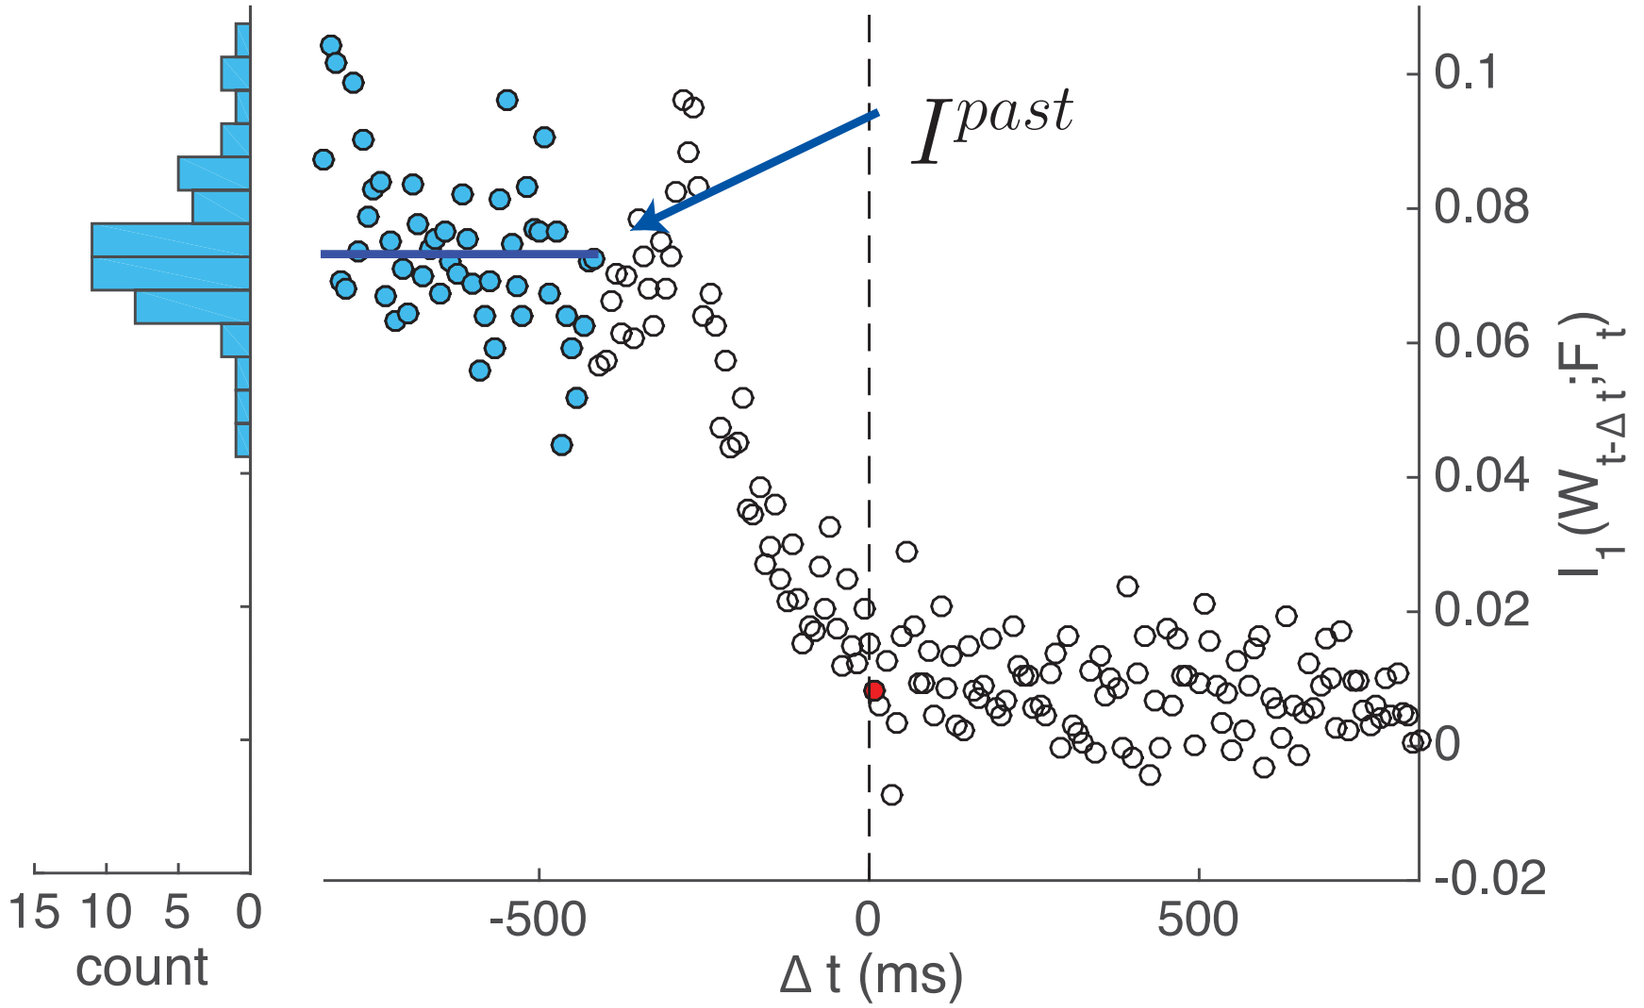

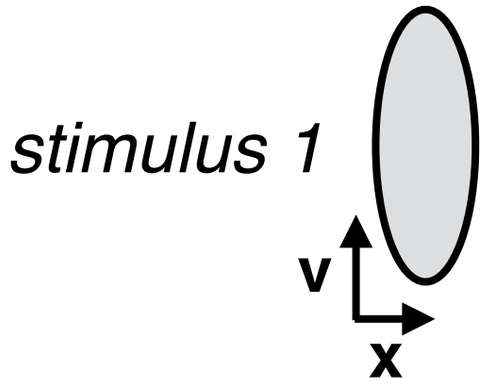

An example group of four cells: