Approximate Bayesian Computation for Parameter Inference

Clayton W. Seitz

May 26, 2022

Approximate Bayesian Computation (ABC)

Beating the curse of dimensionality for parameter inference

We want to avoid needing to train new networks for every parameter value θ (expensive!) then computing the likelihood of the experimental data. Instead we compute the likelihood of simulated data under set of target variational distributions trained on experimental data. This assumes simulations and experimental data are "exchangeable" when computing the posterior

Variational step: we learn N_t target distributions by training a deep network on the experimental data. In this way we have N_t variational target distributions

ABC step: We sample parameters from our prior $\theta \sim \pi(\theta)$, and produce N Monte Carlo trajectories $\mathbf{x}(t)$. We compute the likelihood of the simulated trajectory with a tolerance ϵ (with a tolerance schedule). This replaces the distance metric in ABC with a variational likelihood