A brief introduction to deep generative models

Clayton W. Seitz

February 7, 2022

Outline

The logic of generative modeling

Say we have a set of variables $\mathrm{x} = (x_1, x_2, ..., x_n)$ which might have some statistical dependence

The variable x might be an amino acid sequence, gene expression data, microscopy image, etc.

- Often we are handed a batch of empirical samples $\{\mathrm{x}_i\}_{i=1}^{N}$
- We want to know the generating distribution $p(\mathrm{x})$

In supervised generative learning, we try to explicity learn the joint distribution $p(\mathrm{x}) = \prod_{i=1}^{N-1} p(x_i|x_{i+1:N})p(x_N)$, which is generally more difficult than discriminative learning.

Perks of generative modeling

- Fitting complete multivariate distributions $p(\mathrm{x})$ goes beyond correlation-based or clustering approaches
- Correlations cannot discover partial correlation in the context of other neighbors
- Fitting $p(\mathrm{x})$ permits sampling based inference

Why generative modeling is difficult

When describing a distribution over multiple variables, we may not know the proper normalization Z. That is,

$$p(\mathrm{x}) = \frac{1}{Z}\tilde{p}(\mathrm{x})$$

This very important situation arises in several contexts:

1. In Bayesian inference where $p(x_1|x_2) = p(x_2|x_1)p(x_1)/p(x_2)$ is intractable due to $Z = p(x_2) = \int p(x_2|x_1)p(x_1)dx_1$. This integral can be very difficult or impossible to compute.

2. In models from statistical physics, e.g. the Ising model, we only know $\tilde{p}(\mathrm{x}) = e^{-H(\mathrm{x})}$ where $H(\mathrm{x})$ is the Hamiltonian

Variational autoencoders (VAEs)

A variational solution to generative modeling

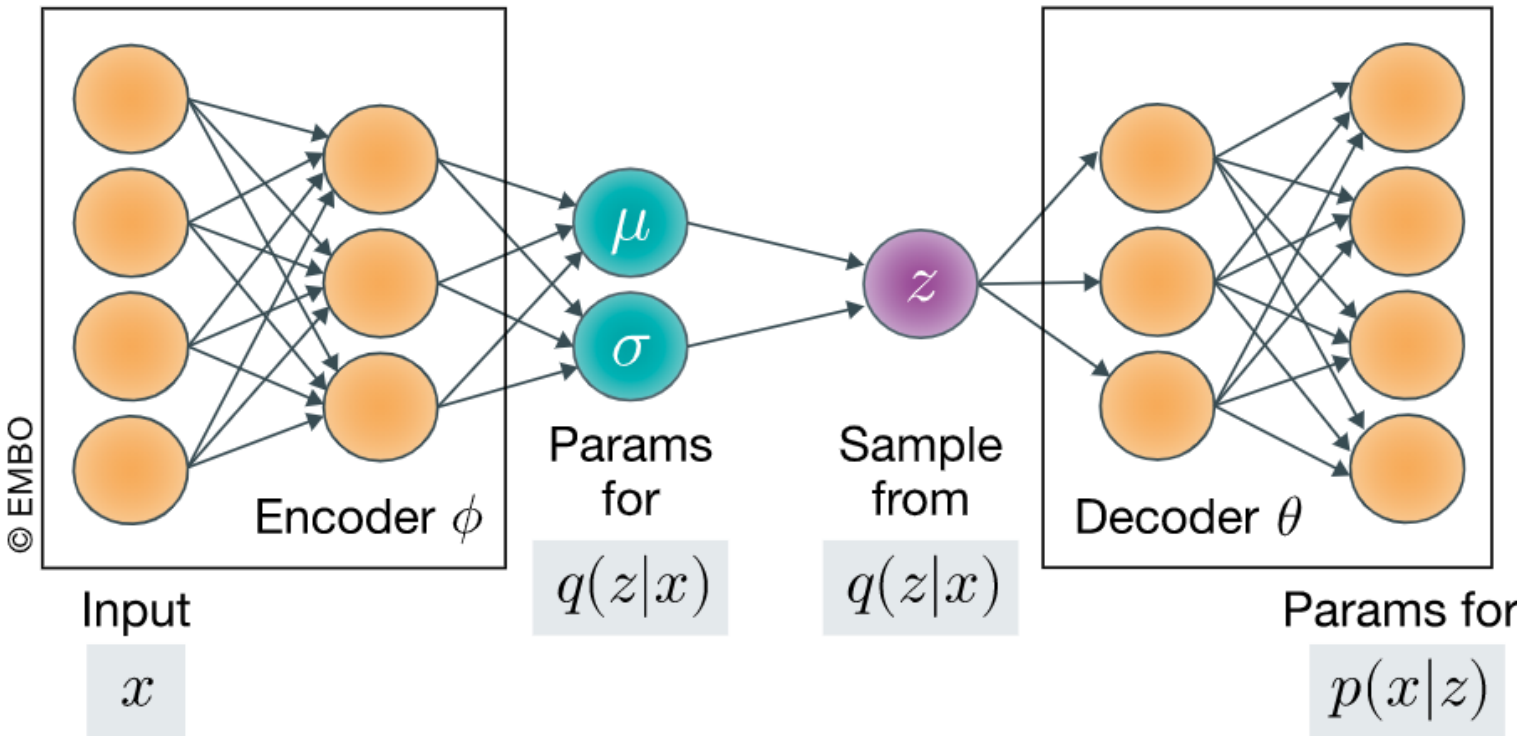

Figure 1: **Variational autoencoder architecture** Taken from Lopez 2020 in EMBO

Bayesian inference

The variable x has a latent representation or code z. We often say that z is the *causal source* of x. Ultimately, we would like to know the distribution $P_\Phi(\mathrm{x})$

$$P_\phi(\mathrm{x}) = \frac{P_\phi(\mathrm{x}|\mathrm{z})P_\phi(\mathrm{z})}{Q_\psi(\mathrm{z}|\mathrm{x})}$$

in order to find the model parameters that maximize the likelihood of the observed data:

$$\Phi^* = \underset{\Phi}{\operatorname{argmin}} - \log P_\Phi(\mathrm{x})$$

but we generally do not know $P_\psi(\mathrm{z}|\mathrm{x})$ due to the intractable integral $Z = \int P_\phi(\mathrm{x}|\mathrm{z})P_\phi(\mathrm{z})d\mathrm{z}$ (see slide 5)

Computing the evidence

We can rewrite the evidence as

$$
\begin{aligned}
P_\phi(\mathrm{x}) &= \int P_\phi(\mathrm{z}) P_\phi(\mathrm{x}|\mathrm{z}) d\mathrm{z} \\
&= \int P_\phi(\mathrm{z}) P_\phi(\mathrm{x}|\mathrm{z}) \frac{P_\phi(\mathrm{z}|\mathrm{x})}{P_\phi(\mathrm{z}|\mathrm{x})} d\mathrm{z} \\
&= \mathbb{E}_{\mathrm{z} \sim P_\phi(\mathrm{z}|\mathrm{x})} \frac{P_\phi(\mathrm{z}) P_\phi(\mathrm{x}|\mathrm{z})}{P_\phi(\mathrm{z}|\mathrm{x})}
\end{aligned}
$$

where $P_\phi(\mathrm{z}|\mathrm{x})$ is our model "encoder"

The evidence lower bound (ELBO)

$$
\begin{aligned}
\log P_\phi(\mathrm{x}) &= \log \int_z P(x,z)dx \\
&= \log \int_z P(x,z)\frac{Q(z|x)}{Q(z|x)}dz \\
&= \log \mathbb{E}_{\mathrm{z}\sim P_\phi(\mathrm{z}|\mathrm{x})}\frac{P(x|z)P(z)}{Q(z|x)} \\
&\geq \mathbb{E}_{\mathrm{z}\sim P_\phi(\mathrm{z}|\mathrm{x})} \log \frac{Q(x|z)}{P(z)} + \log P(x|z)
\end{aligned}
$$

$$
-\log P_\phi(\mathrm{x}) \leq \mathbb{E}_{\mathrm{z}\sim P_\phi(\mathrm{z}|\mathrm{x})} \log \frac{Q(x|z)}{P(z)} - \log P(x|z)
$$

The ELBO objective

$$\begin{aligned}\Phi^* &= \mathcal{L}(\Phi) \\ &= \operatorname*{argmin}_{\Phi} \; \mathbb{E}_{x \sim \mathrm{Pop},\, z \sim P_\phi(z|x)} \log \frac{Q_\psi(z|x)}{P(z)} - \log P(x|z)\end{aligned}$$

The ELBO can be rewritten in terms of a KL-divergence and population entropy. We often think of Φ as "position" and the loss $\mathcal{L}$ as an "energy"

Stochastic gradient descent

Applying deep generative models to biological data

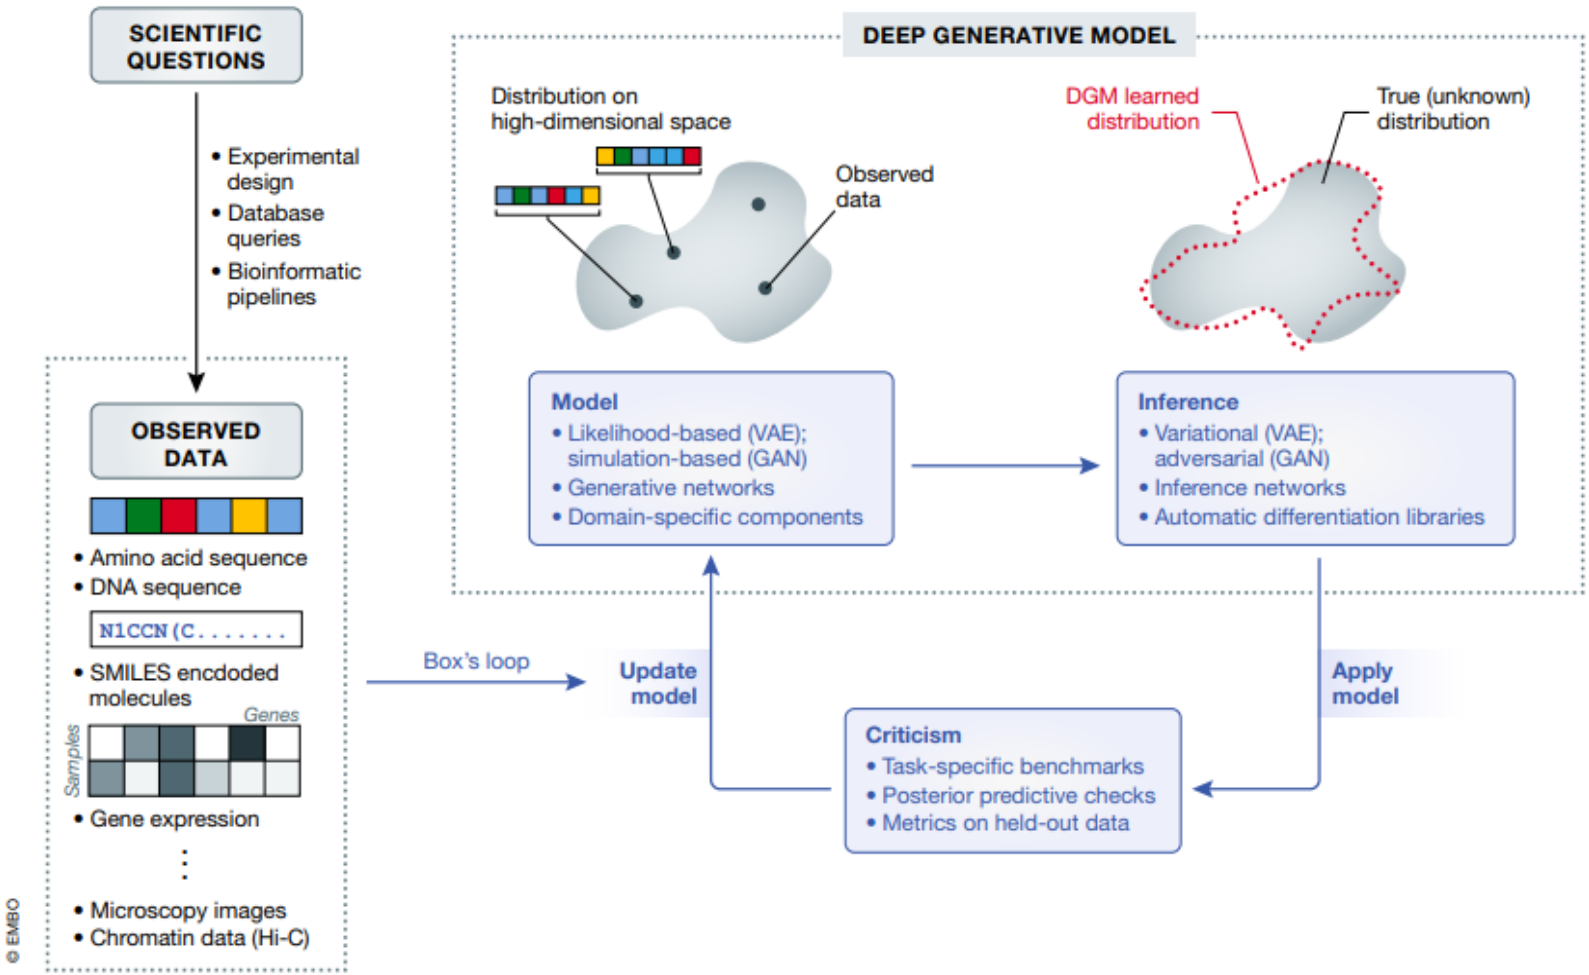

A recent VAE for transcriptomics data: scV1

Recent improvements of interpretability

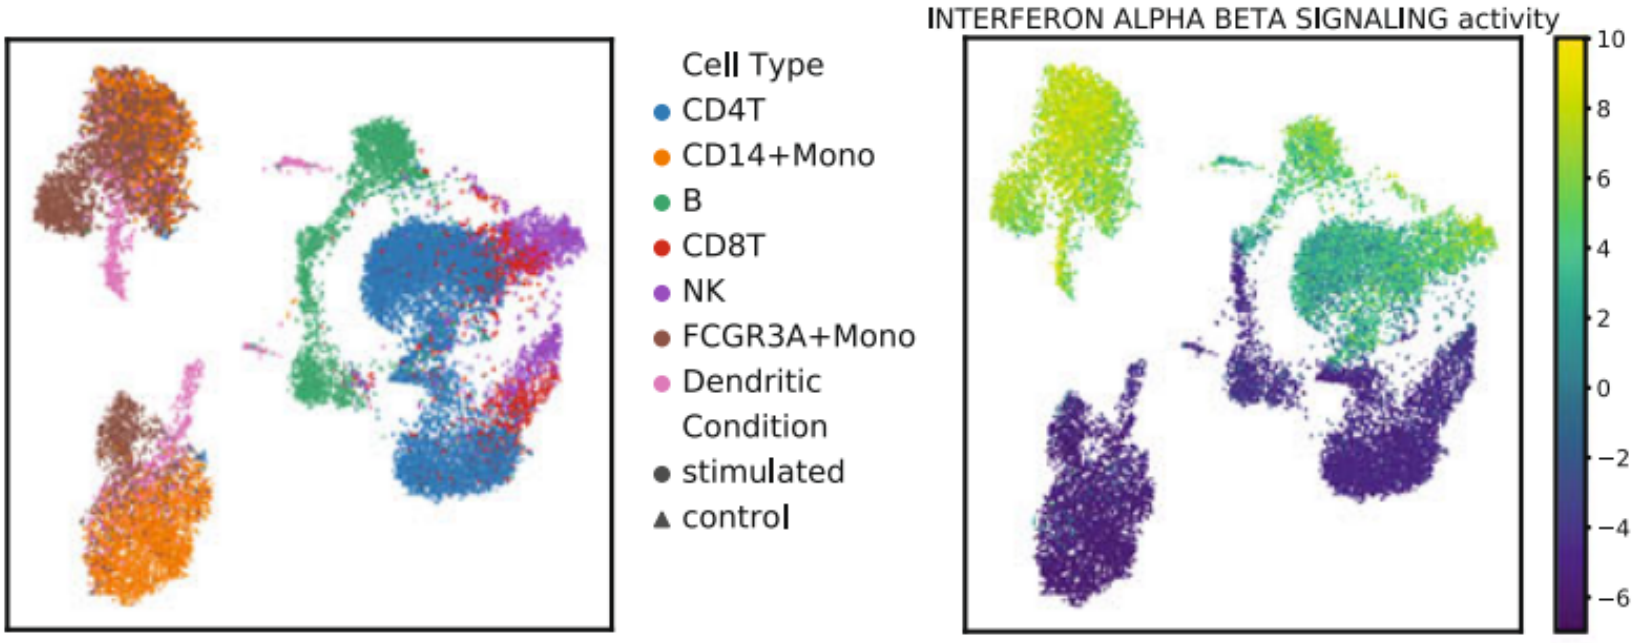

Figure 2: Phenotype segregation using a VAE on single-cell transcriptomics data. Taken from Seninge et al.

Previous studies lack interpretability

Previous studies lack spatial information

Multiplexed RNA imaging

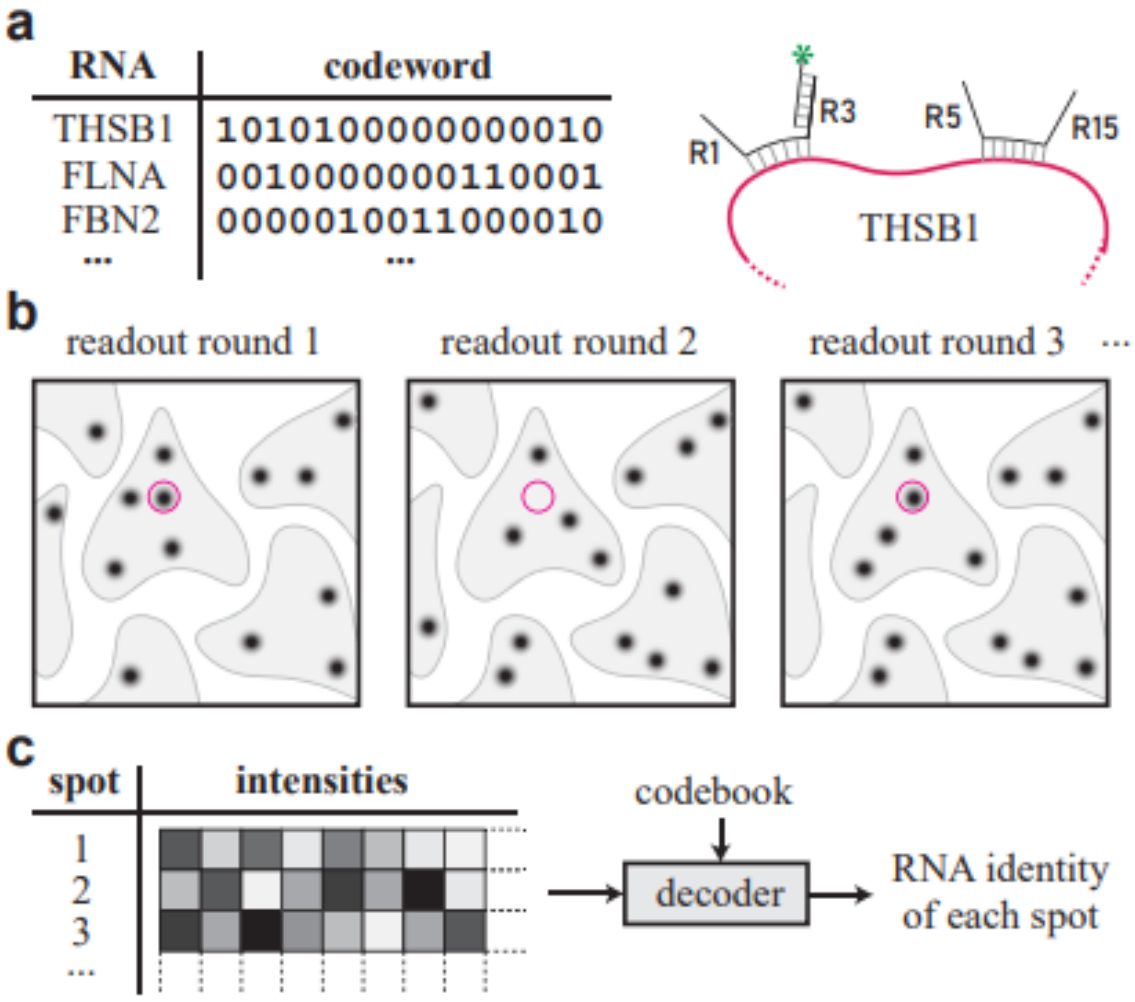

Figure 3:

Minimum Hamming Distance (MHD) codes

Very basic codes that generate codewords $x \in \mathcal{C}$ based on Hamming distance:

$$D = \sum_{n=1}^{L} \mathbb{I}(x_n = \hat{x}_n)$$

2^L possible code words (with $2^L - 1$ usable ones)

Binary symmetric channel (BSC)

References I