

TTIC 31230 Fundamentals of Deep Learning

AlphaZero Problems.

Background. A version of AlphaZero can be defined as follows.

To select a move in the game, first construct a search tree over possible moves to evaluate options.

The tree is grown by running “simulations”. Each simulation descends into the tree from the root selecting a move from each position until the selected move has not been explored before. When the simulation reaches an unexplored move it expands the tree by adding a node for that move. Each simulation returns the value $V_\Phi(s)$ for the newly added node s .

Each node in the search tree represents a board position s and stores the following information which can be initialized by running the value and policy networks on position s .

- $V_\Phi(s)$ — the value network value for the position s .
- For each legal move a from s , the policy network probability $\pi_\Phi(s, a)$.
- For each legal move a from s , the number $N(s, a)$ of simulations that have taken move a from s . This is initially zero.
- For each legal move a from s with $N(s, a) > 0$, the average $\hat{\mu}(s, a)$ of the values of the simulations that have taken move a from position s .

In descending into the tree, simulations select the move $\operatorname{argmax}_a U(s, a)$ where we have

$$U(s, a) = \begin{cases} \lambda_u \pi_\Phi(s, a) & \text{if } N(s, a) = 0 \\ \hat{\mu}(s, a) + \lambda_u \pi_\Phi(s, a) / N(s, a) & \text{otherwise} \end{cases} \quad (1)$$

When the search is completed, we must select a move from the root position. For this we use a post-search stochastic policy

$$\pi_{s_{\text{root}}}(a) \propto N(s_{\text{root}}, a)^\beta \quad (2)$$

where β is a temperature hyperparameter.

For training we construct

$$\text{a replay buffer of triples } (s_{\text{root}}, \pi_{s_{\text{root}}}, z) \quad (3)$$

accumulated from self play where s_{root} is a root position from a search during a game, $\pi_{s_{\text{root}}}$ is the post-search policy constructed for s_{root} , and z is the outcome of that game.

Training is then done by SGD on the following objective function.

$$\Phi^* = \underset{\Phi}{\operatorname{argmin}} E_{(s,\pi,z) \sim \text{Replay}, a \sim \pi} \begin{pmatrix} (V_\Phi(s) - z)^2 \\ -\lambda_1 \log \pi_\Phi(a|s) \\ +\lambda_2 \|\Phi\|^2 \end{pmatrix} \quad (4)$$

Problem 1. AlphaZero for BLEU Translation Score.

A version of AlphaZero can be defined as follows.

To select a move in the game, first construct a search tree over possible moves to evaluate options.

The tree is grown by running “simulations”. Each simulation descends into the tree from the root selecting a move from each position until the selected move has not been explored before. When the simulation reaches an unexplored move it expands the tree by adding a node for that move. Each simulation returns the value $V_\Phi(s)$ for the newly added node s .

Each node in the search tree represents a board position s and stores the following information which can be initialized by running the value and policy networks on position s .

- $V_\Phi(s)$ — the value network value for the position s .
- For each legal move a from s , the policy network probability $\pi_\Phi(s, a)$.
- For each legal move a from s , the number $N(s, a)$ of simulations that have taken move a from s . This is initially zero.
- For each legal move a from s with $N(s, a) > 0$, the average $\hat{\mu}(s, a)$ of the values of the simulations that have taken move a from position s .

In descending into the tree, simulations select the move $\operatorname{argmax}_a U(s, a)$ where we have

$$U(s, a) = \begin{cases} \lambda_u \pi_\Phi(s, a) & \text{if } N(s, a) = 0 \\ \hat{\mu}(s, a) + \lambda_u \pi_\Phi(s, a) / N(s, a) & \text{otherwise} \end{cases} \quad (1)$$

When the search is completed, we must select a move from the root position. For this we use a post-search stochastic policy

$$\pi_{s_{\text{root}}}(a) \propto N(s_{\text{root}}, a)^\beta \quad (2)$$

where β is a temperature hyperparameter.

For training we construct

$$\text{a replay buffer of triples } (s_{\text{root}}, \pi_{s_{\text{root}}}, z) \quad (3)$$

accumulated from self play where s_{root} is a root position from a search during a game, $\pi_{s_{\text{root}}}$ is the post-search policy constructed for s_{root} , and z is the outcome of that game.

Training is then done by SGD on the following objective function.

$$\Phi^* = \underset{\Phi}{\operatorname{argmin}} E_{(s, \pi, z) \sim \text{Replay}, a \sim \pi} \begin{pmatrix} (V_{\Phi}(s) - z)^2 \\ -\lambda_1 \log \pi_{\Phi}(a|s) \\ +\lambda_2 \|\Phi\|^2 \end{pmatrix} \quad (4)$$

Reformulate this algorithm to optimizing BLEU score in machine translation. More specifically,

(a) Define the optimization of the BLEU score as a tree search problem. What are the nodes and what are the moves?

Solution: The tree is defined by sequentially choosing words. Each node correspond to a sequences of moves — a node at depth d in the tree corresponds to a prefix of length d (the first d words) of a possible translation. The moves from a given node correspond to the choice of the next word.

(b) AlphaZero has three levels — complete games resulting in a final outcome z — move selection at each position in a complete game based on UTC algorithm — and simulations within the UTC algorithm. Describe how each of these can be interpreted in an algorithm for optimizing BLEU score in machine translation.

Solution: A “complete game” consists of selecting a move (word) at each position in the game ending one specific translation.

Move selection corresponds to selecting the next word using the UCT tree search algorithm.

We can use the same simulation algorithm as in AlphaZero. It might also be possible to let each simulation generate a full translation, although there is a danger of the simulation getting very deep.

(c) What do we take z to be in the replay buffer and in equation (4)?

Solution: z should be taken to be the BLEU score of the final translation generated in the “game”.

Problem 2. Tuning λ_u

(a) If we increase λ_u encourage more diversity or less diversity in the actions selected by simulations? Explain your answer.

Solution: Increasing λ_u leads to more diversity because larger values of λ_u cause $U(s, a)$ to decline as $N(s, a)$ increases leading to the simulation moving away from a as $N(s, a)$ increases.

(b) Consider the case of very large λ_u so that the term $\lambda_u \pi_\Phi(a|s)/N(s, a) \gg \hat{\mu}(s, a)$. Equivalently we can change the definition of $U(s, a)$ to be

$$U(s, a) = \frac{\pi_\Phi(a|s)}{\min(1, N(s, a))} \quad (5)$$

After running a some number of simulations from s define a^* by

$$a^* = \operatorname{argmax}_a U(s, a)$$

In other words a^* is the move that would be expanded in the next simulation visiting S . Consider a move a other than a^* . Give a lower bound on $N(s, a)$ in terms of $U(s, a^*)$ and $\pi_\Phi(s, a)$ where $U(s, a)$ is defined by (5).

Solution:

$$U(s, a^*) \geq U(s, a)$$

$$U(s, a^*) \geq \frac{\pi_\Phi(a|s)}{N(s, a)}$$

$$N(s, a) \geq \frac{\pi_\Phi(s|a)}{U(s, a^*)} \quad (6)$$

Note that (6) holds for all a . So for very large λ_u the number of simulations visiting an action a is proportional to $\pi_\Phi(a|s)$ for some top k actions.

Problem 2. Replacing the Policy with a Q -function. We consider replacing the policy network π_Φ with a Q -value network Q_Φ so that each node s stores the Q -values $Q_\Phi(s, a)$ rather than the policy probabilities $\pi_\Phi(s, a)$. We then replace (1) with

$$U(s, a) = \begin{cases} \lambda_u Q_\Phi(s, a) & \text{if } N(s, a) = 0 \\ \hat{\mu}(s, a) + \lambda_u Q_\Phi(s, a)/N(s, a) & \text{otherwise} \end{cases} \quad (1')$$

and leave (2) and (3) unchanged. Rewrite (4) to train $Q_\Phi(s, a)$ by minimizing a squared “Bellman Error” between $Q_\Phi(s, a)$ and the outcome z over actions drawn from the replay buffer’s stored policy. Presumably this version does not work as well.

Solution:

$$\Phi^* = \operatorname{argmin}_{\Phi} E_{(s, \pi, z) \sim \text{Replay}, a \sim \pi} \begin{pmatrix} (V_\Phi(s) - z)^2 \\ -\lambda_1 (Q(s, a) - z)^2 \\ +\lambda_2 \|\Phi\|^2 \end{pmatrix}$$

Problem 3. An advantage actor-critic version. We consider replacing the policy network π_Φ with an advantage network A_Φ so that each node s stores the A -values $A_\Phi(s, a)$ rather than the policy probabilities $\pi_\Phi(s, a)$. We now have each node s also store $\hat{\mu}(s)$ which equals the average value of the simulations that go through state s . We then replace (1) with

$$U(s, a) = \begin{cases} \lambda_u(A_\Phi(s, a) + V_\Phi(s)) & \text{if } N(s, a) = 0 \\ \hat{\mu}(s, a) + \lambda_u(A_\Phi(s, a) + V_\Phi(s))/N(s, a) & \text{otherwise} \end{cases} \quad (1'')$$

and leave (2) unchanged and replace (3) by

$$\text{a replay buffer of tuples } (s_{\text{root}}, \pi_{s_{\text{root}}}, z, \hat{\mu}(s_{\text{root}}), \hat{\mu}(s_{\text{root}}, a)) \quad (3')$$

Rewrite (4) to train $A_\Phi(s, a)$ by minimizing a squared “Bellman Error” between $A_\Phi(s, a)$ and $\hat{A}(s, a)$ defined as follows

$$\hat{A}(s, a) = \begin{cases} A_\Phi(s, a) & \text{if } N(s, a) = 0 \\ \hat{\mu}(s, a) - \hat{\mu}(s) & \text{otherwise} \end{cases} \quad (5)$$

Solution:

$$\Phi^* = \underset{\Phi}{\operatorname{argmin}} E_{(s, \pi, R) \sim \text{Replay}, a \sim \pi} \left(\begin{array}{l} (V_\Phi(s) - z)^2 \\ \lambda_A (A_\Phi(s, a) - \hat{A}(s, a))^2 \\ + \lambda_R \|\Phi\|^2 \end{array} \right) \quad (2)$$

Although a valiant attempt, this version also presumably also does not work as well.