

TTIC 31230, Fundamentals of Deep Learning

David McAllester, Winter 2020

Representing Functions with Programs

Python, Assembler, and the Turing Tarpit

Representing Functions with Programs

Neural Turing Machines Alex Graves, Greg Wayne, Ivo Danihelka, 2014

(Actually a differentiable Von Neumann architecture)

The machine undergoes continuous state, discrete time, state transitions defined a differentiable feed-forward circuit.

Compositional Attention Networks for Machine Reasoning

Hudson and Manning, ICLR 2018

The MAC cell is similar to a gated RNN cell used as the decoder in translation.

It is also similar to a Neural Turing Machine.

It was applied to image-based question answering and uses attention over the image and the question during multi-step “decoding”.

What about Python?

High level scripting languages such as Python seem to be the most productive programming languages for human programmers.

Does Python represent a particularly effective universal learning bias?

Productivity in programming seems to be greatly enhanced by functional expressions (functional programming) and object-oriented programming (objects, classes and inheritance).

This seems crucial if we want to somehow achieve I. J. Good's intelligence explosion.

The Turing Tarpit

But in theory the choice of programming language does not matter.

For any two Turing universal languages, say Python and Assembler, there exists an interpreter I for Python written in Assembler where we write $I(h)$ for the assembler interpreter I applied to Python program h . We then get

$$|I(h)|_{\text{Assembler}} = |h|_{\text{Python}} + |I|_{\text{Assembler}}$$

Bootstrapping layers of language can make the interpreter small.

The Turing Tarpit

$$|I(h)|_{\text{Assembler}} = |h|_{\text{Python}} + |I|_{\text{Assembler}}$$

Up to the additive constant of the interpreter, assembler gives just as good a learning bias as Python.

Yet we know that the choice of language does matter — Python is clearly better than assembler.

END